

Tracking learning

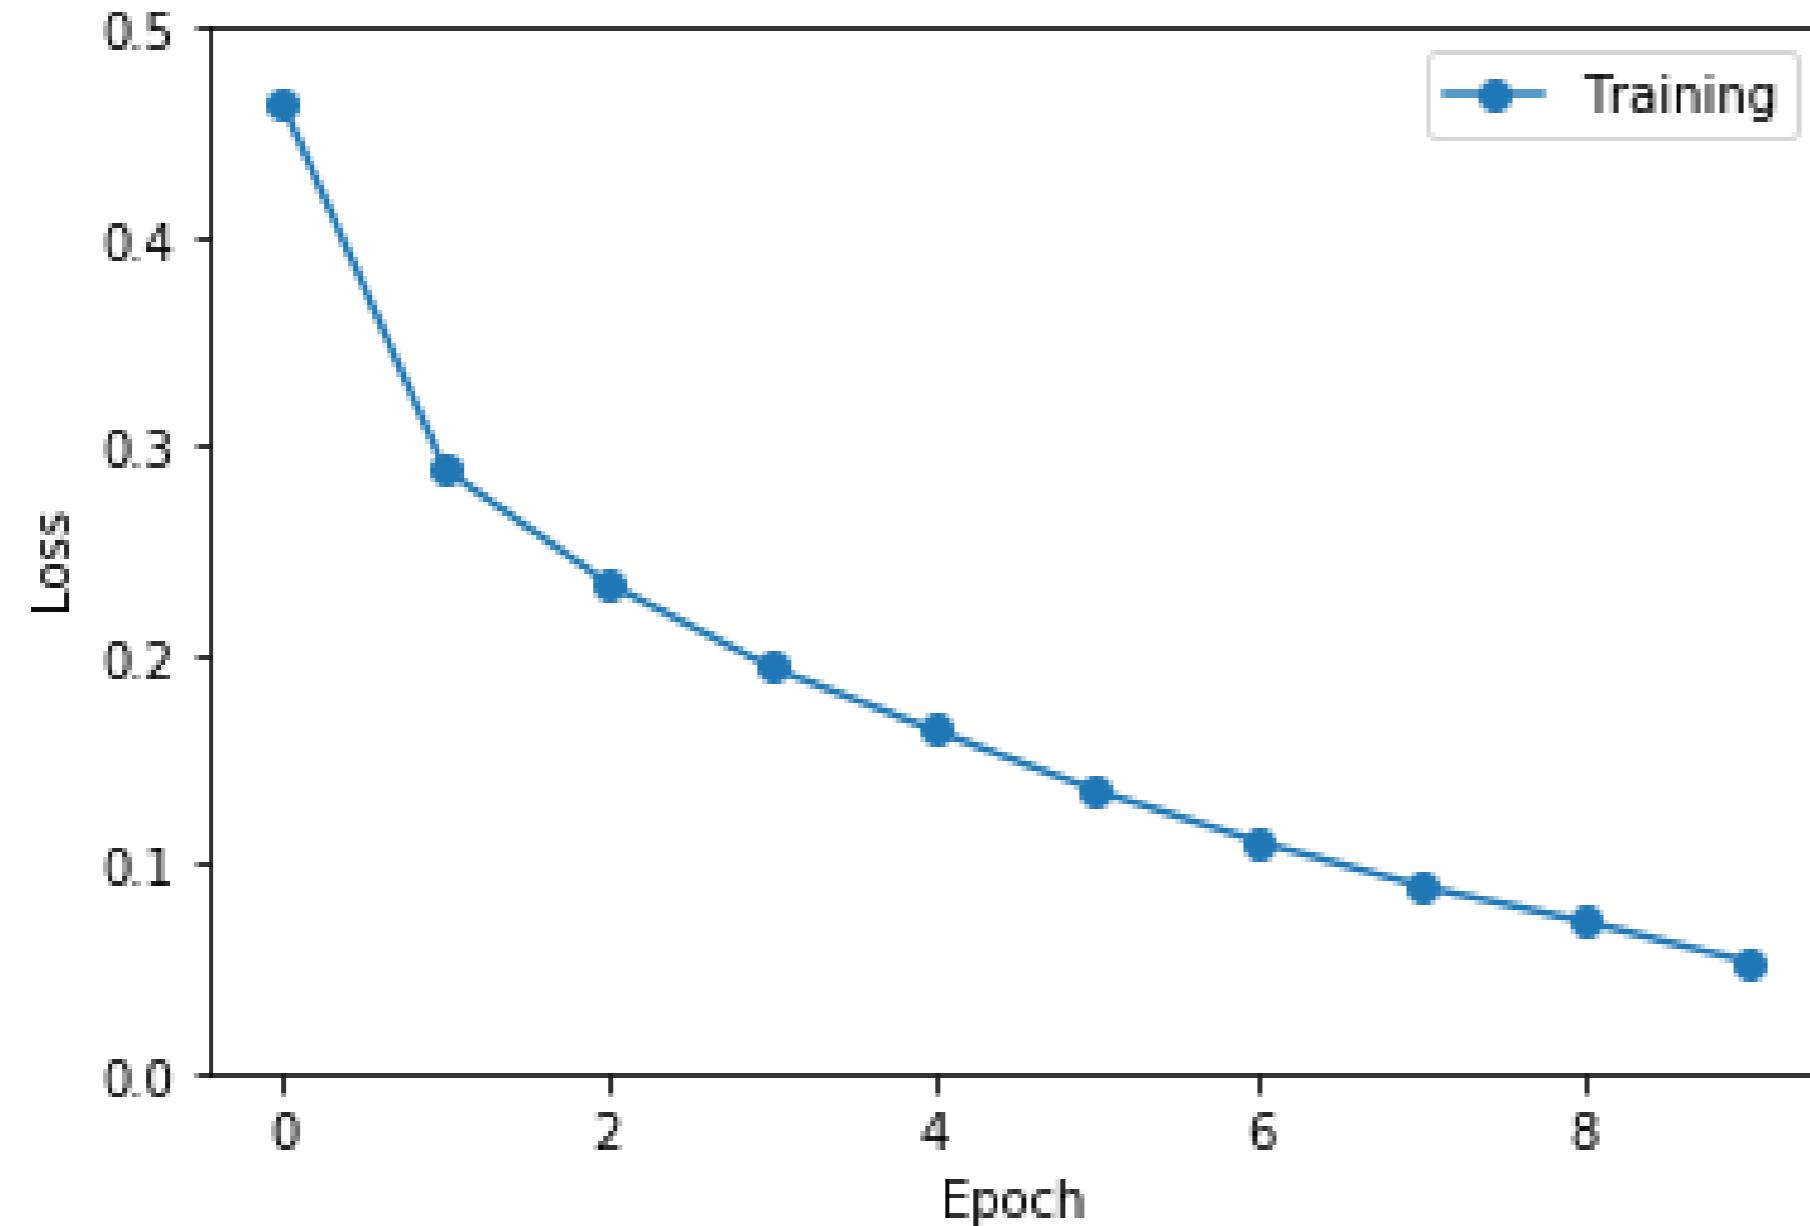
IMAGE MODELING WITH KERAS




Ariel Rokem

Senior Data Scientist, University of
Washington


Learning curves: training




Learning curves: validation



Learning curves: overfitting



```
training = model.fit(train_data, train_labels,  
                      epochs=3, validation_split=0.2)  
  
import matplotlib.pyplot as plt  
plt.plot(training.history['loss'])  
plt.plot(training.history['val_loss'])  
plt.show()
```



Storing the optimal parameters

```
from keras.callbacks import ModelCheckpoint

# This checkpoint object will store the model parameters
# in the file "weights.hdf5"
checkpoint = ModelCheckpoint('weights.hdf5', monitor='val_loss',
                             save_best_only=True)

# Store in a list to be used during training
callbacks_list = [checkpoint]
# Fit the model on a training set, using the checkpoint as a
#callback
model.fit(train_data, train_labels, validation_split=0.2,
          epochs=3, callbacks=callbacks_list)
```

Loading stored parameters

```
model.load_weights('weights.hdf5')  
model.predict_classes(test_data)  
array([2, 2, 1, 2, 0, 1, 0, 1, 2, 0])
```

Let's practice!

IMAGE MODELING WITH KERAS

Neural network regularization

IMAGE MODELING WITH KERAS



Ariel Rokem


Senior Data Scientist, University of
Washington

Dropout


In each learning step:

- Select a subset of the units
- Ignore it in the forward pass
- And in the back-propagation of error

Dropout



(a) Standard Neural Net



(b) After applying dropout.

Dropout in Keras

```
from keras.models import Sequential
from keras.layers import Dense, Conv2D, Flatten, Dropout
model = Sequential()
model.add(Conv2D(5, kernel_size=3, activation='relu',
                 input_shape=(img_rows, img_cols, 1)))
model.add(Dropout(0.25))
model.add(Conv2D(15, kernel_size=3, activation='relu'))
model.add(Flatten())
model.add(Dense(3, activation='softmax'))
```

Batch normalization

- Rescale the outputs

Batch Normalization in Keras

```
from keras.models import Sequential
from keras.layers import Dense, Conv2D, Flatten, BatchNormalization

model = Sequential()
model.add(Conv2D(5, kernel_size=3, activation='relu',
                input_shape=(img_rows, img_cols, 1)))
model.add(BatchNormalization())
model.add(Conv2D(15, kernel_size=3, activation='relu'))
model.add(Flatten())
model.add(Dense(3, activation='softmax'))
```

Be careful when using them together!

The disharmony between dropout and batch normalization

Let's practice!

IMAGE MODELING WITH KERAS

Interpreting the model

IMAGE MODELING WITH KERAS



Ariel Rokem

Senior Data Scientist, University of
Washington

Selecting layers

```
model.layers
```

```
[<keras.layers.convolutional.Conv2D at 0x109f10c18>,
 <keras.layers.convolutional.Conv2D at 0x109ec5ba8>,
 <keras.layers.core.Flatten at 0x1221ffcc0>,
 <keras.layers.core.Dense at 0x1221ffff0>]
```

Getting model weights

```
conv1 = model.layers[0]  
weights1 = conv1.get_weights()  
len(weights1)
```

2

```
kernels1 = weights1[0]  
kernels1.shape
```


(3, 3, 1, 5)

```
kernel1_1 = kernels1[:, :,  
                      0, 0]  
kernel1_1.shape
```

(3, 3)


Visualizing the kernel

```
plt.imshow(kernel1_1)
```




Visualizing the kernel responses

```
test_image = test_data[3, :, :, 0]  
plt.imshow(test_image)
```




Visualizing the kernel responses

```
filtered_image = convolution(test_image, kernel1_1)  
plt.imshow(filtered_image)
```




Visualizing the kernel responses

```
test_image = test_data[4, :, :, 1]  
plt.imshow(test_image)
```




Visualizing the kernel responses

```
filtered_image = convolution(test_image, kernel1_1)  
plt.imshow(filtered_img)
```



Visualizing the kernel responses

```
kernel1_2 = kernels[:, :, 0, 1]  
filtered_image = convolution(test_image, kernel1_2)  
plt.imshow(filtered_img)
```



Let's practice!

IMAGE MODELING WITH KERAS

Wrapping up

IMAGE MODELING WITH KERAS



Ariel Rokem

Senior Data Scientist, University of
Washington

What did we learn?


- Image classification
- Convolutions
- Reducing the number of parameters
 - Tweaking your convolutions
 - Adding pooling layers
- Improving your network
 - Regularization
- Understanding your network
 - Monitoring learning
 - Interpreting the parameters

Model interpretation


<https://distill.pub/2017/feature-visualization/>



Feature visualization answers questions about what a network—or parts of a network—are looking for by generating examples.



Attribution¹ studies what part of an example is responsible for the network activating a particular way.

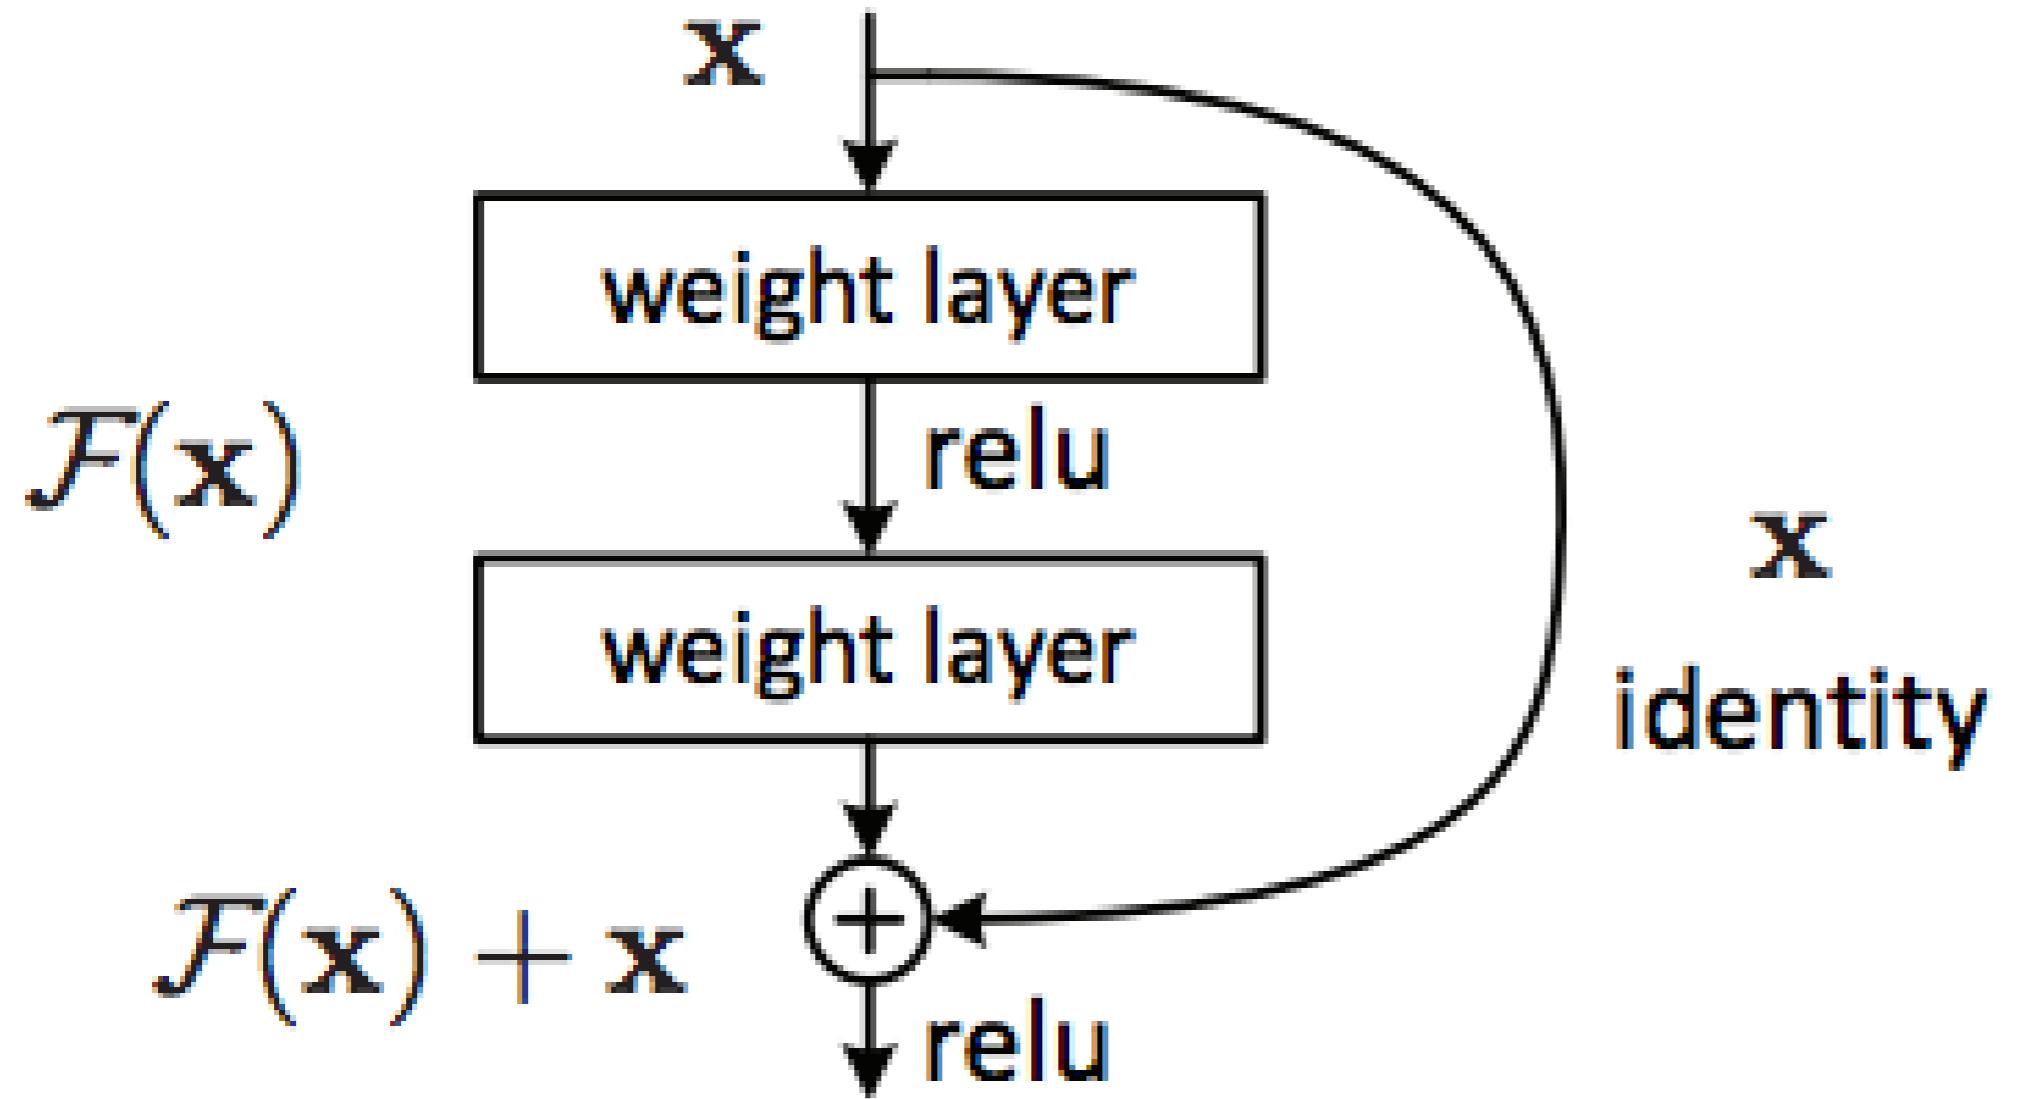




What next?

- Even deeper networks
- Residual networks


Residual networks




What next?

- Even deeper networks
- Residual networks
- Transfer learning
- Fully convolutional networks

Fully convolutional networks



Generative adversarial networks



What next?

- Even deeper networks
- Residual networks
- Transfer learning
- Fully convolutional networks
- Generative adversarial networks
- ...

Good luck!

IMAGE MODELING WITH KERAS