Monitoring ML services

MLOPS DEPLOYMENT AND LIFE CYCLING

Nemanja Radojkovic
Senior Machine Learning Engineer

datacamp

Maintaining quality

- Paying customers == expectations of quality
- Quality assurance starts with quality control
- Monitoring

Performance indicators

Fundamental health indicators

- Service up and running?
- Number of requests in time?
- Latency distribution?

Ultimate quality metric

- Predictive performance

How do ML models deteriorate?

Features and labels at training time

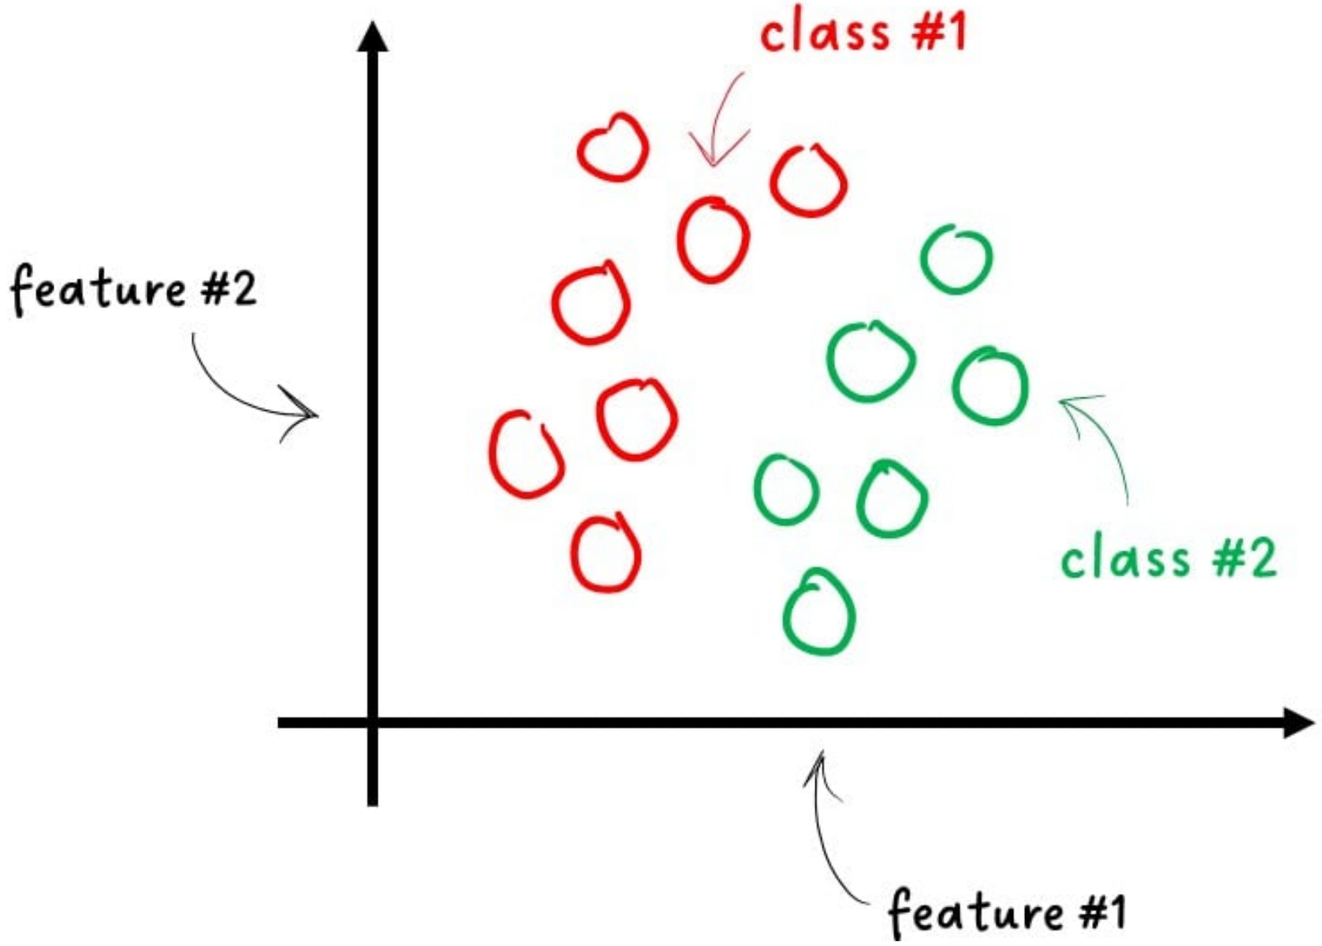

Features and labels at training time

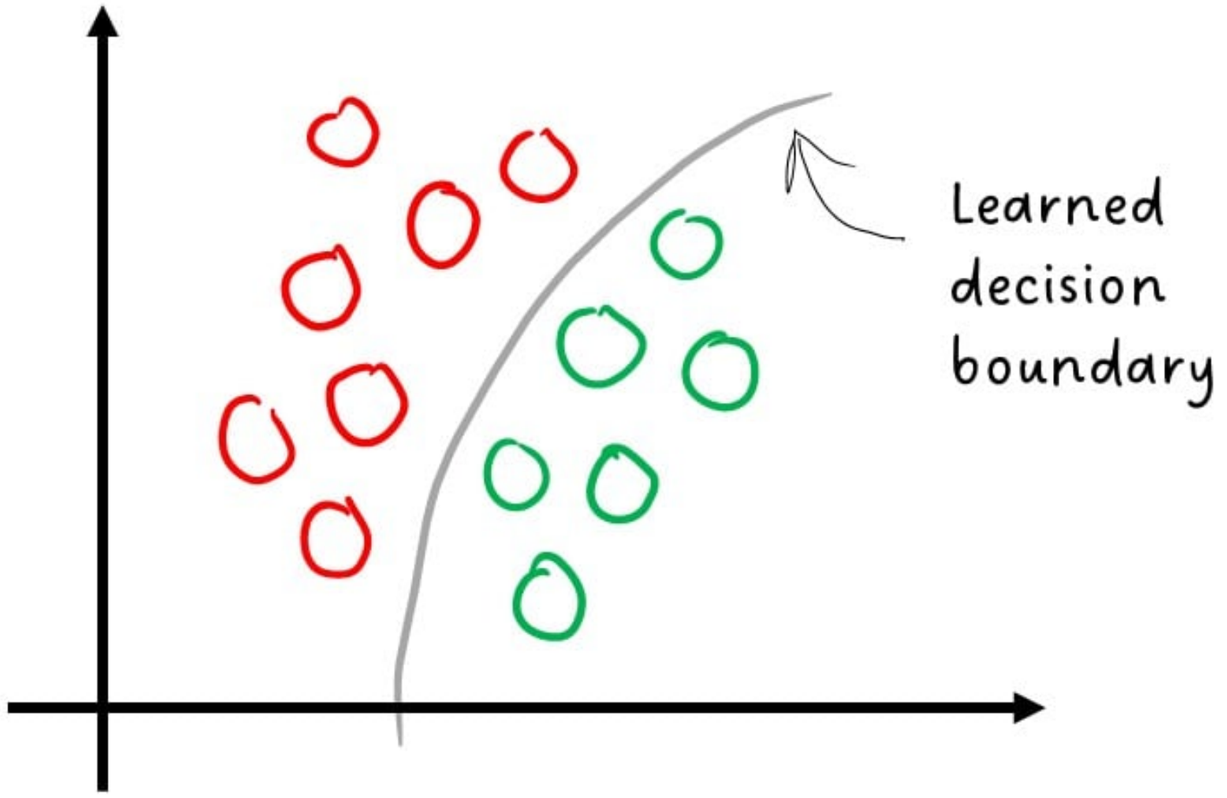

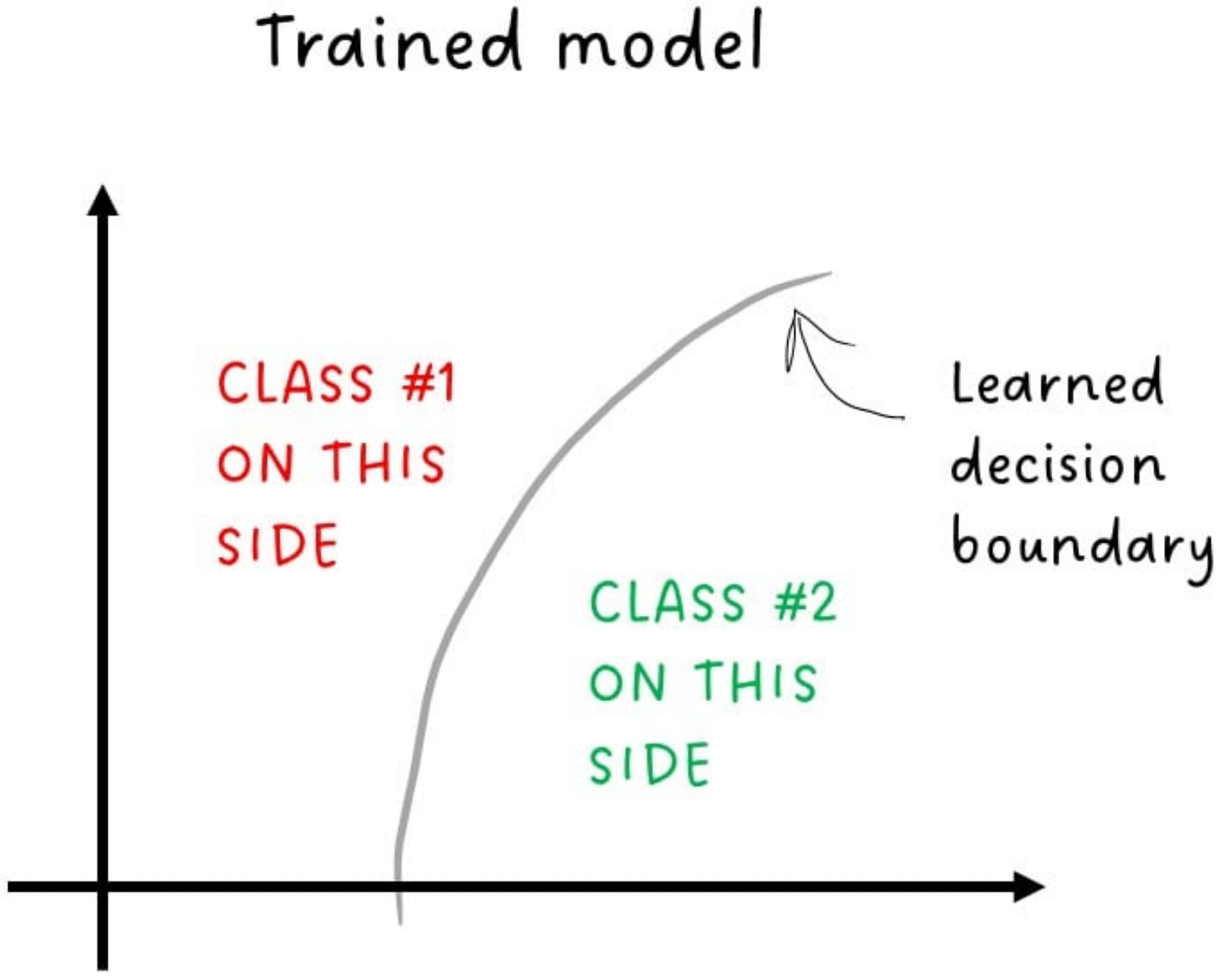
Trained model
CLASS #1
ON THIS
SIDE
CLASS #2
ON THIS
SIDE
Learned
decision
boundary

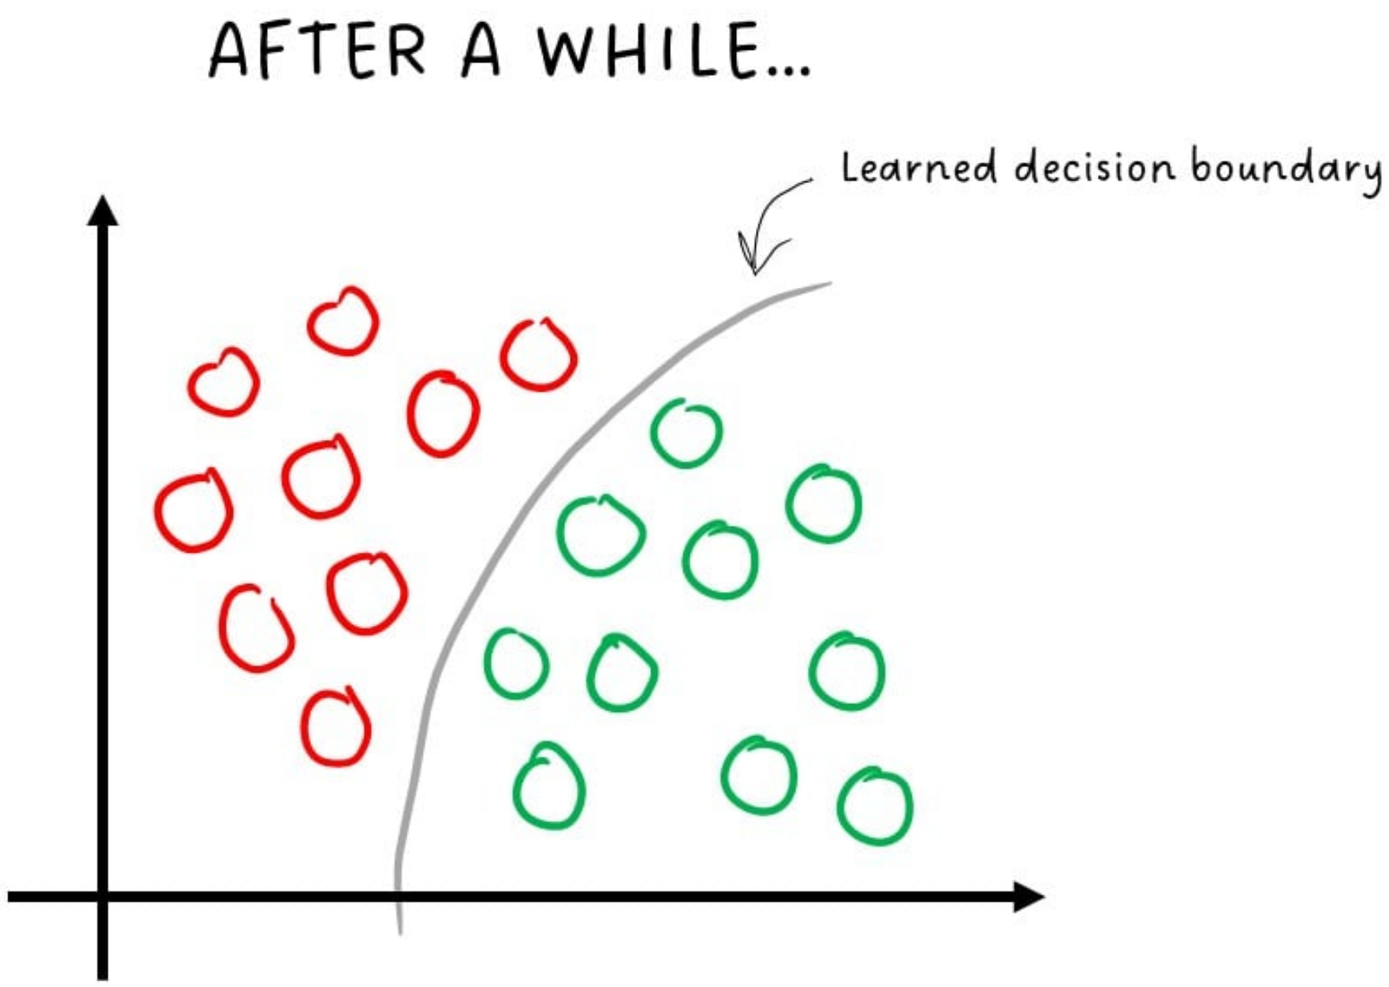
AFTER A WHILE...
Learned decision boundary

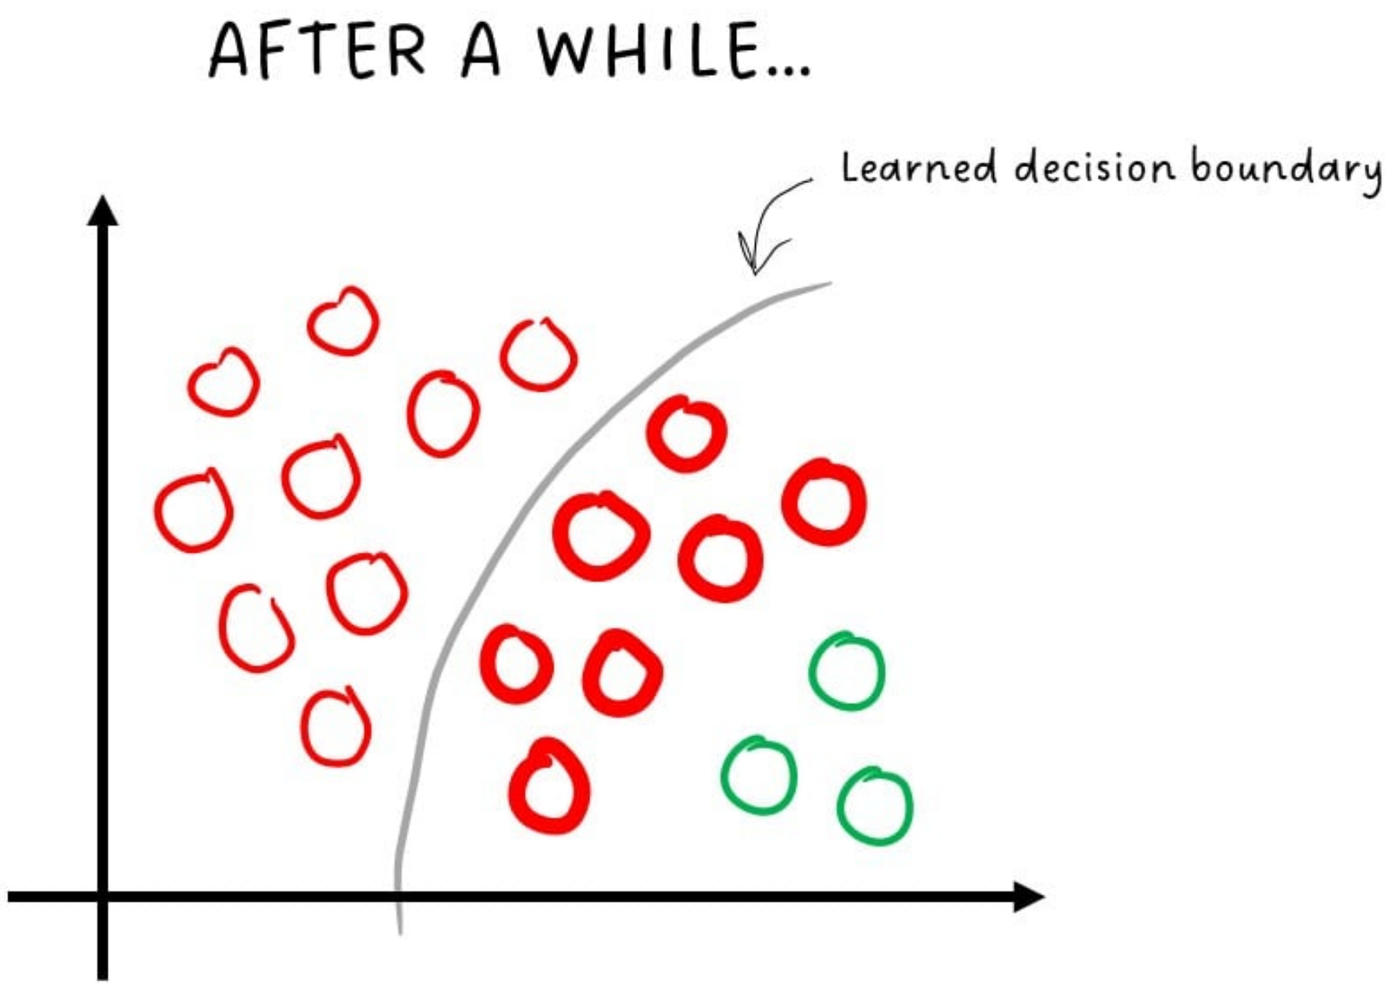
AFTER A WHILE...
Learned decision boundary

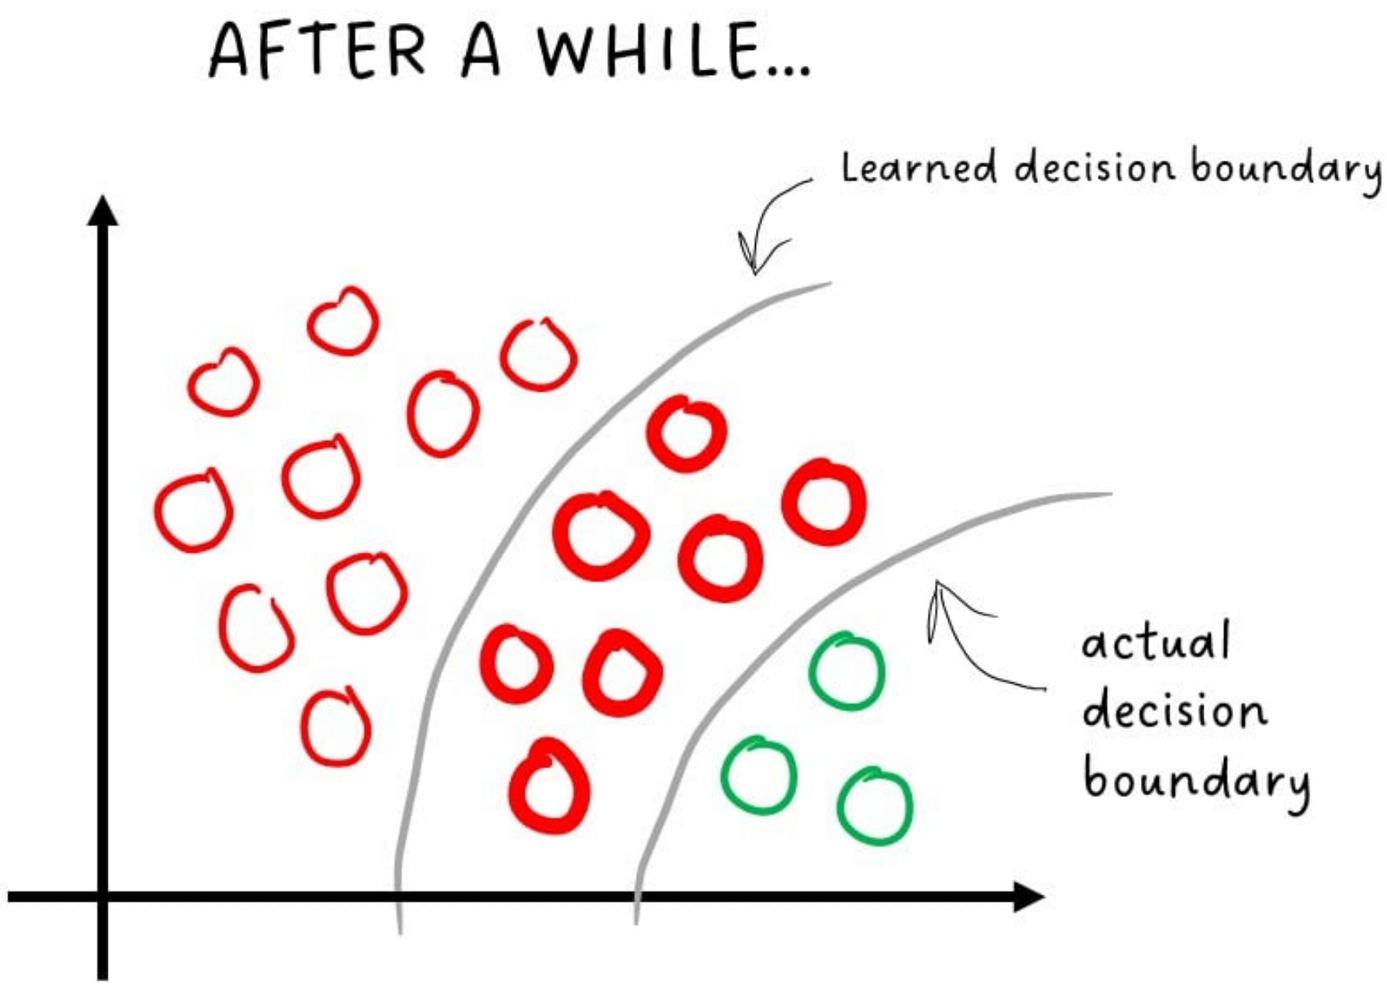
AFTER A WHILE...
Learned decision boundary
actual
decision
boundary

AFTER A WHILE...

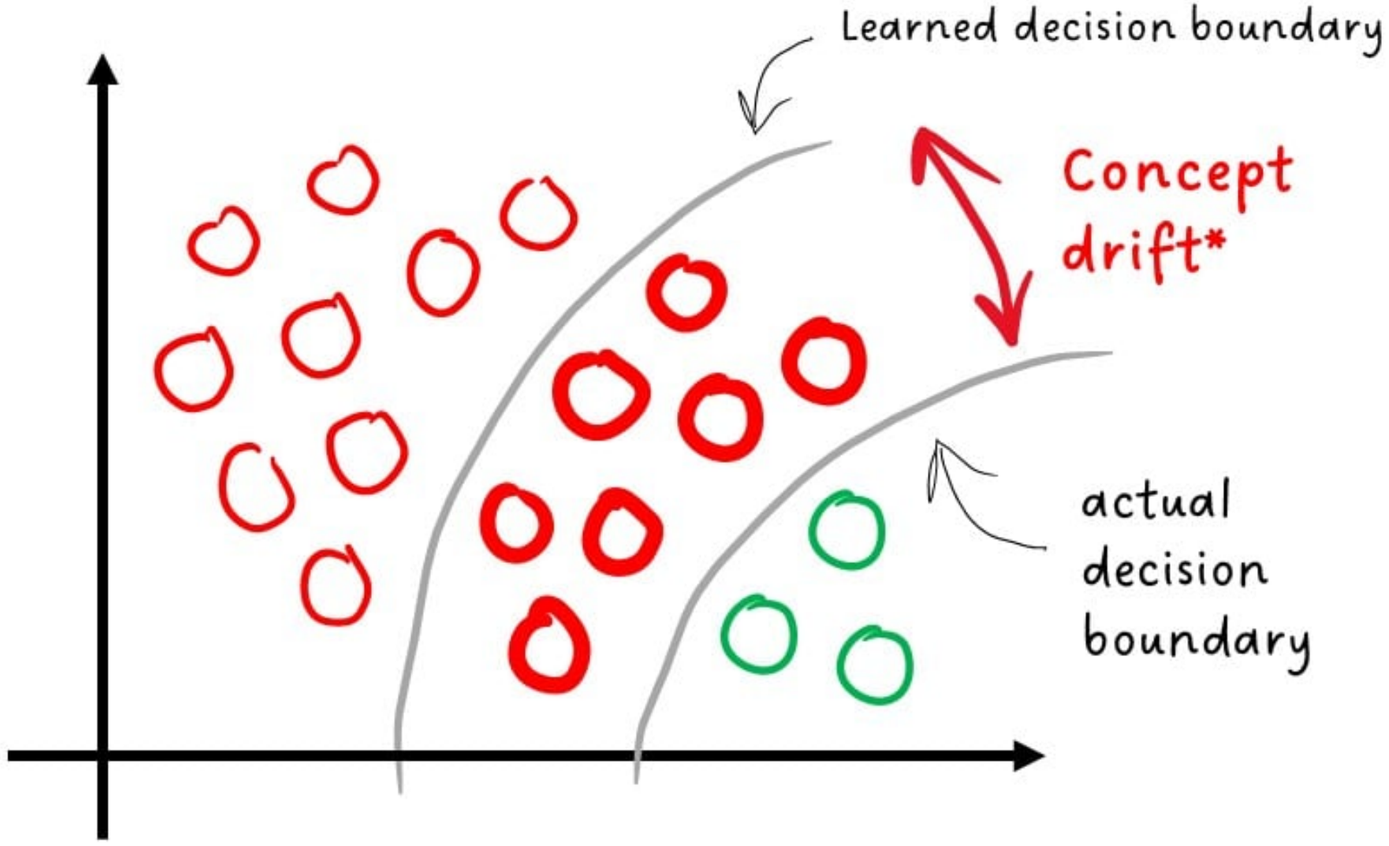

*concept drift == significant change in the actual relationship between the input and output features

AFTER A WHILE...

The model "deteriorates" by staying the same, while the reality changes!

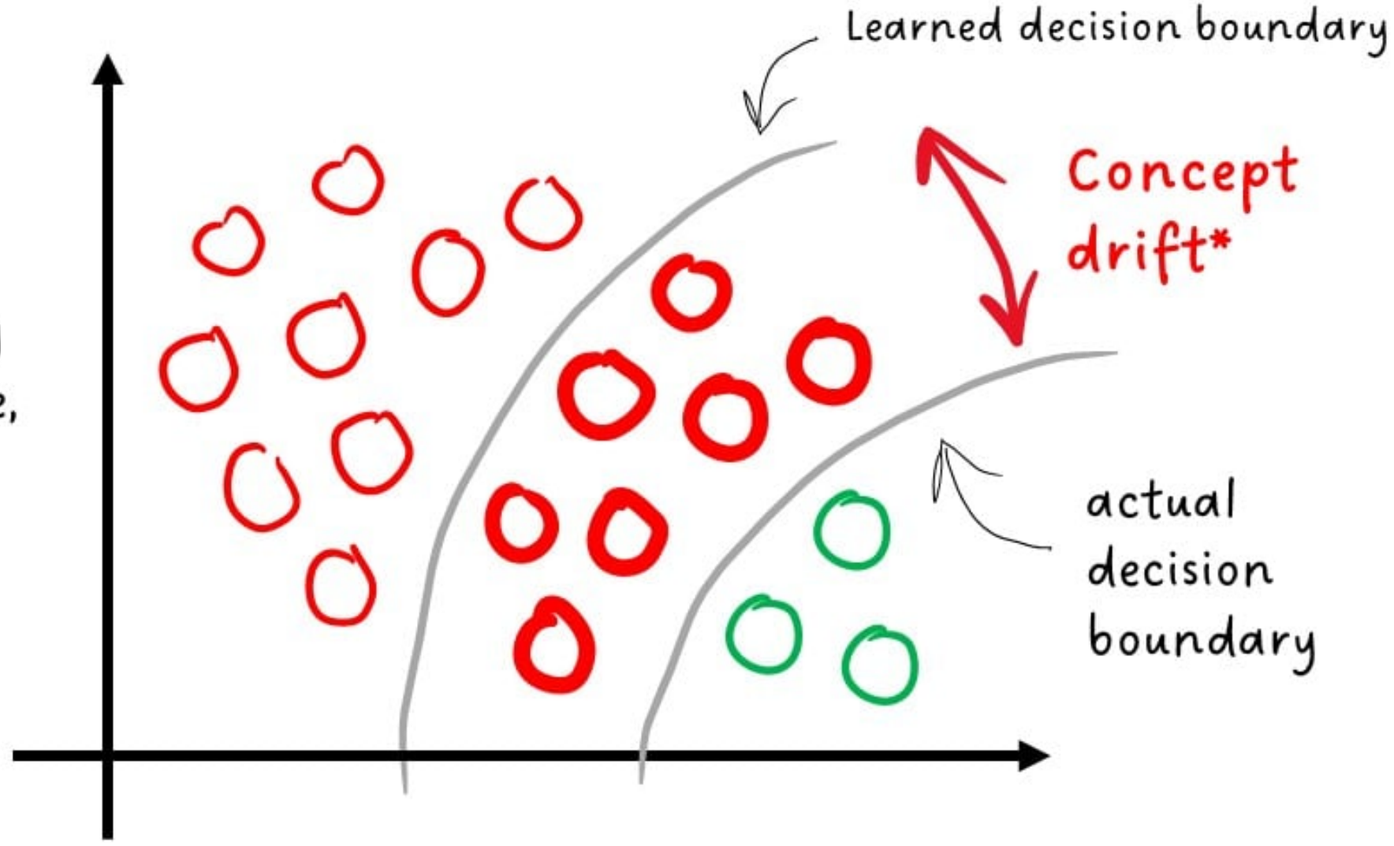

*concept drift == significant change in the actual relationship between the input and output features

How to detect concept drift?

How to detect concept drift?

```
IF prediction != ground_truth
MORE_OFTEN_THAN_EXPECTED
THEN WARNING("CONCEPT DRIFT!")
```

The catch:

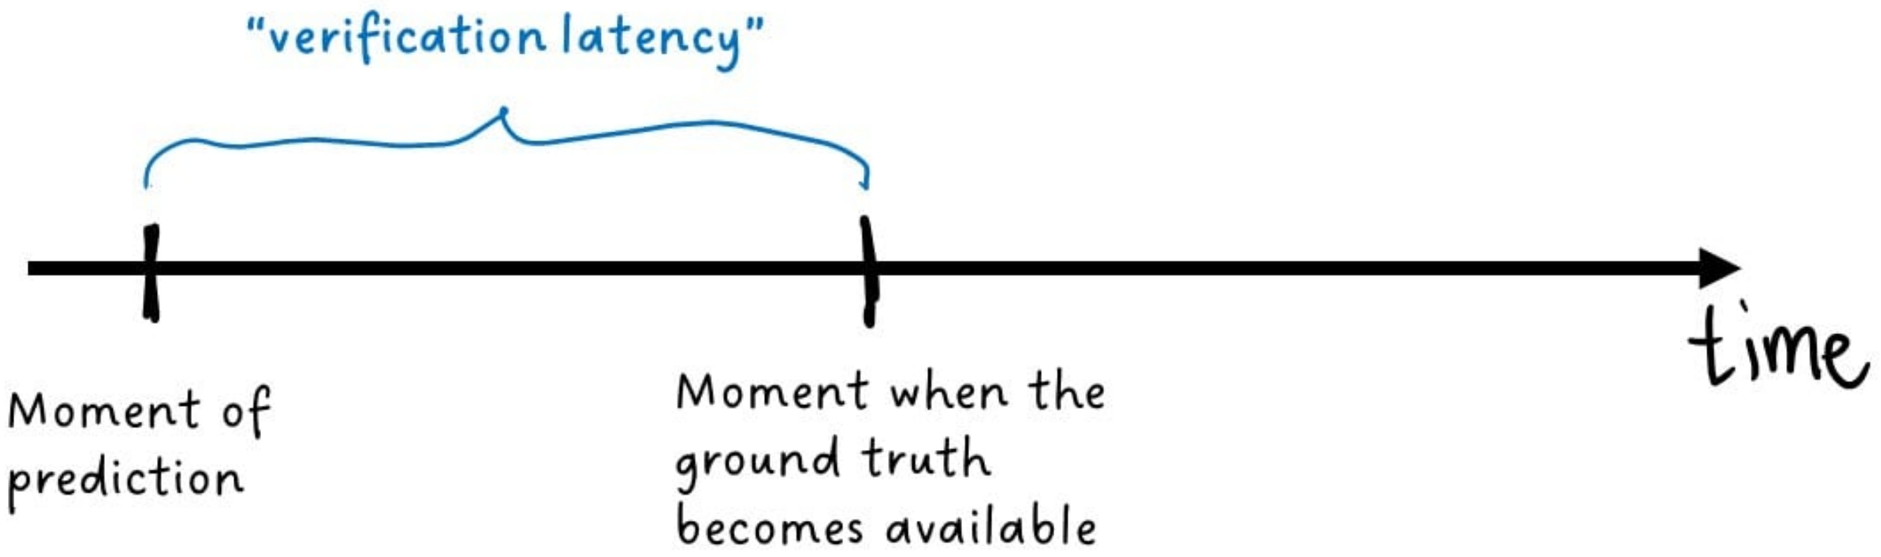
"verification latency"
time
Moment of prediction
Moment when the ground truth becomes available

The catch:

Stock prediction: < 1 second
Fraud detection: ~ months
Credit scoring: ∞

If a loan is **denied** based on our model, we will **never** know if the decision was in fact wrong!

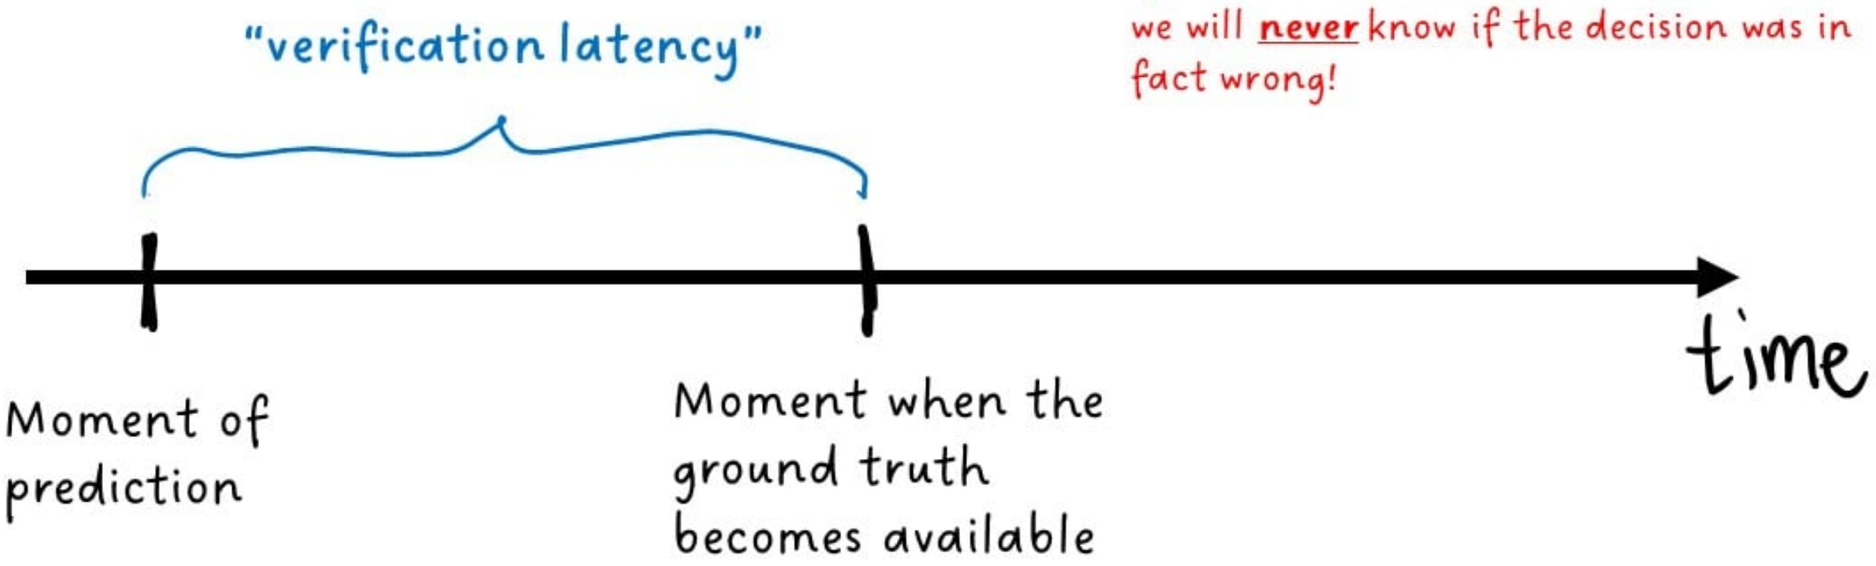

The catch:

Stock prediction: < 1 second
Fraud detection: ~ months
Credit scoring: ∞

If a loan is **denied** based on our model, we will **never** know if the decision was in fact wrong!

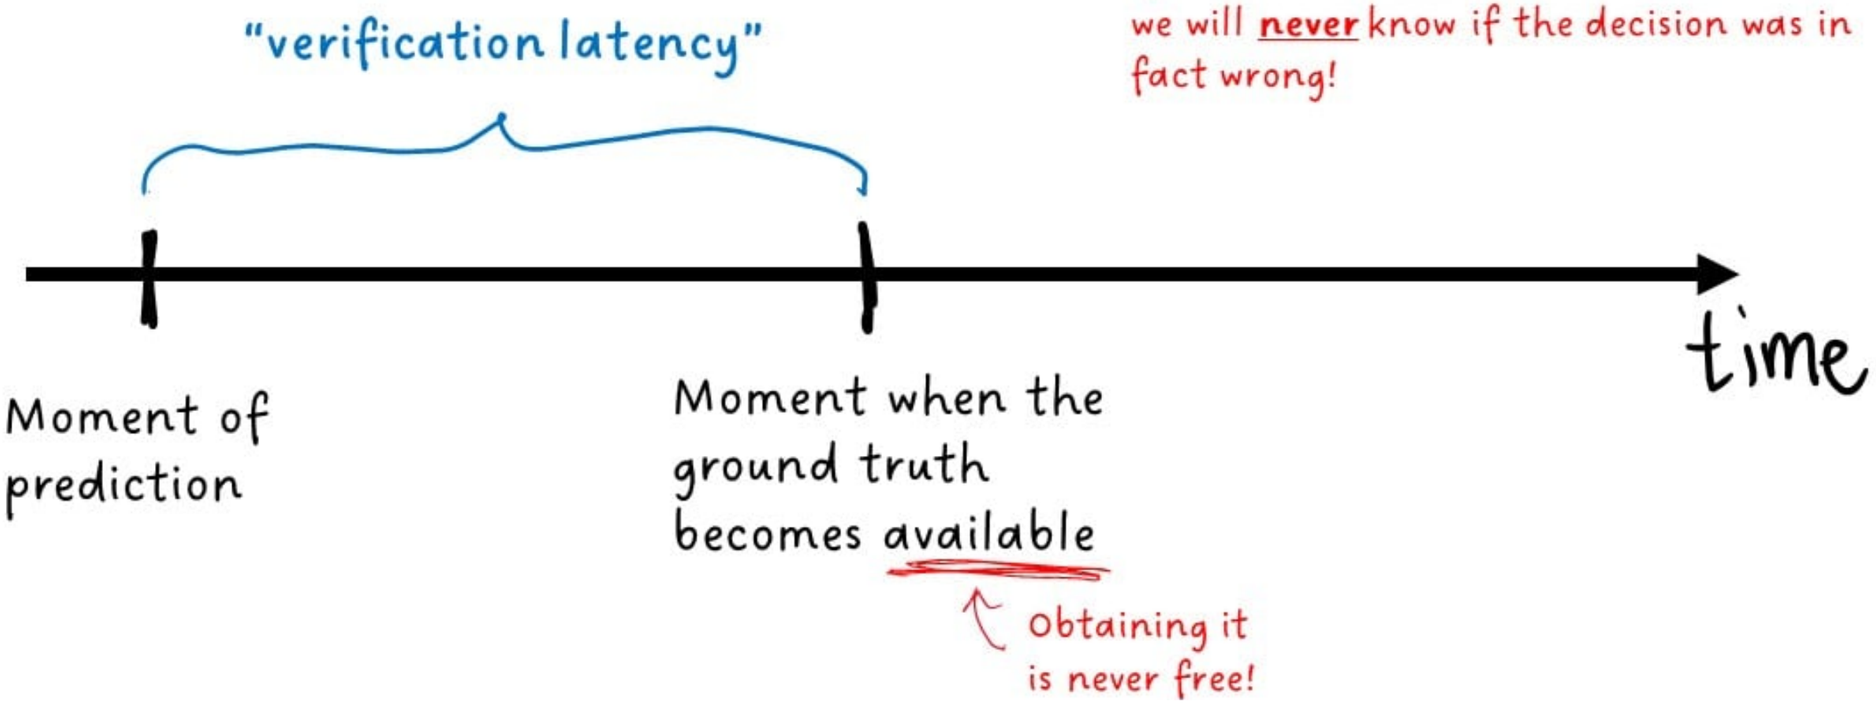

Indirect approach:

Monitor what you have: input features!

Indirect approach:

Monitor what you have: input features!

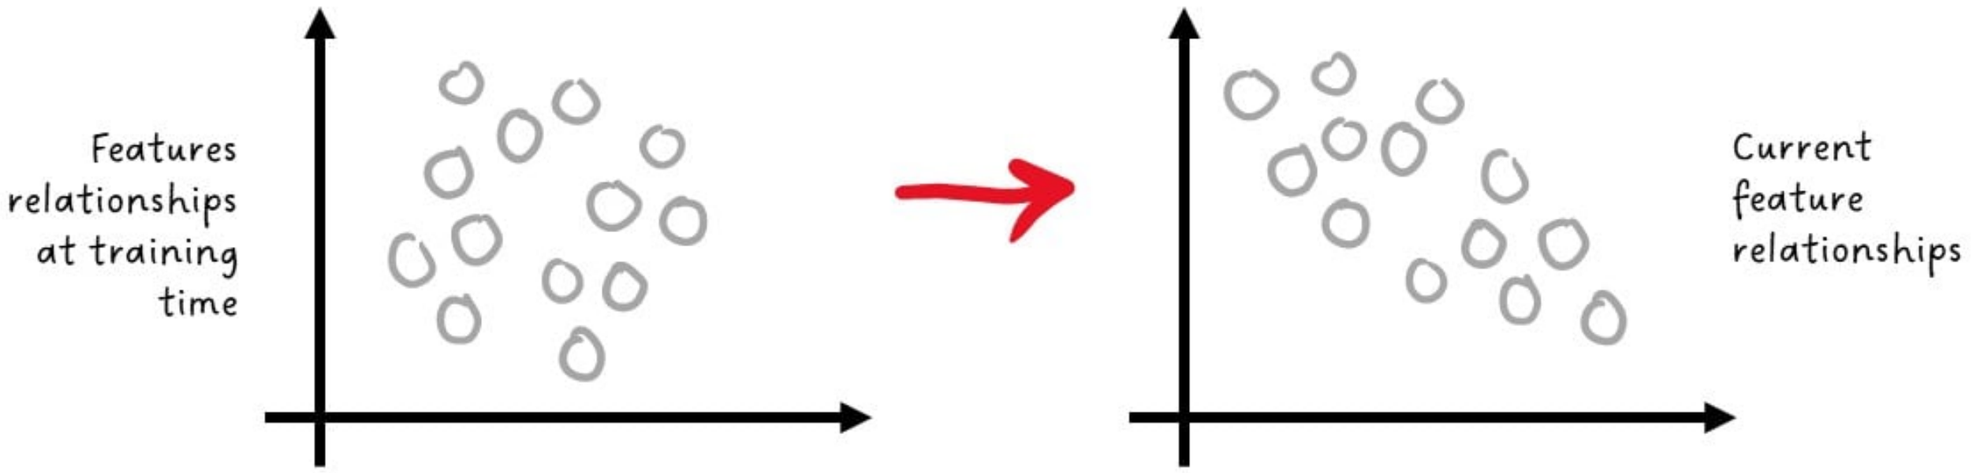

Indirect approach:

Monitor what you have: input features!

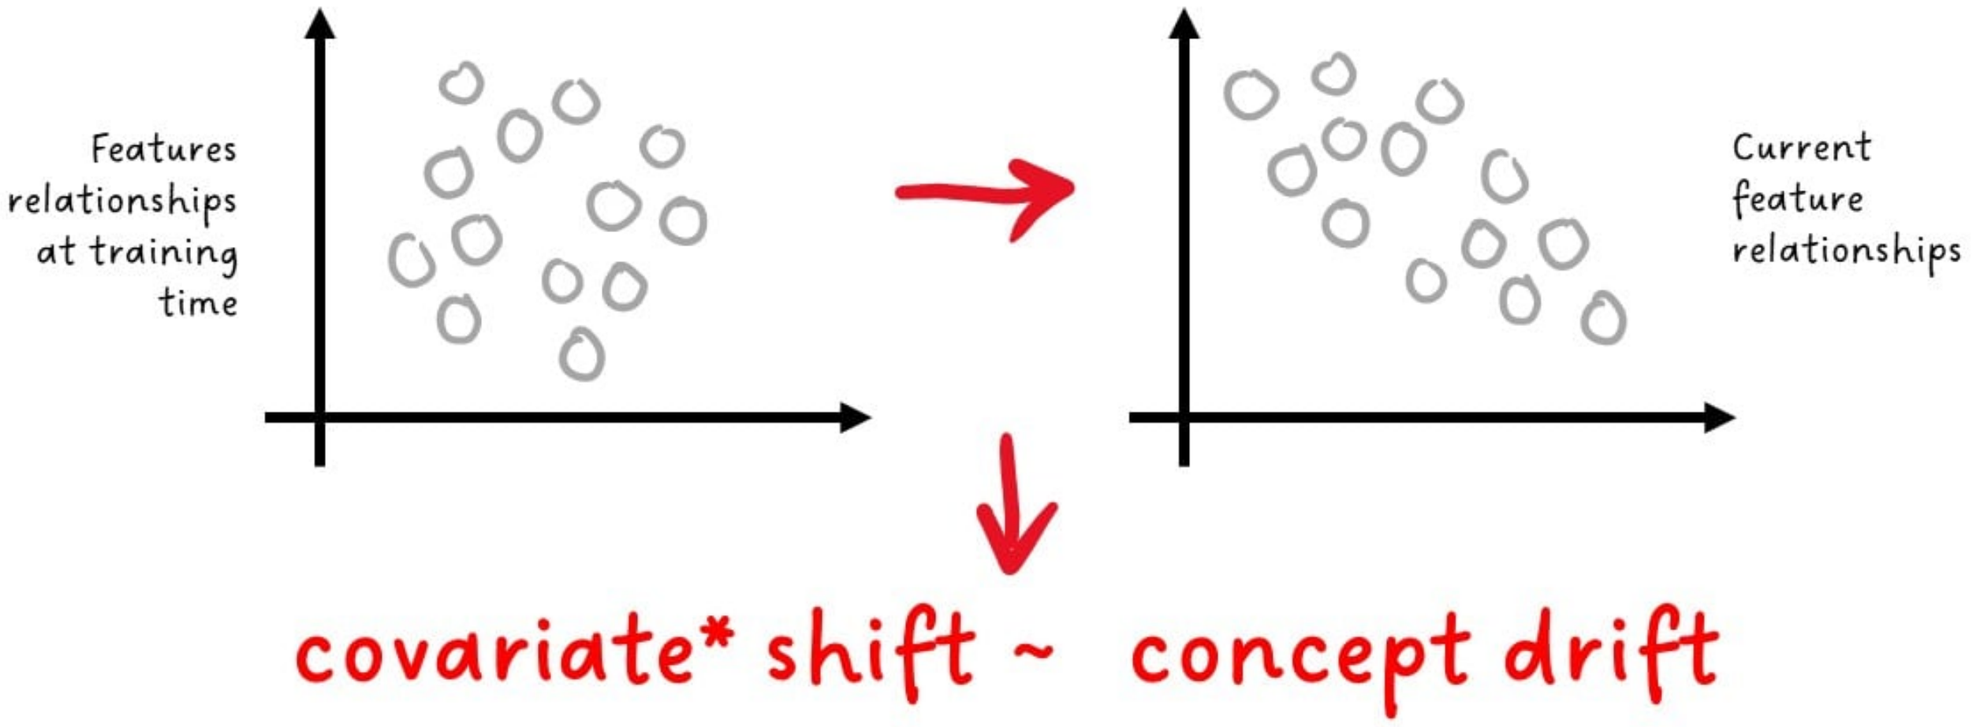

* features == "covariates" in statistical jargon

Limitations of input monitoring

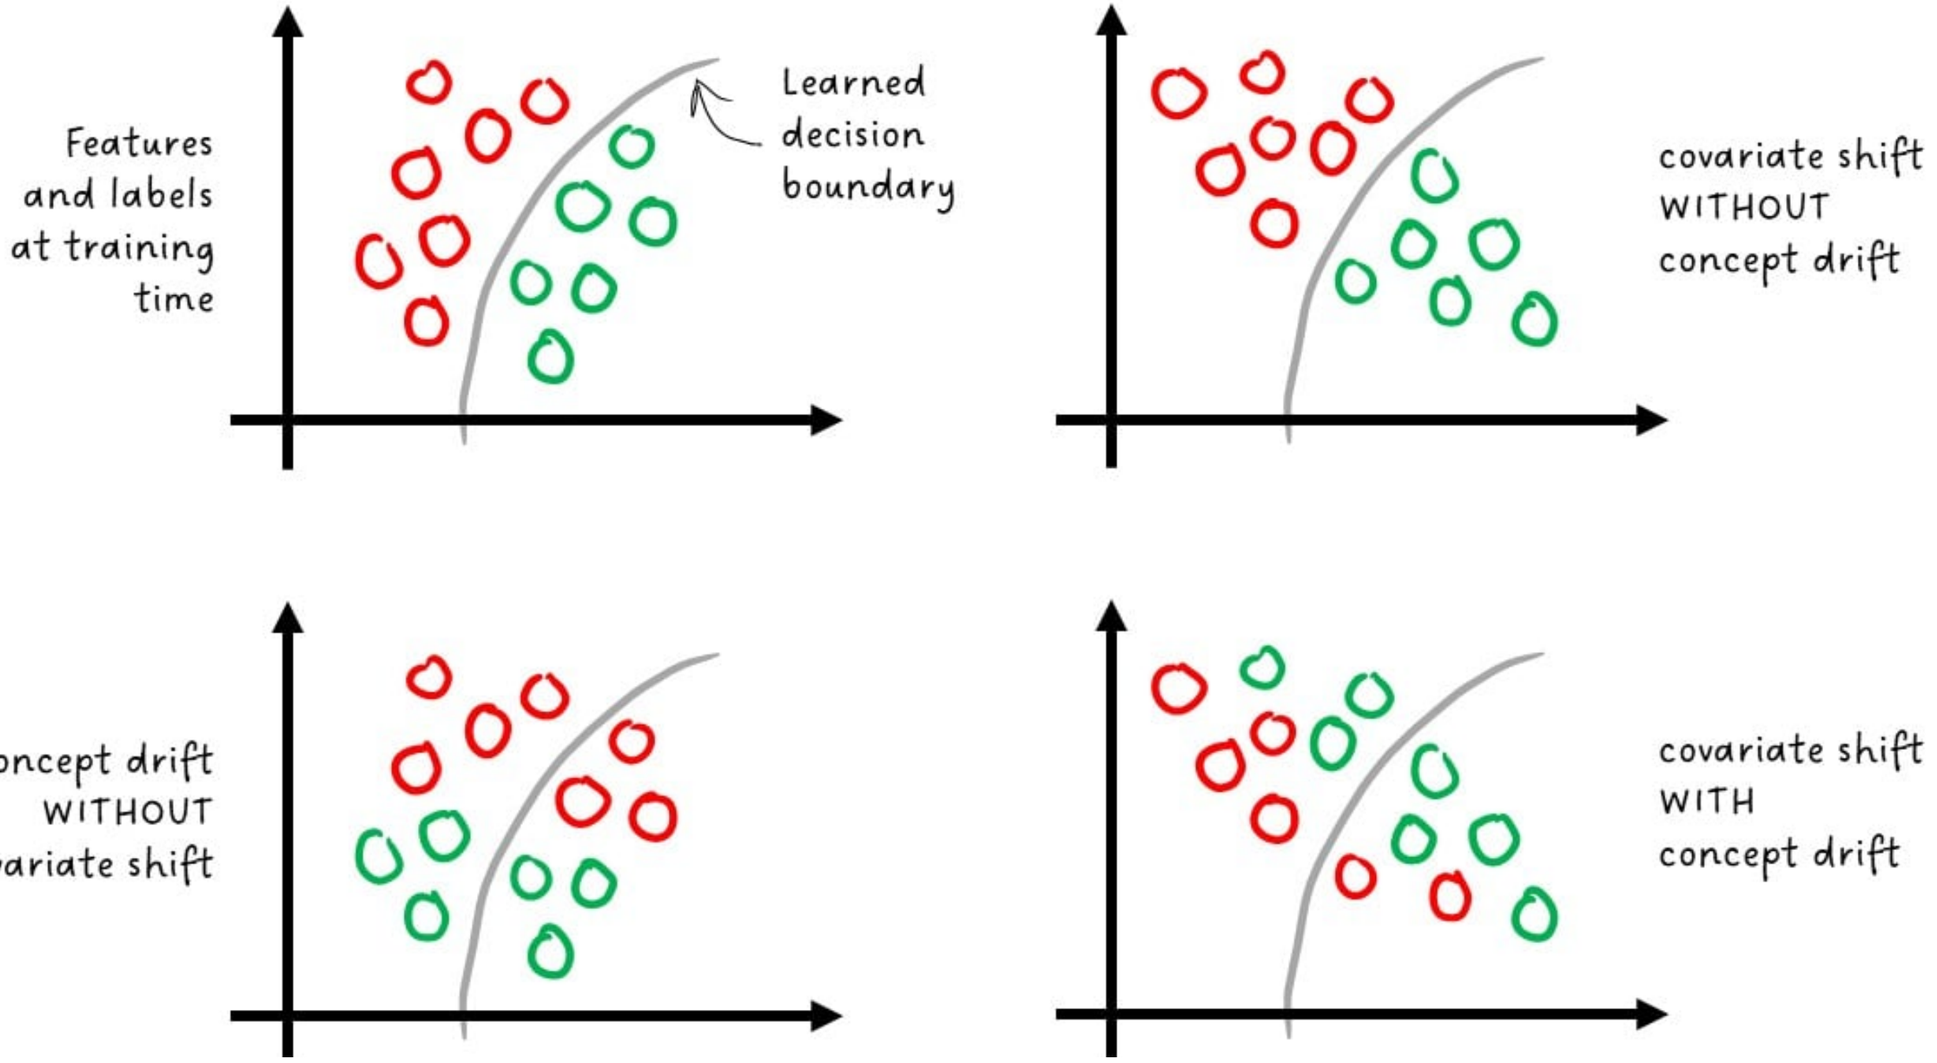

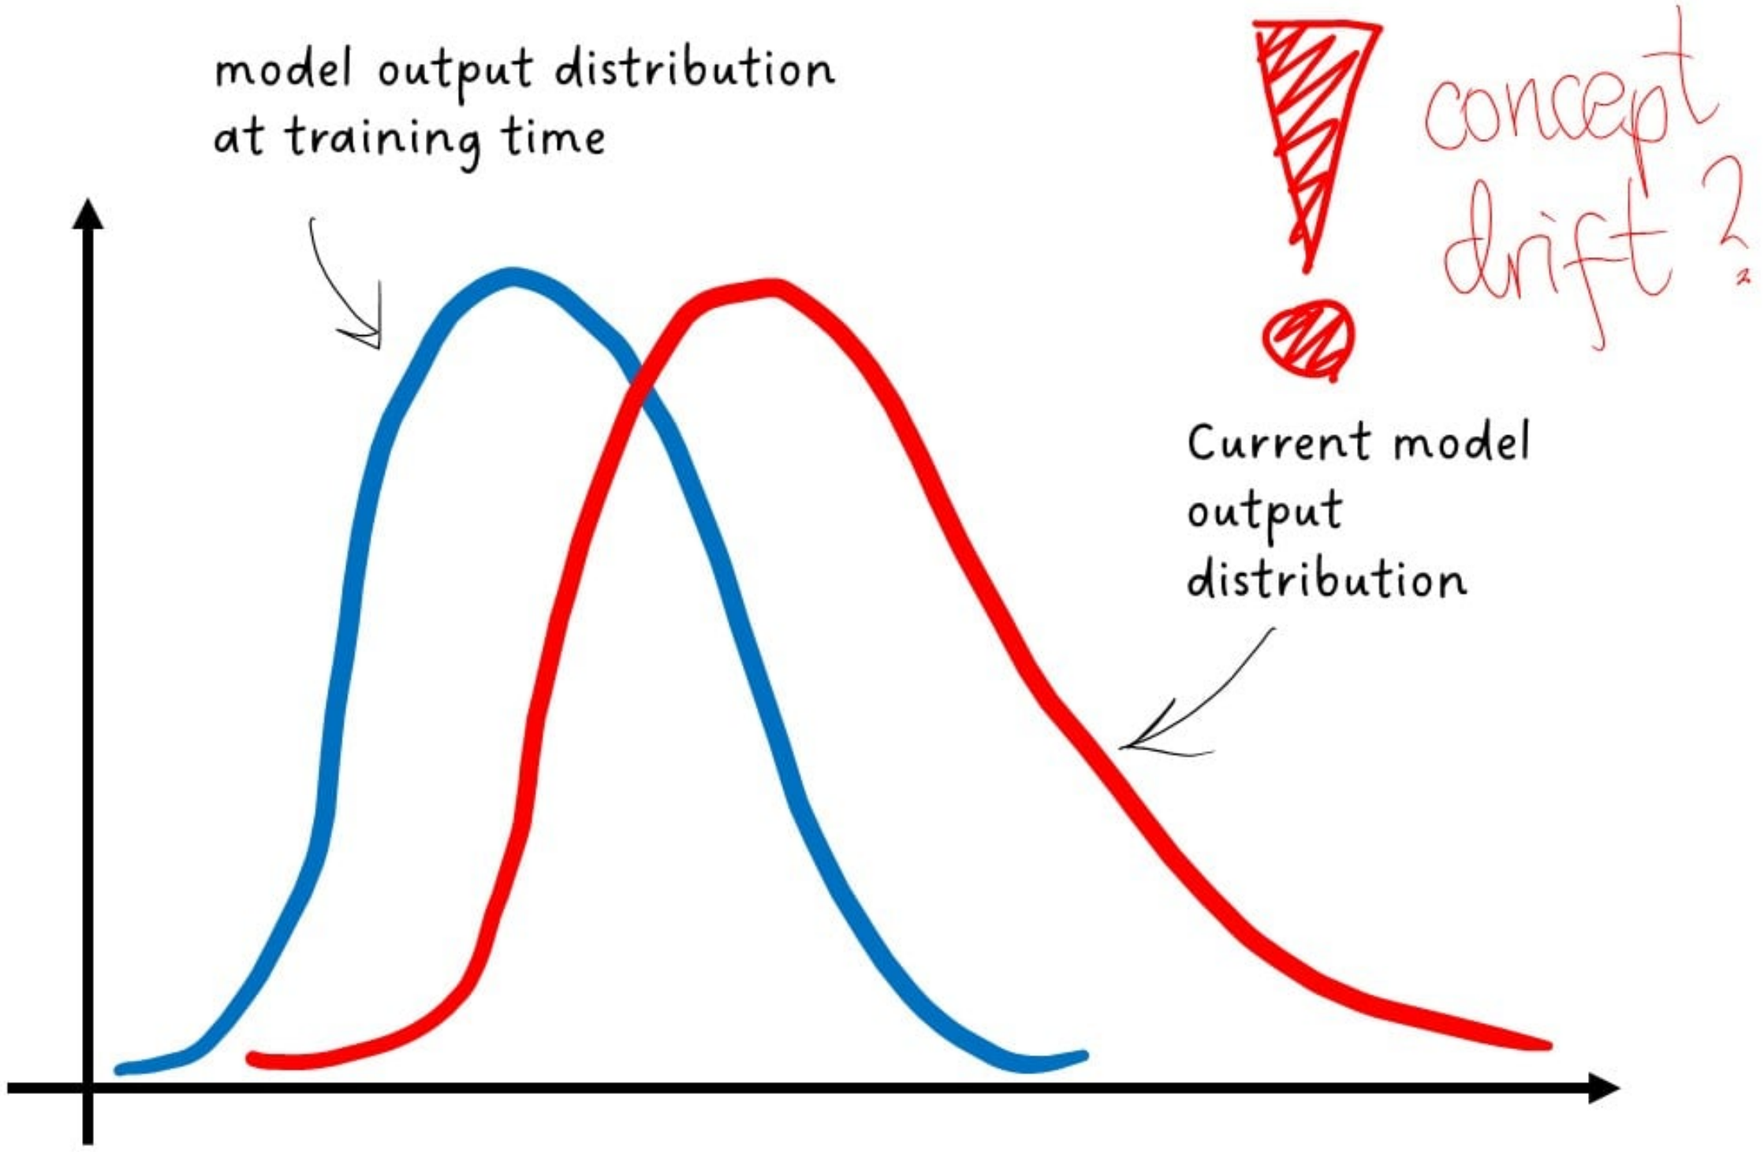
model output distribution
at training time
concept
drift ?
Current model
output
distribution

Let's practice!

MLOPS DEPLOYMENT AND LIFE CYCLING

Monitoring and alerting

MLOPS DEPLOYMENT AND LIFE CYCLING

Nemanja Radojkovic
Senior Machine Learning Engineer

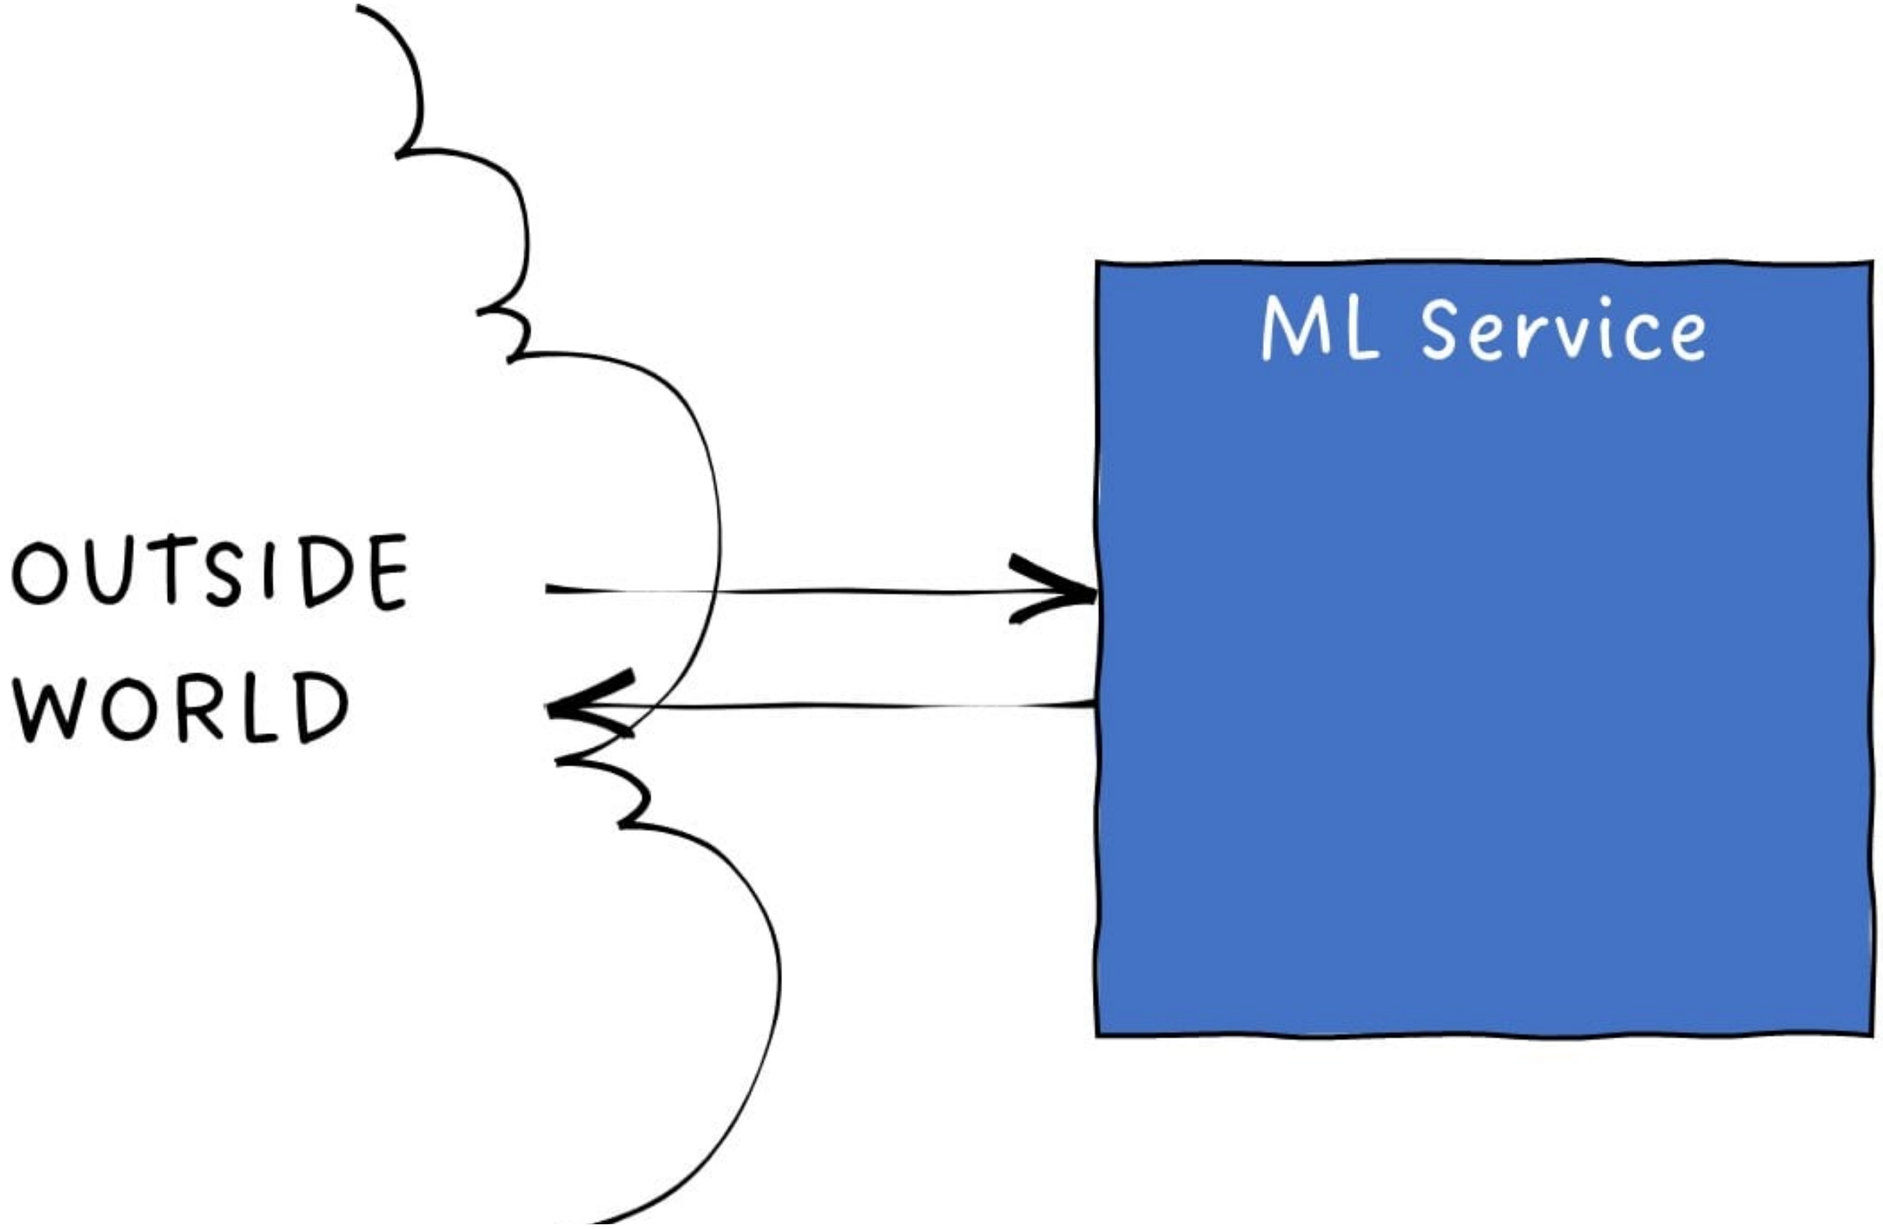
ML Service
OUTSIDE
WORLD

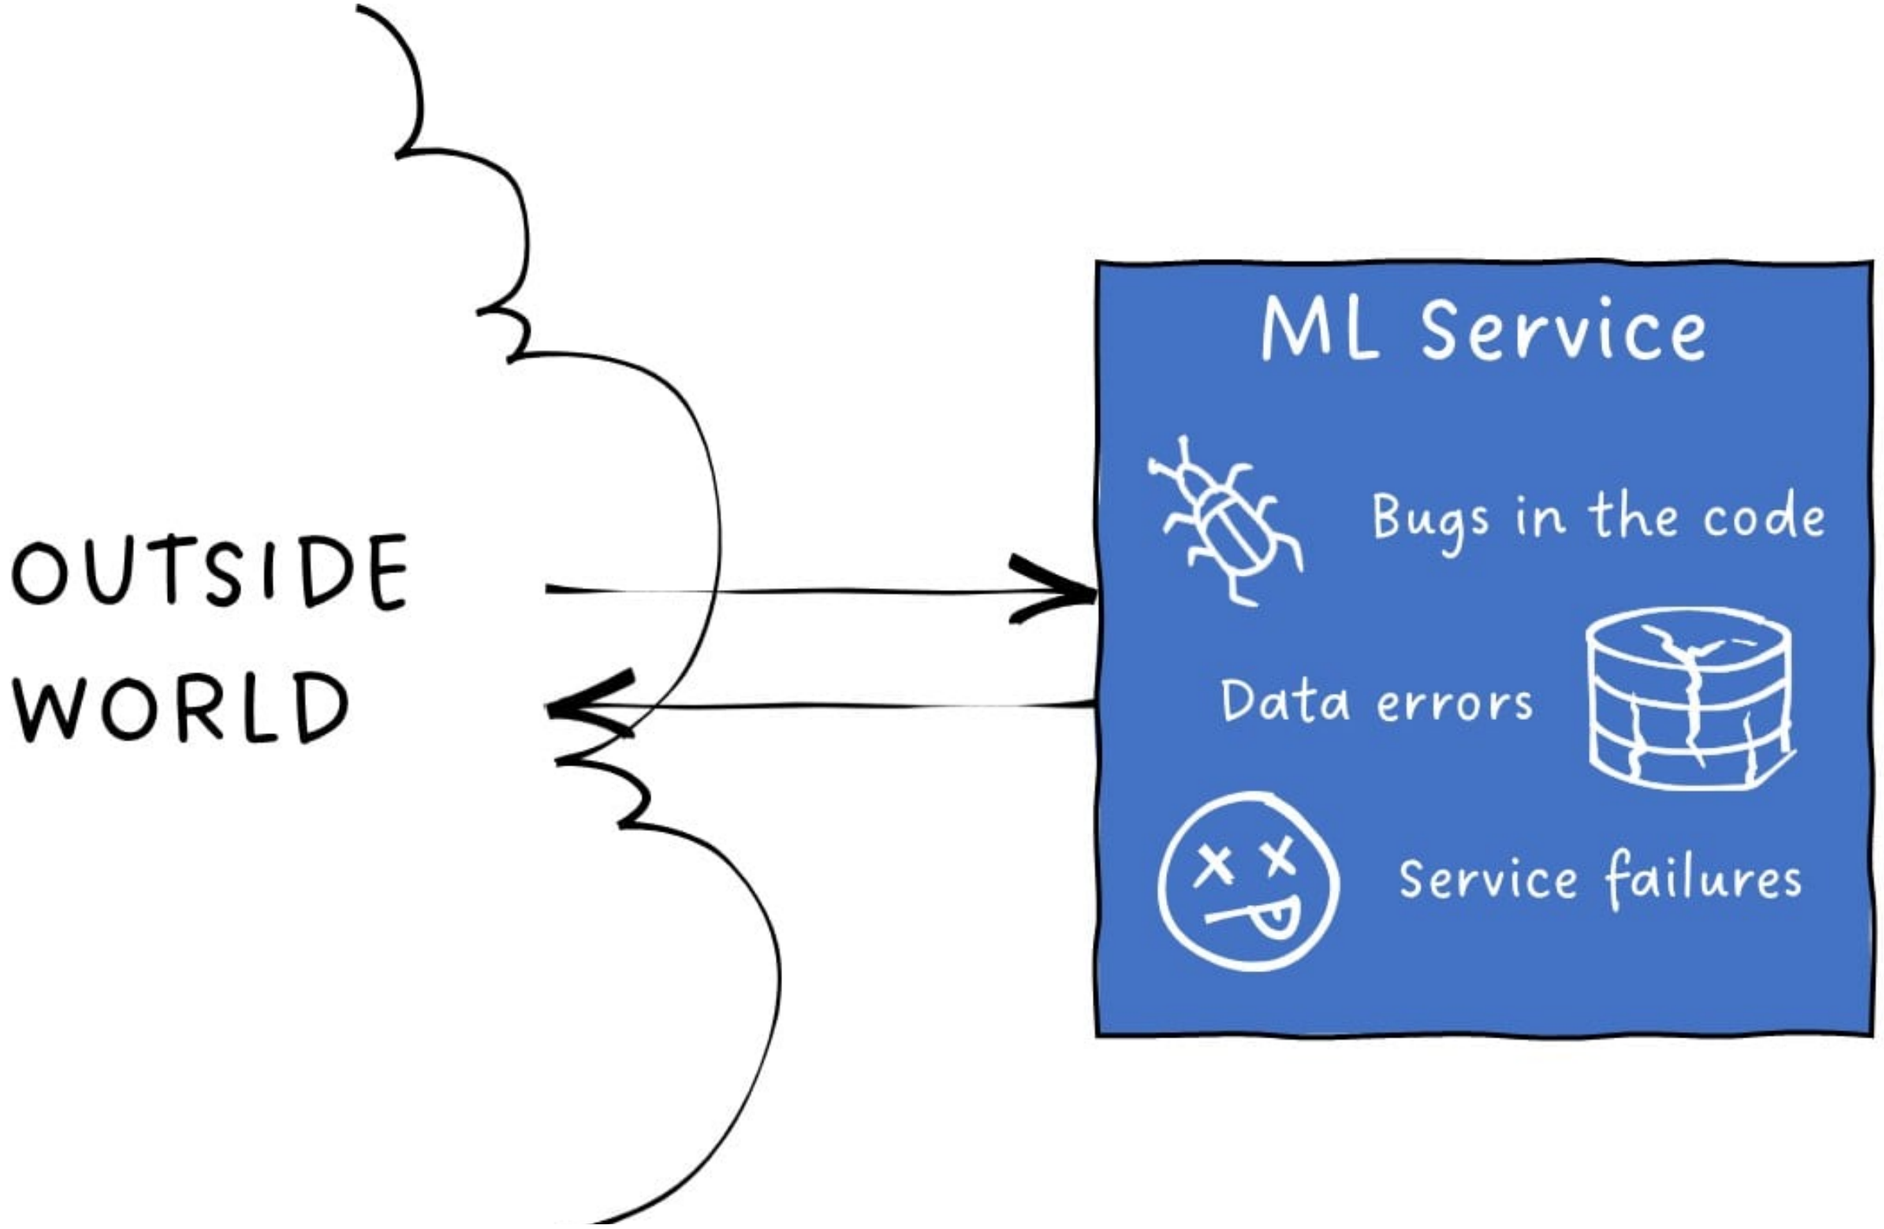
OUTSIDE
WORLD
ML Service
Bugs in the code
Data errors
Service failures

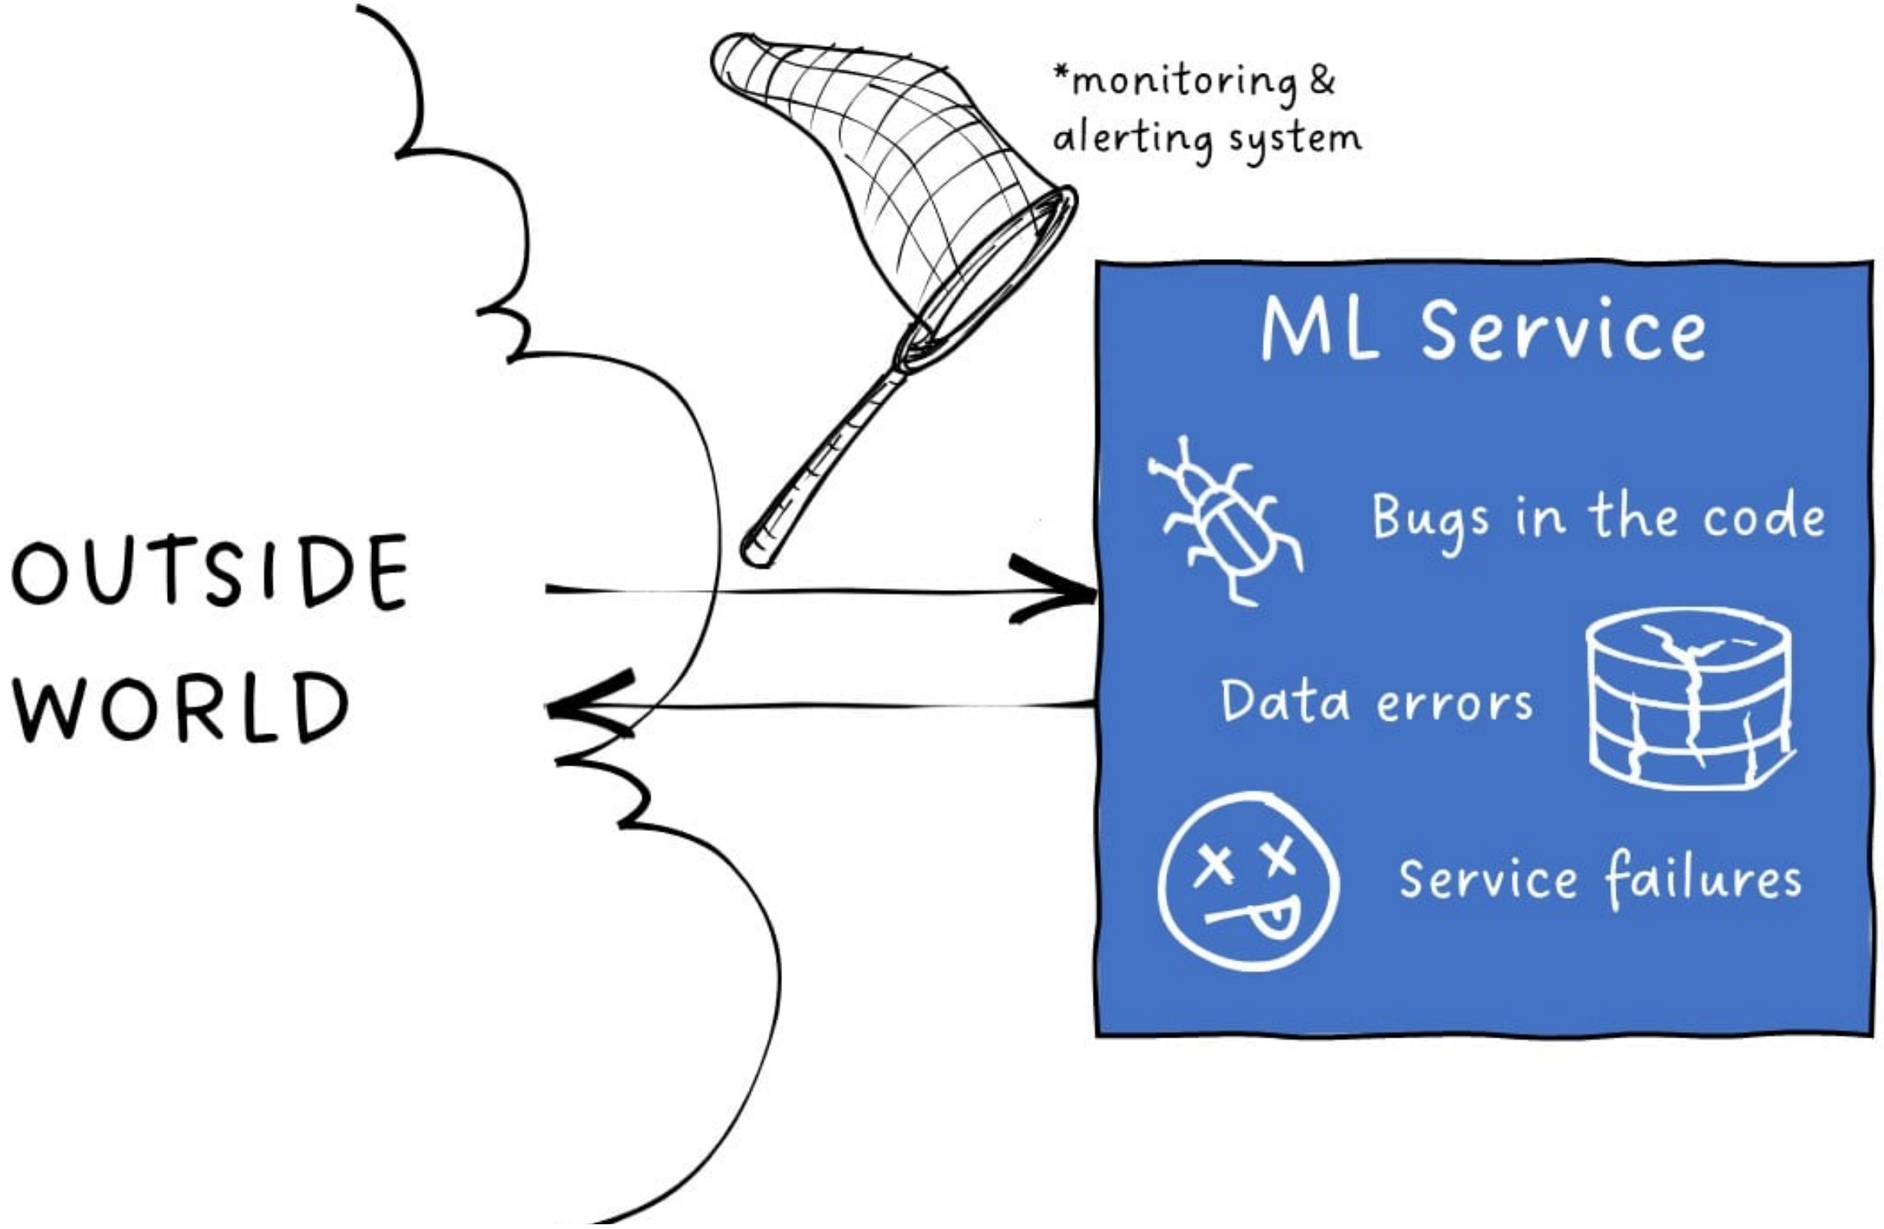
*monitoring &
alerting system
OUTSIDE
WORLD
ML Service
Bugs in the code
Data errors
Service failures

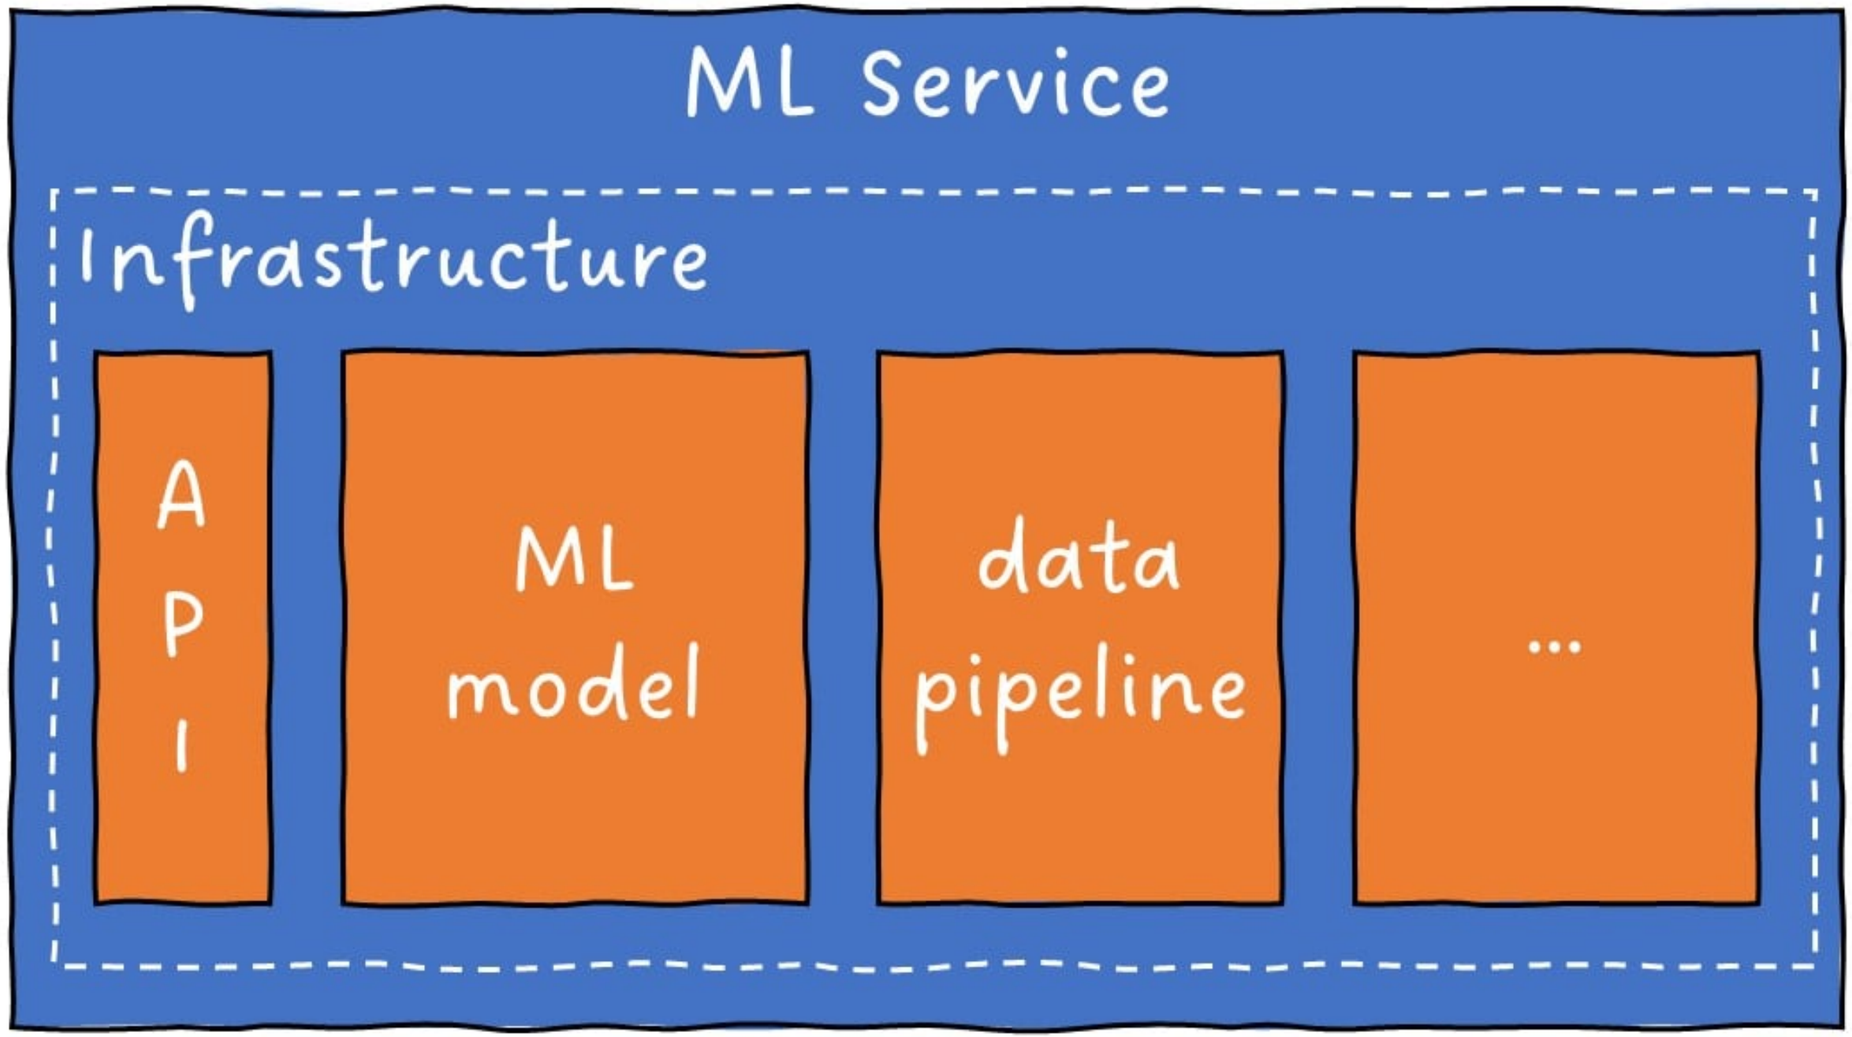
ML Service
Infrastructure
API
ML model
data pipeline
...

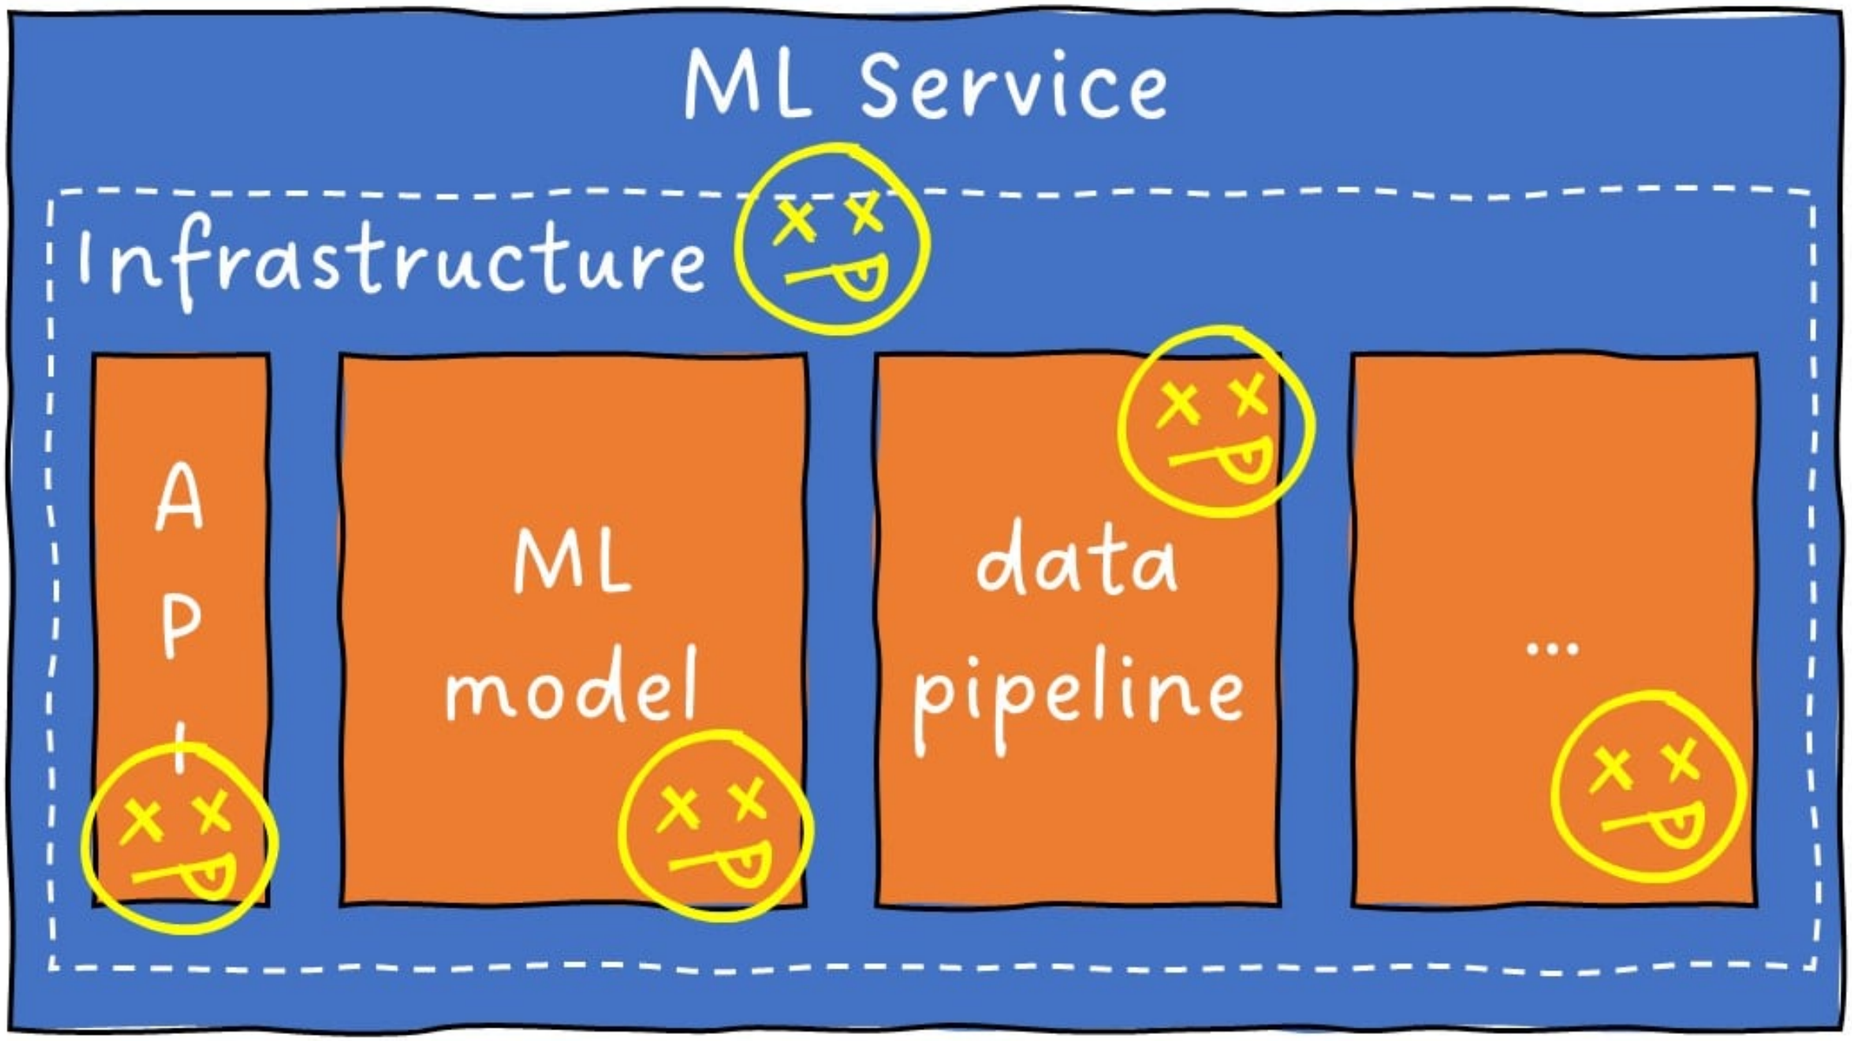
ML Service
Infrastructure
API
ML model
data pipeline
...

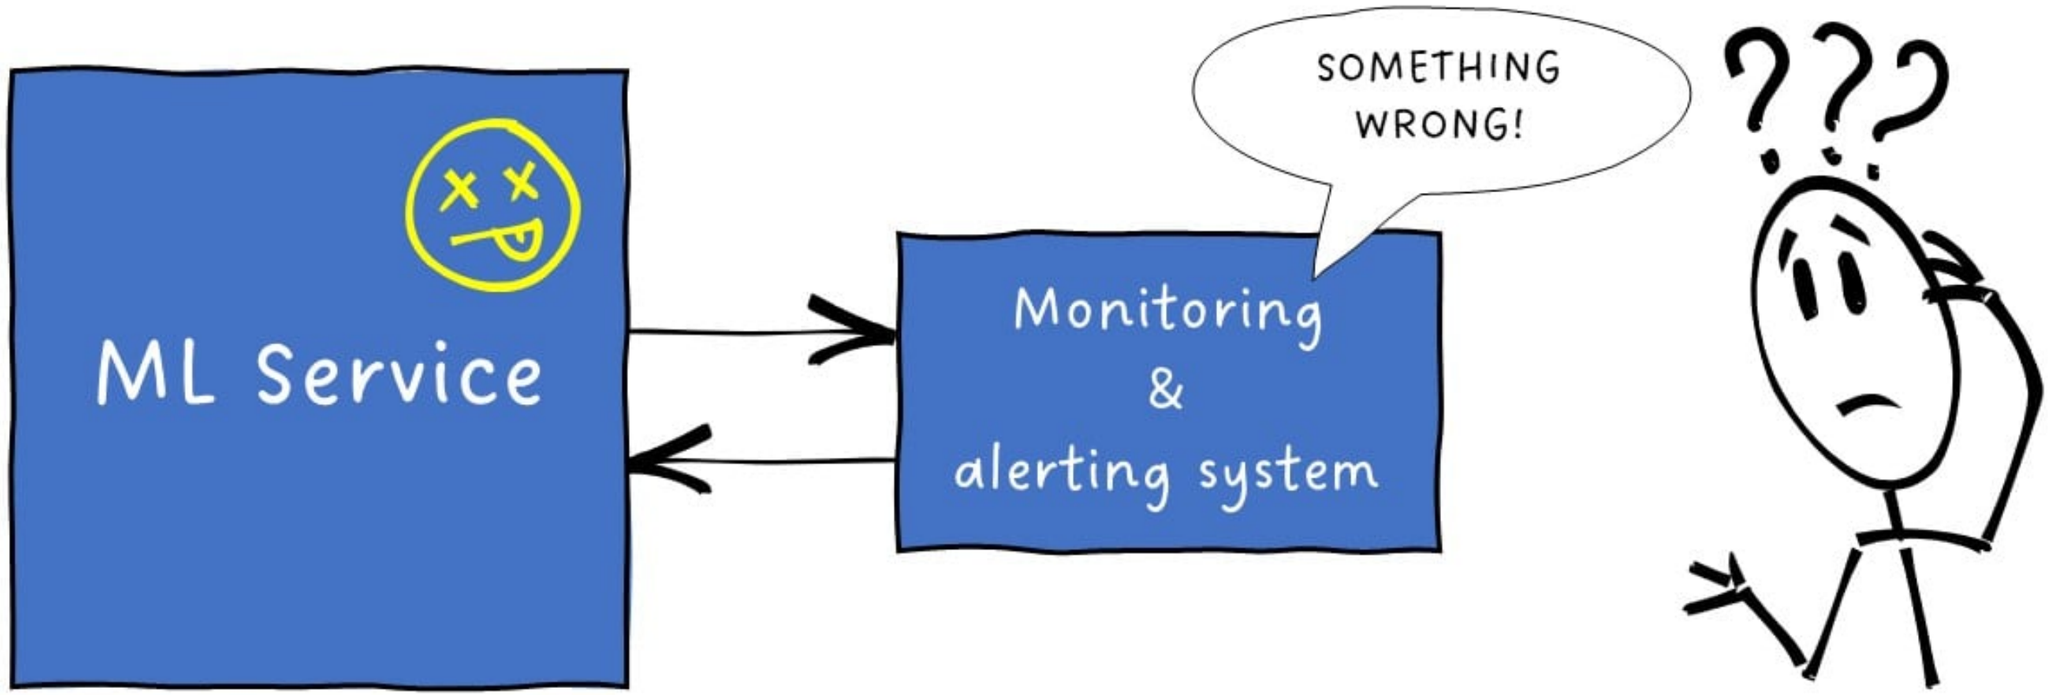
ML Service
Monitoring
&
alerting system
SOMETHING
WRONG!
???

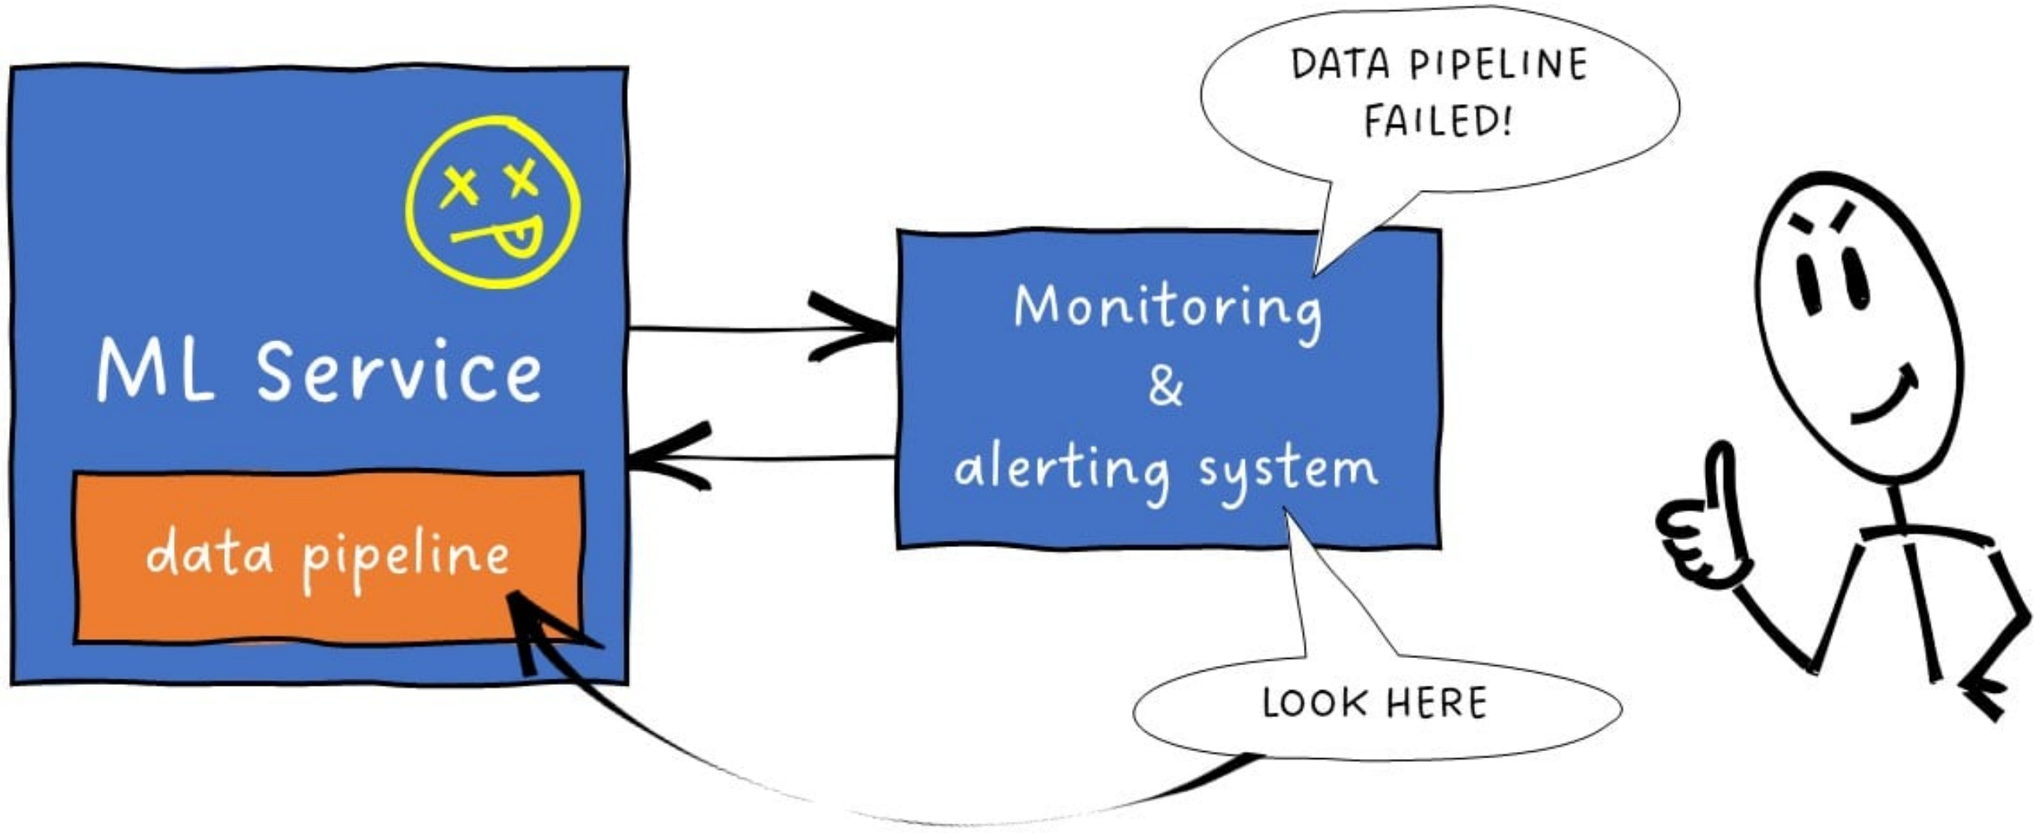
ML Service
data pipeline
Monitoring
&
alerting system
DATA PIPELINE
FAILED!
LOOK HERE

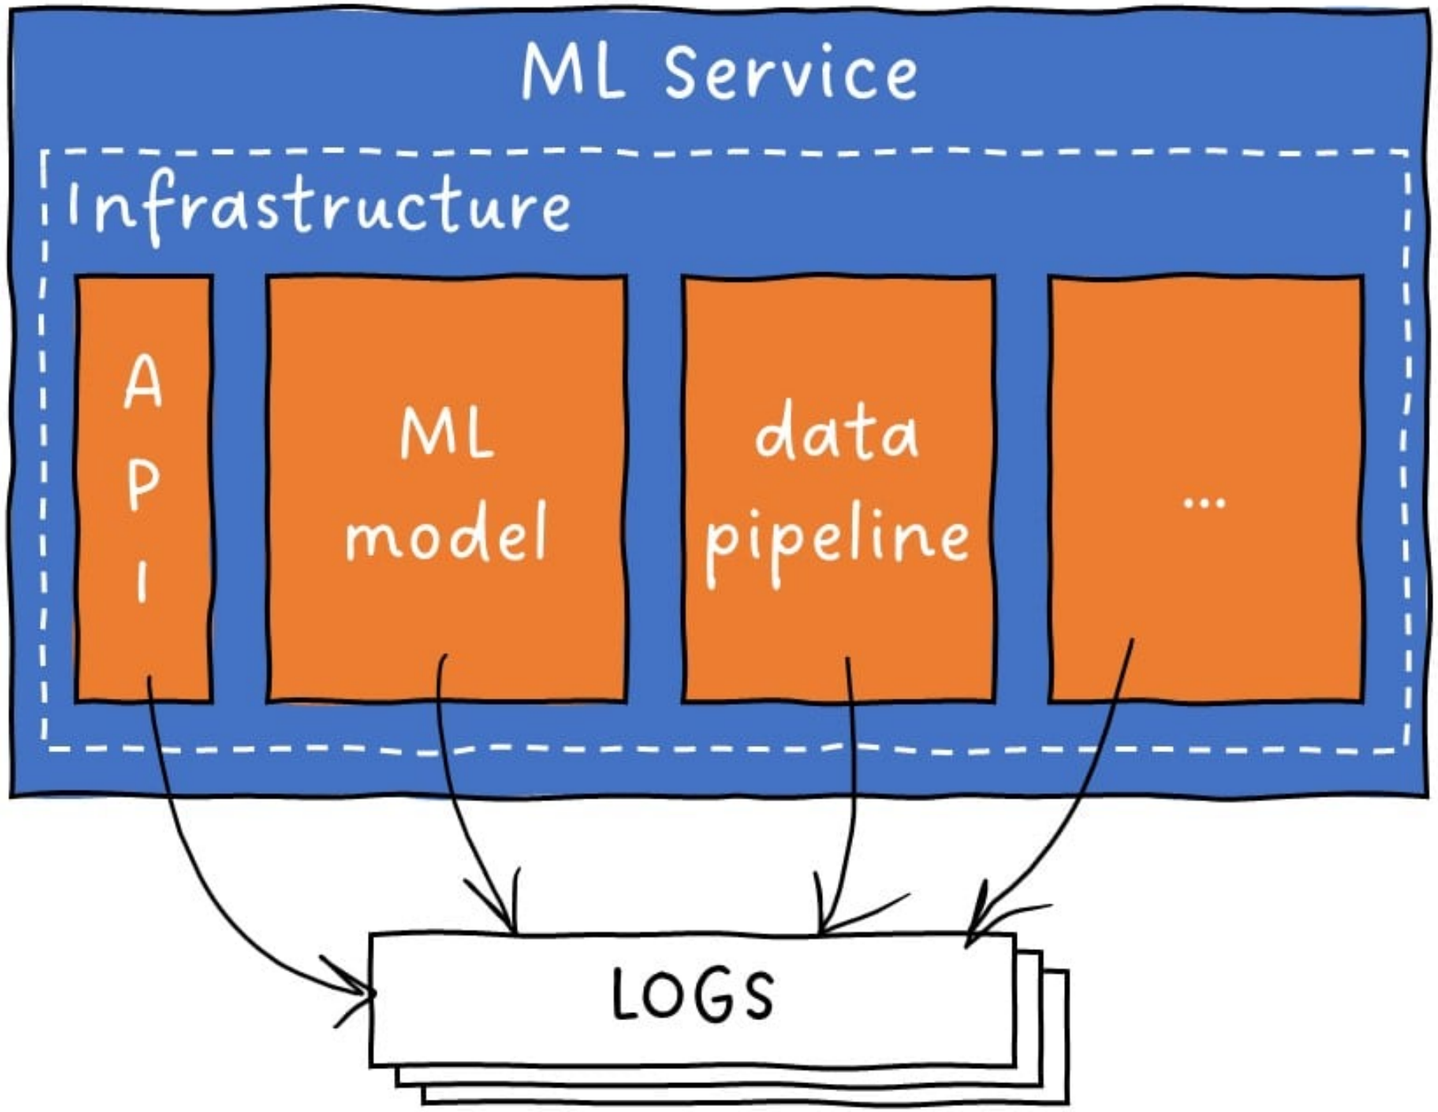
ML Service
Infrastructure
API
ML model
data pipeline
...
LOGS

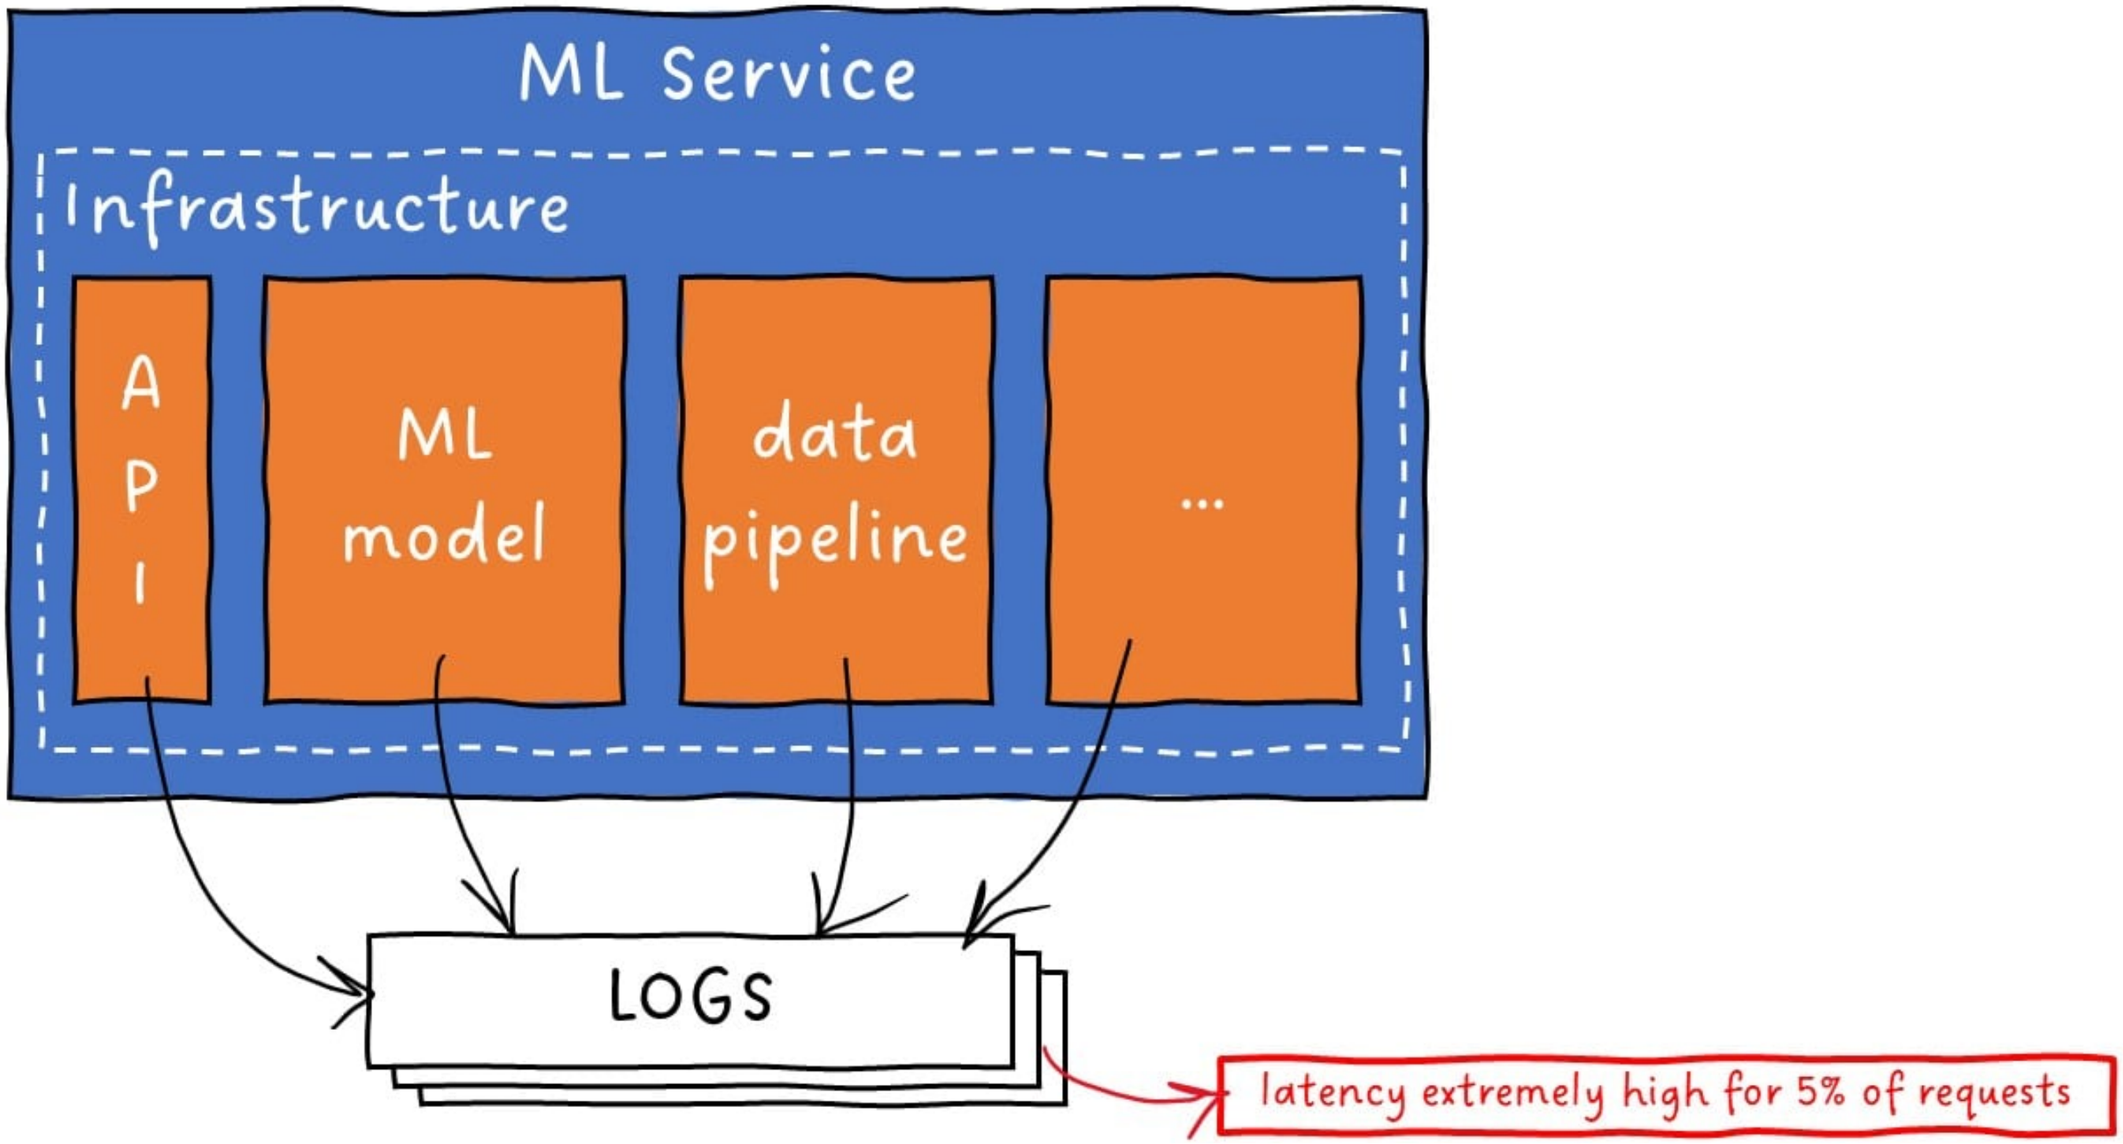
ML Service
Infrastructure
API
ML model
data pipeline
...
LOGS
latency extremely high for 5% of requests

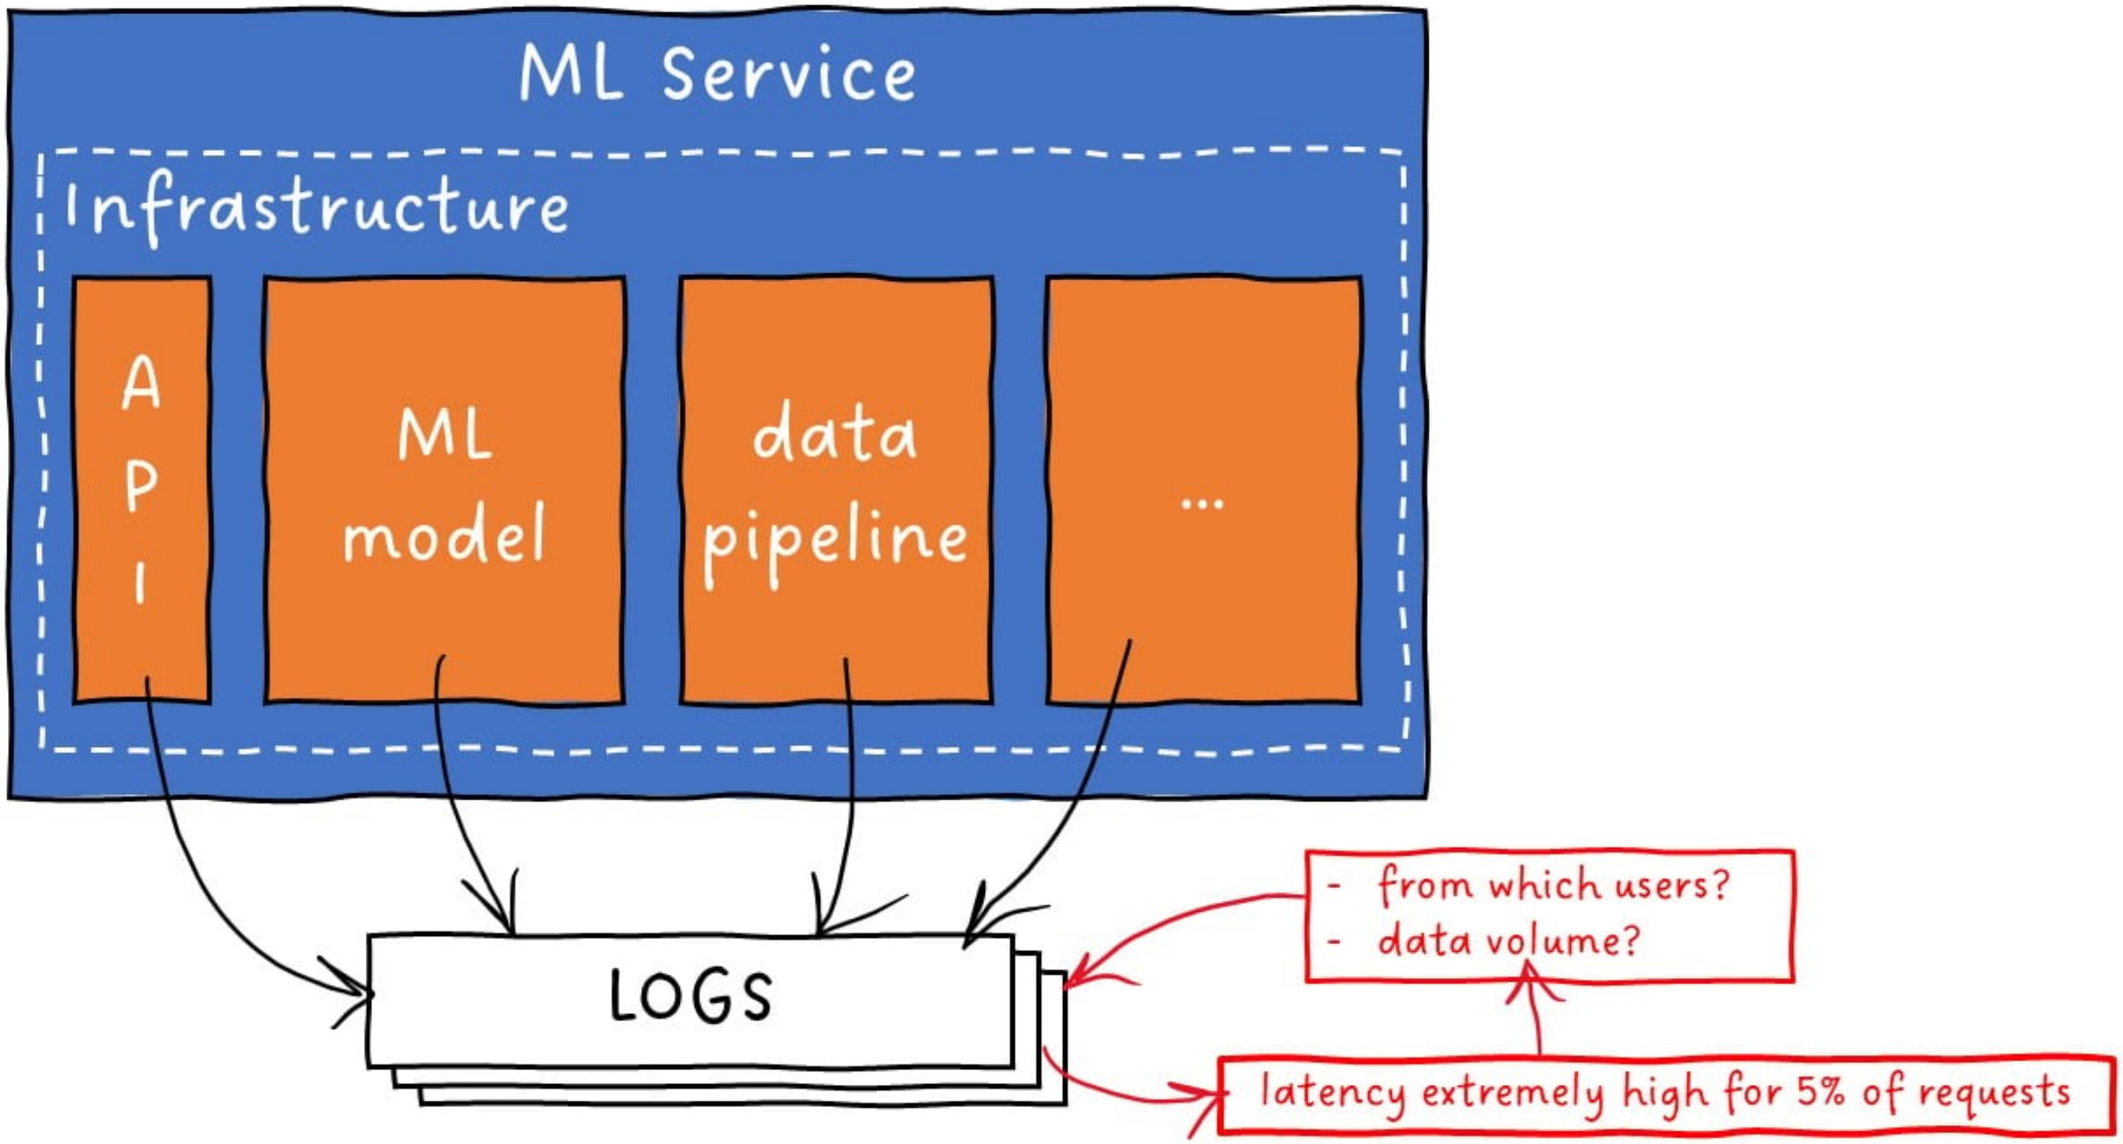
ML Service
Infrastructure
API
ML model
data pipeline
...
LOGS
- from which users?
- data volume?
latency extremely high for 5% of requests

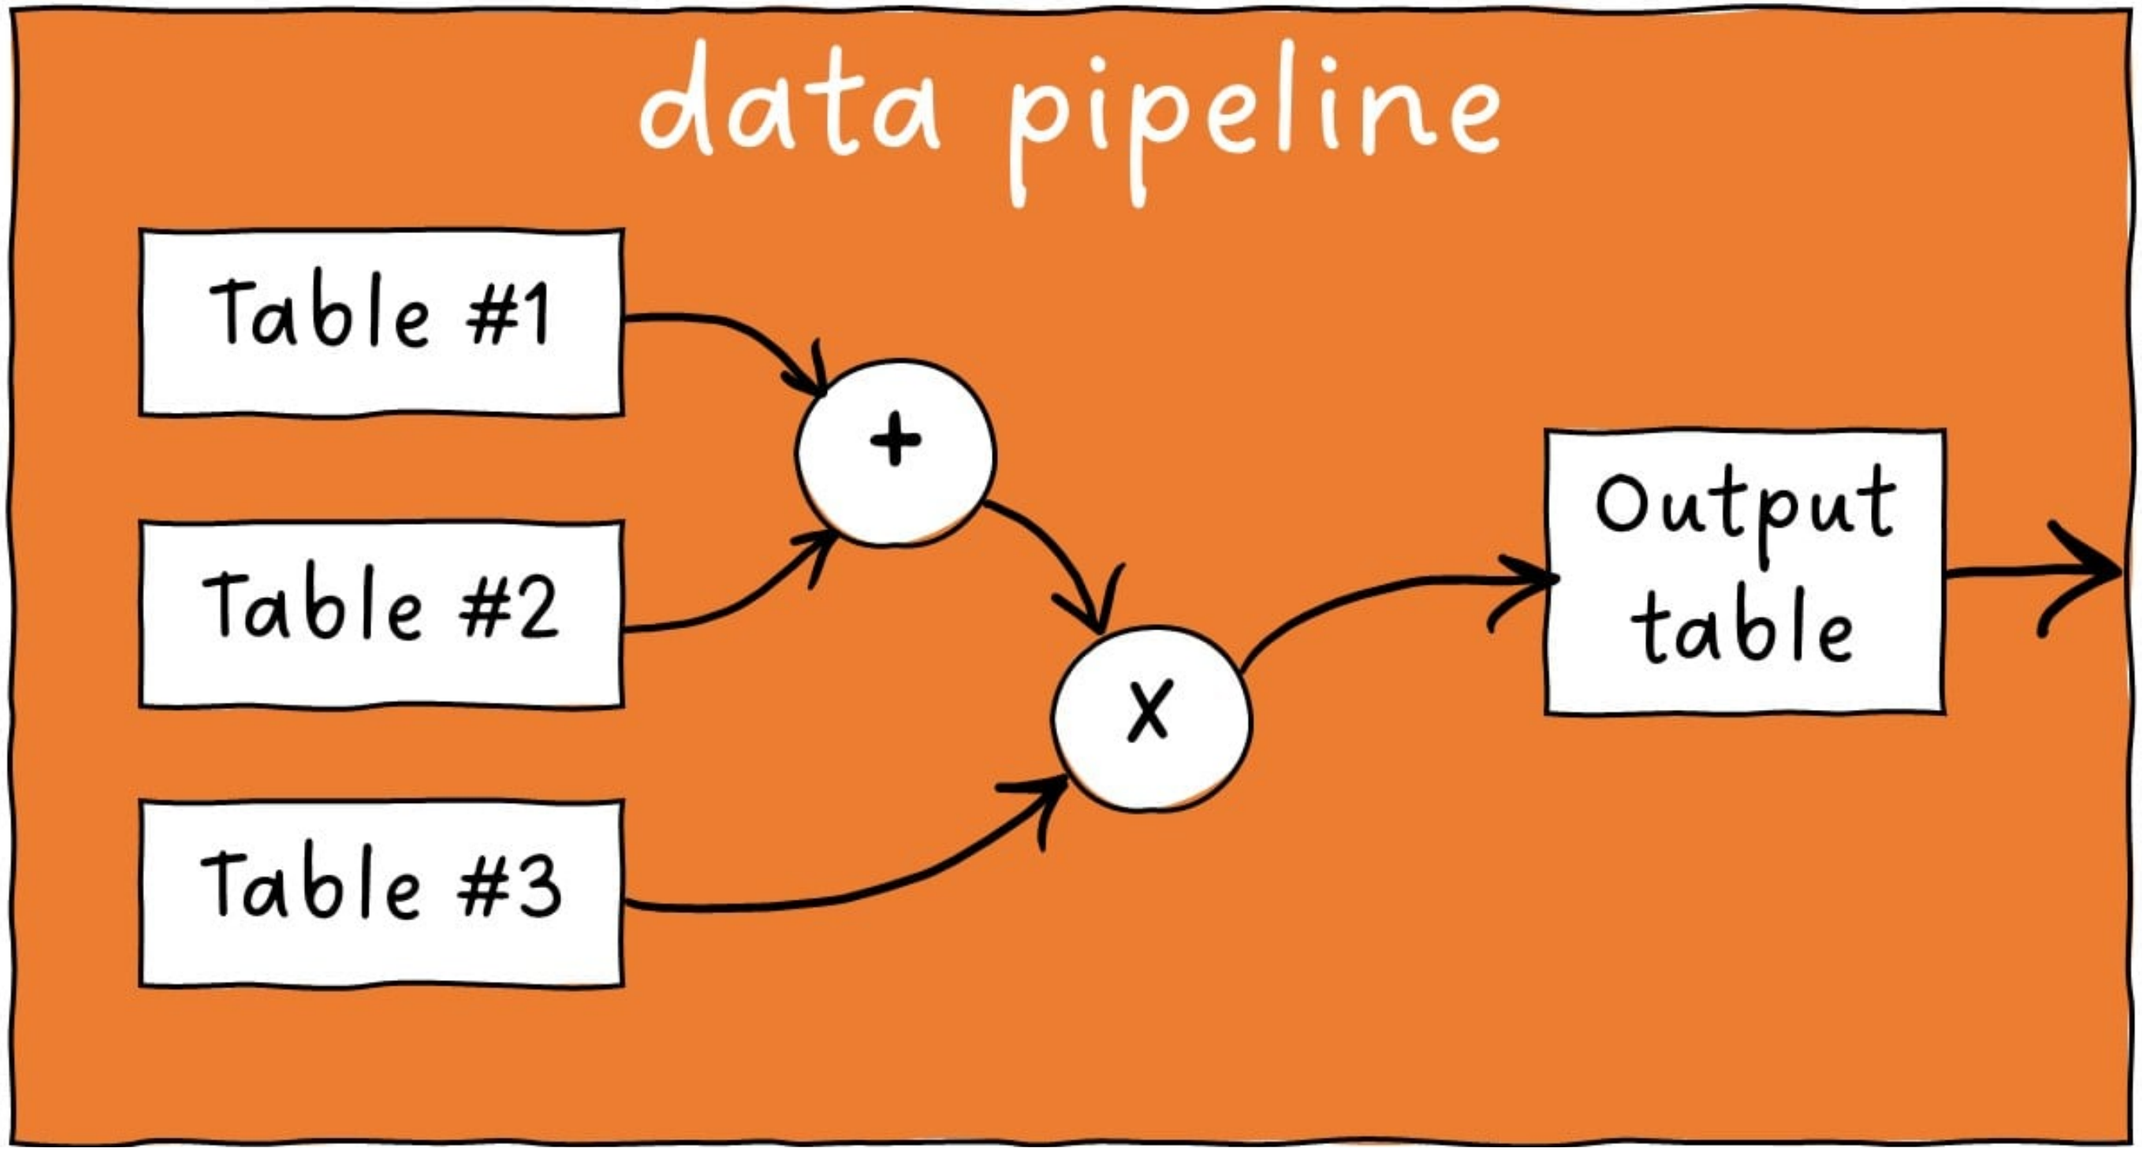
data pipeline
Table #1
Table #2
Table #3
+
X
Output table

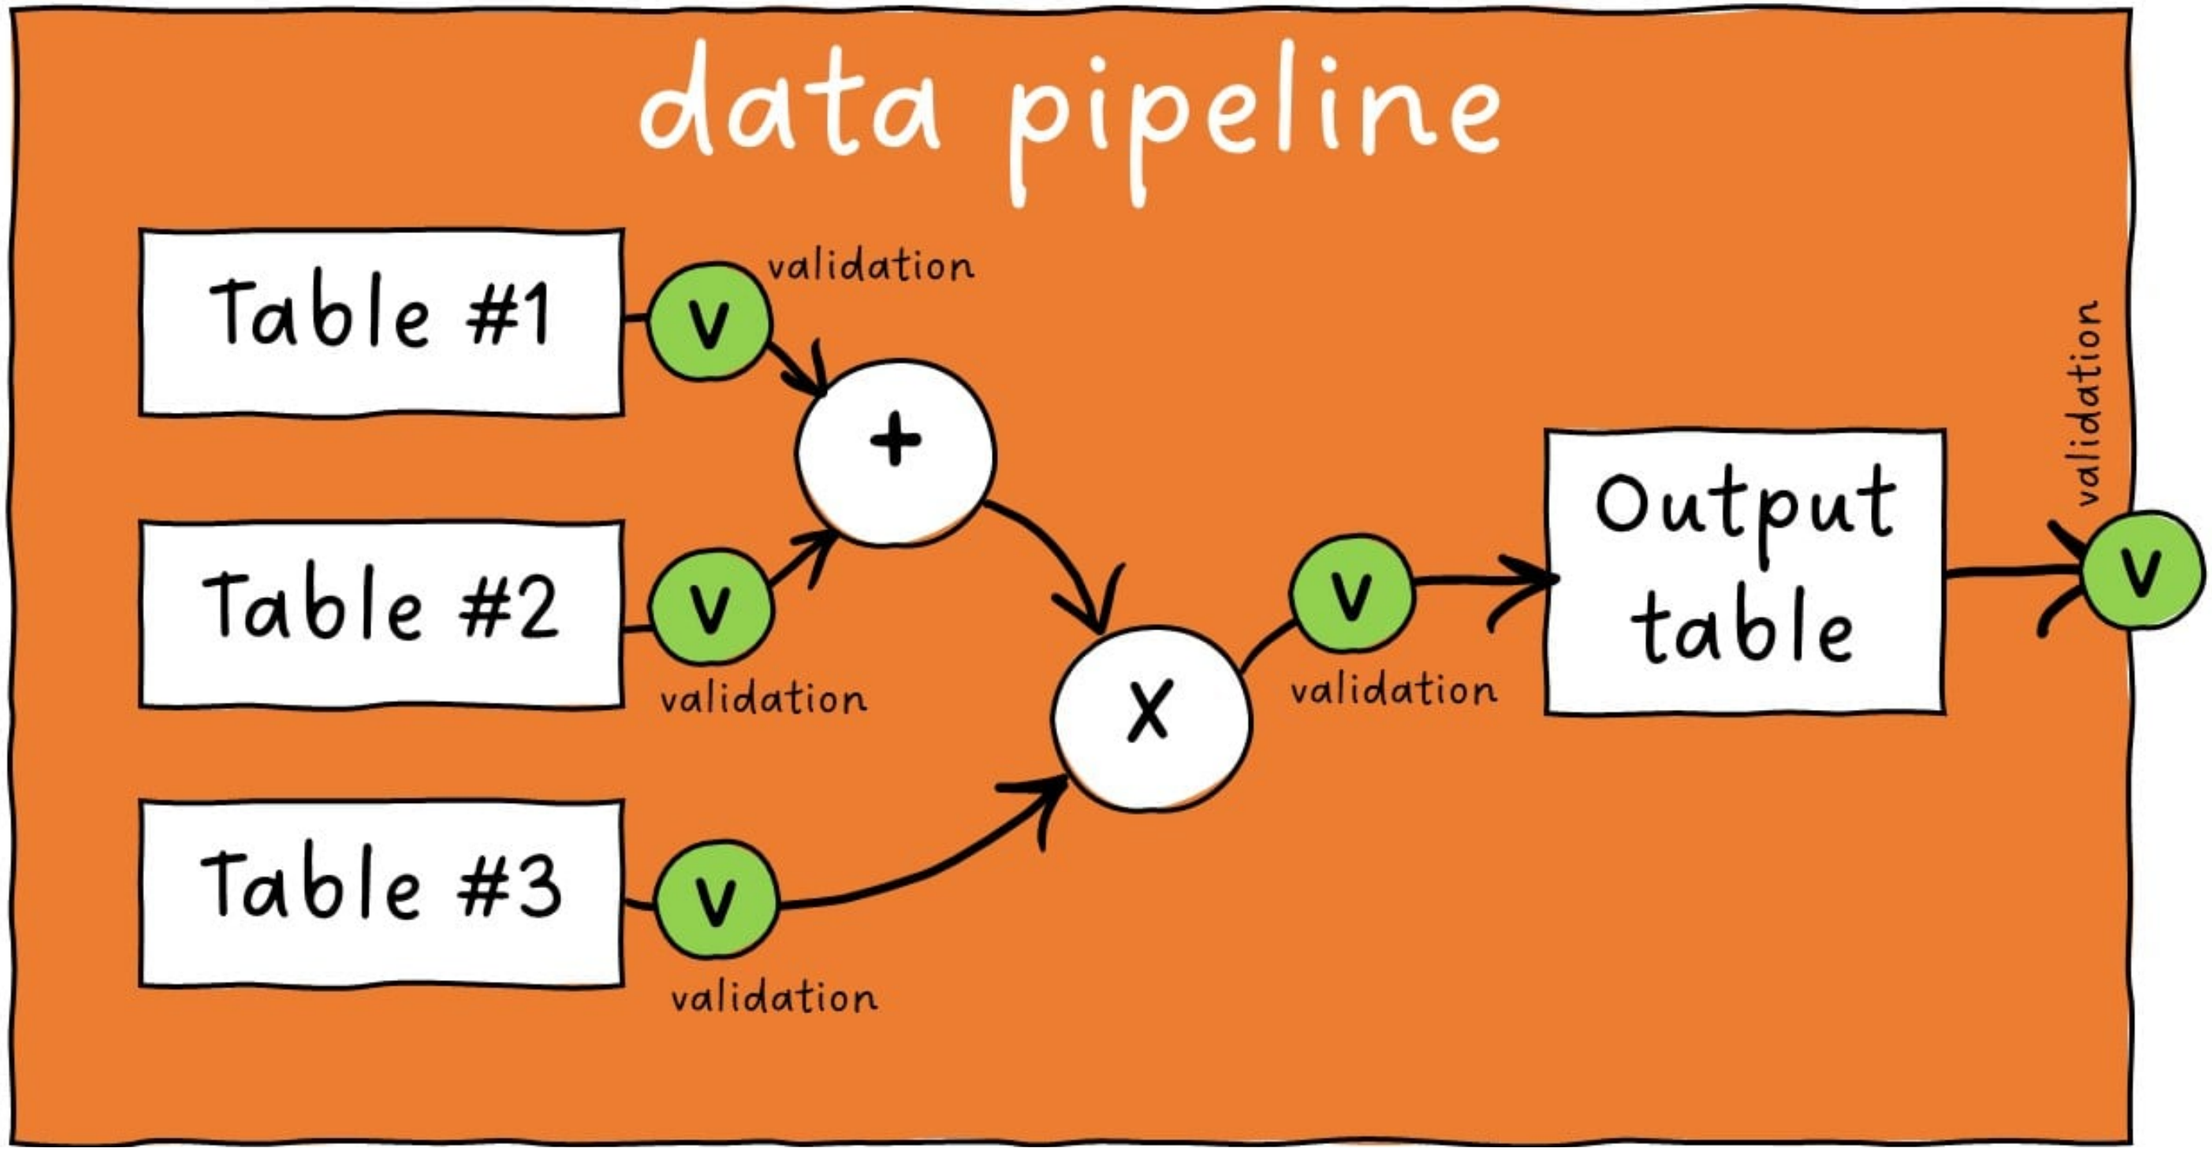
data pipeline
Table #1
Table #2
Table #3
validation
validation
validation
+
X
validation
Output table
validation

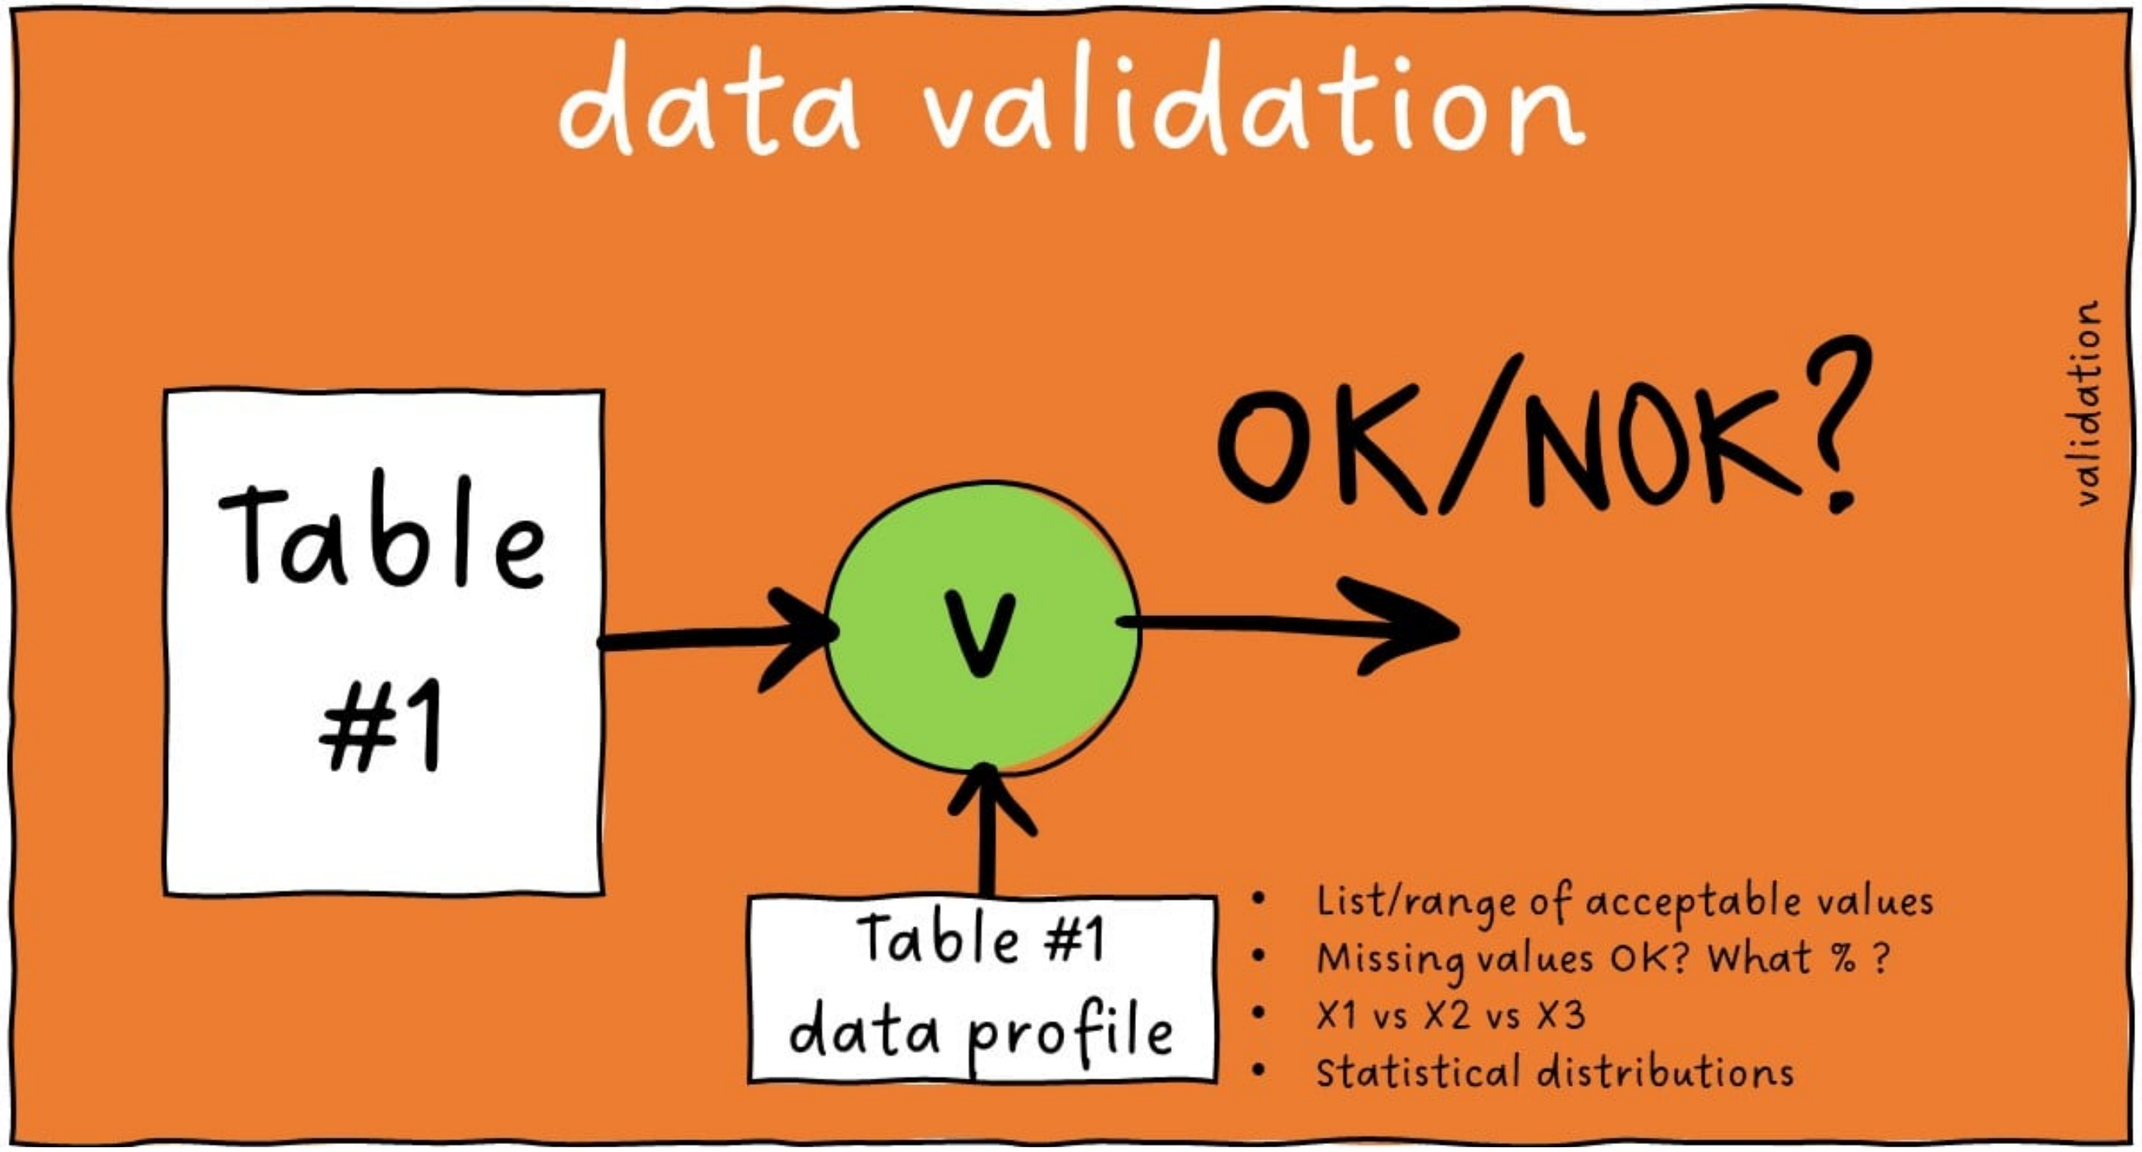
data validation
Table #1
V
OK/NOK?
validation
Table #1 data profile
List/range of acceptable values
Missing values OK? What % ?
X1 vs X2 vs X3
Statistical distributions

Statistical validation

Can be:

- too sensitive
- not informative enough

Risk

- Too many alerts
- "Alert fatigue"
- Important alerts going unnoticed

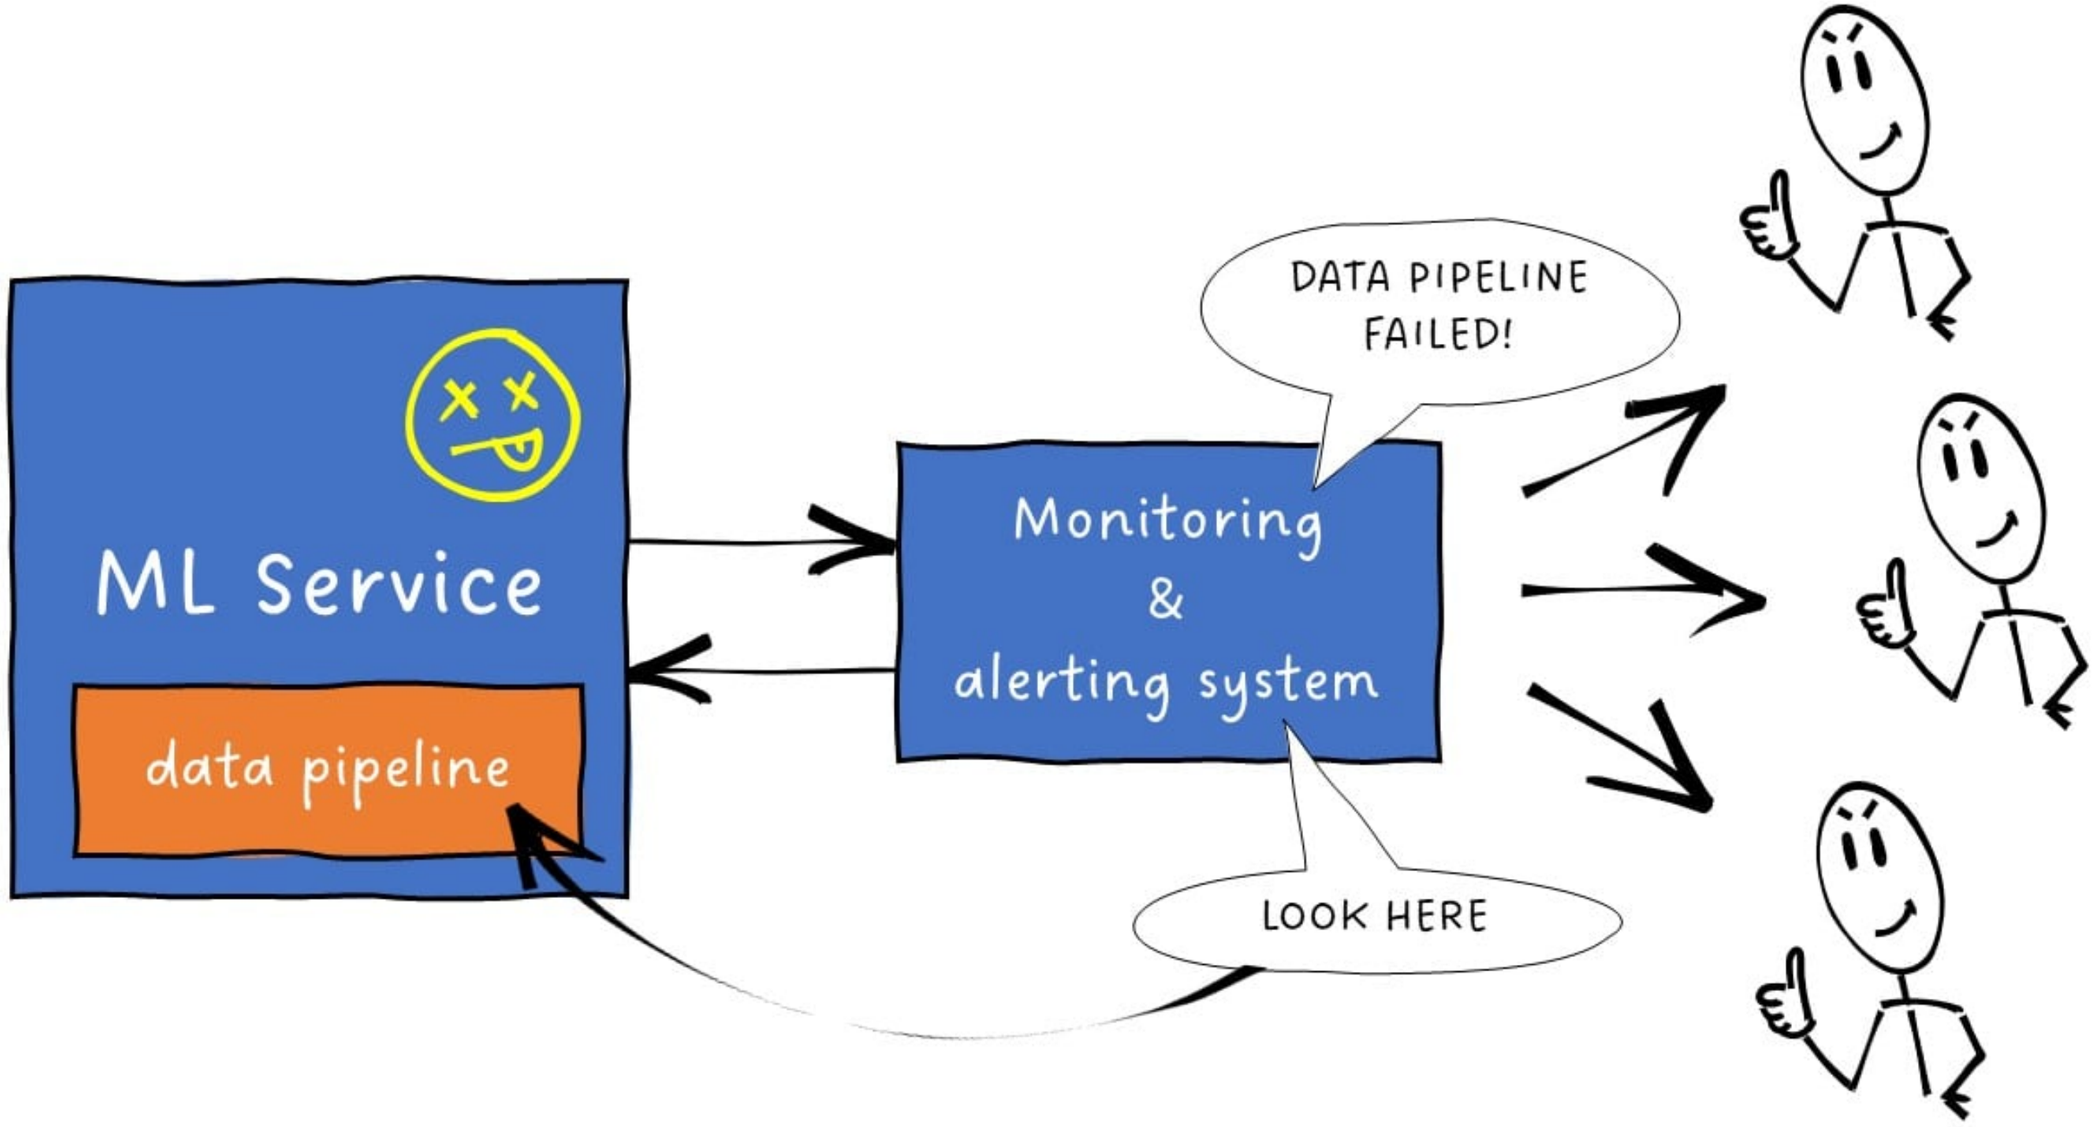
ML Service
data pipeline
Monitoring
&
alerting system
DATA PIPELINE
FAILED!
LOOK HERE

Learn from your history

After treating the incident => Record root cause and resolution steps

Example from Google[1]:

- 10 years of incidents recorded an analyzed
- > 2/3 were not ML-related!

[1] How ML Breaks: A Decade of Outages for One Large ML Pipeline, https://www.usenix.org/conference/opml20/presentation/papasian

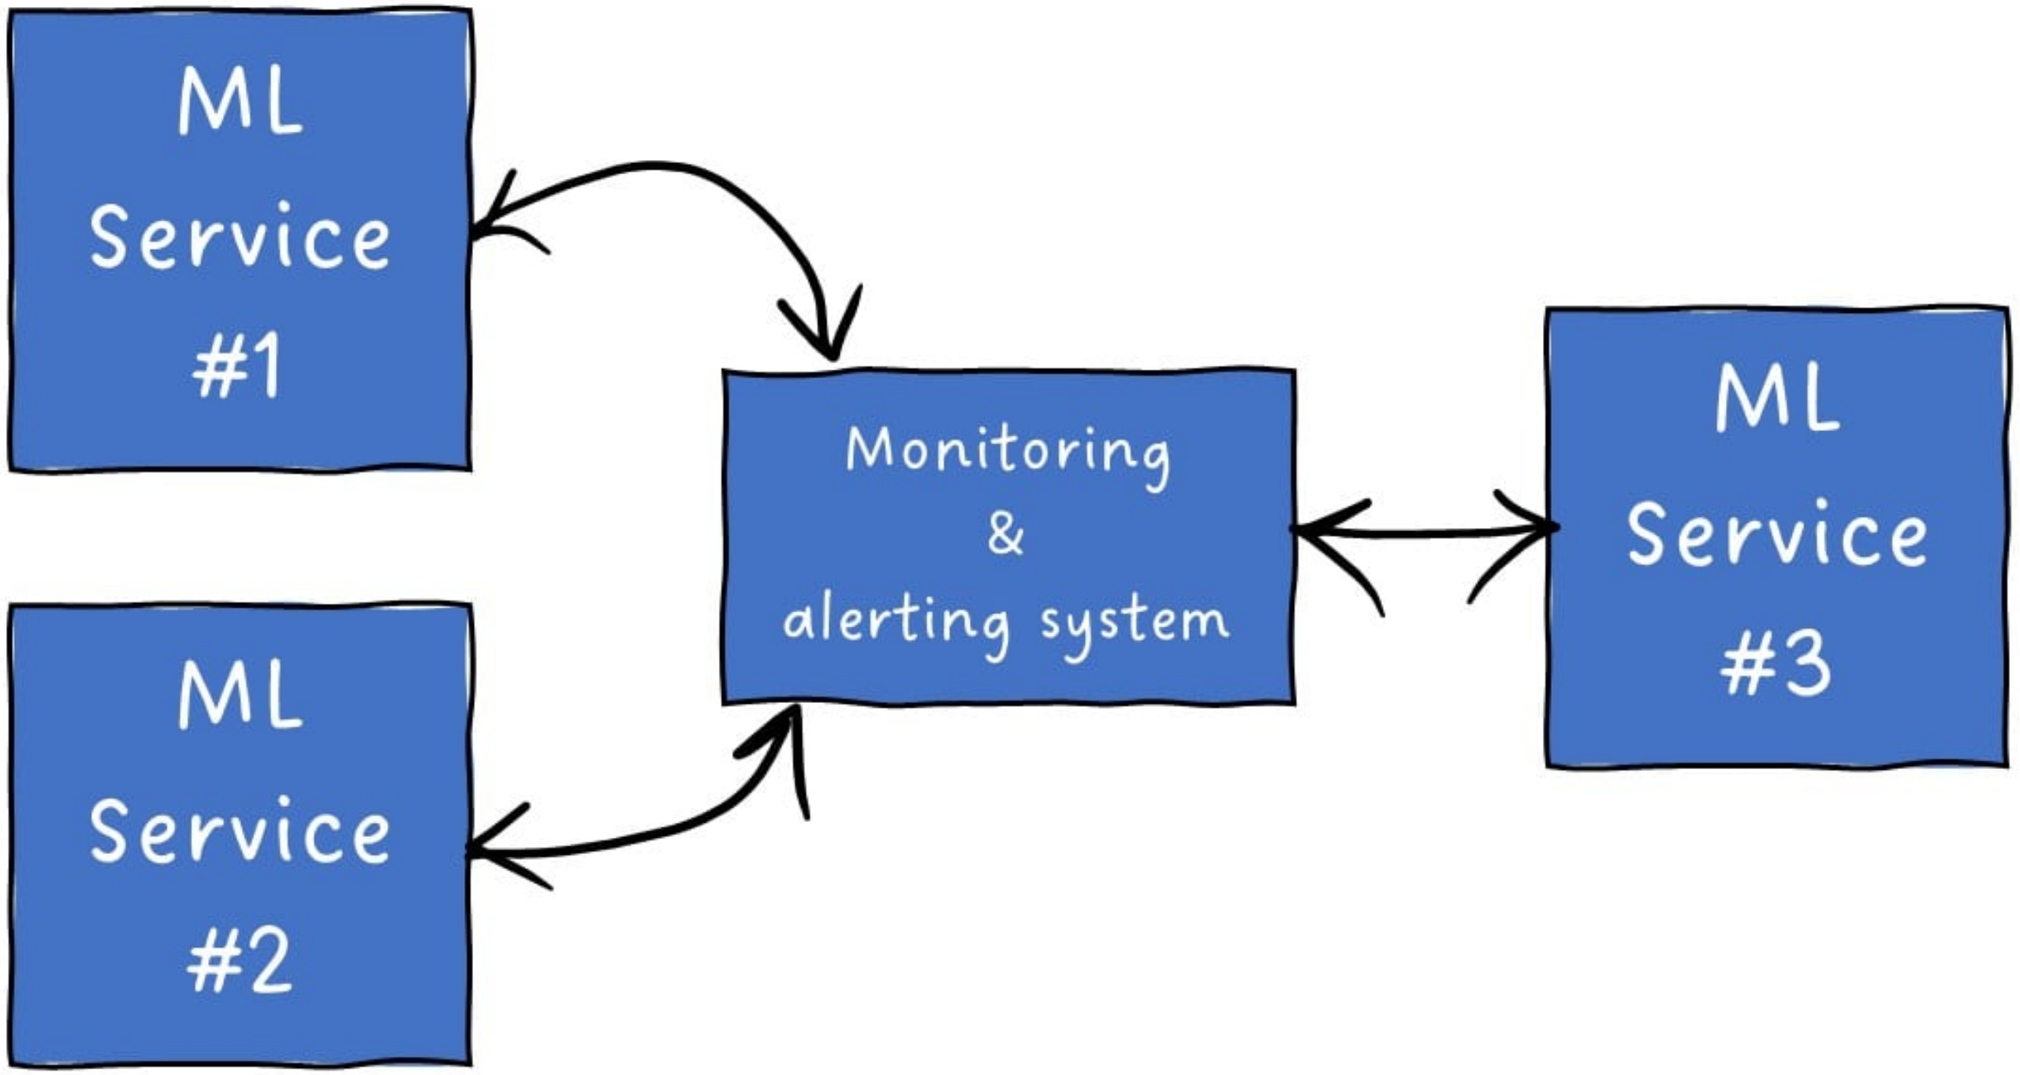
ML
Service
#1
ML
Service
#2
Monitoring
&
alerting system
ML
Service
#3

Let's practice!

MLOPS DEPLOYMENT AND LIFE CYCLING

Model maintenance

MLOPS DEPLOYMENT AND LIFE CYCLING

Nemanja Radojkovic
Senior Machine Learning Engineer

datacamp

Previously...

- How to catch anomalies and disturbances?

Focus of this lesson: ML model failure

Model performance deteriorated?

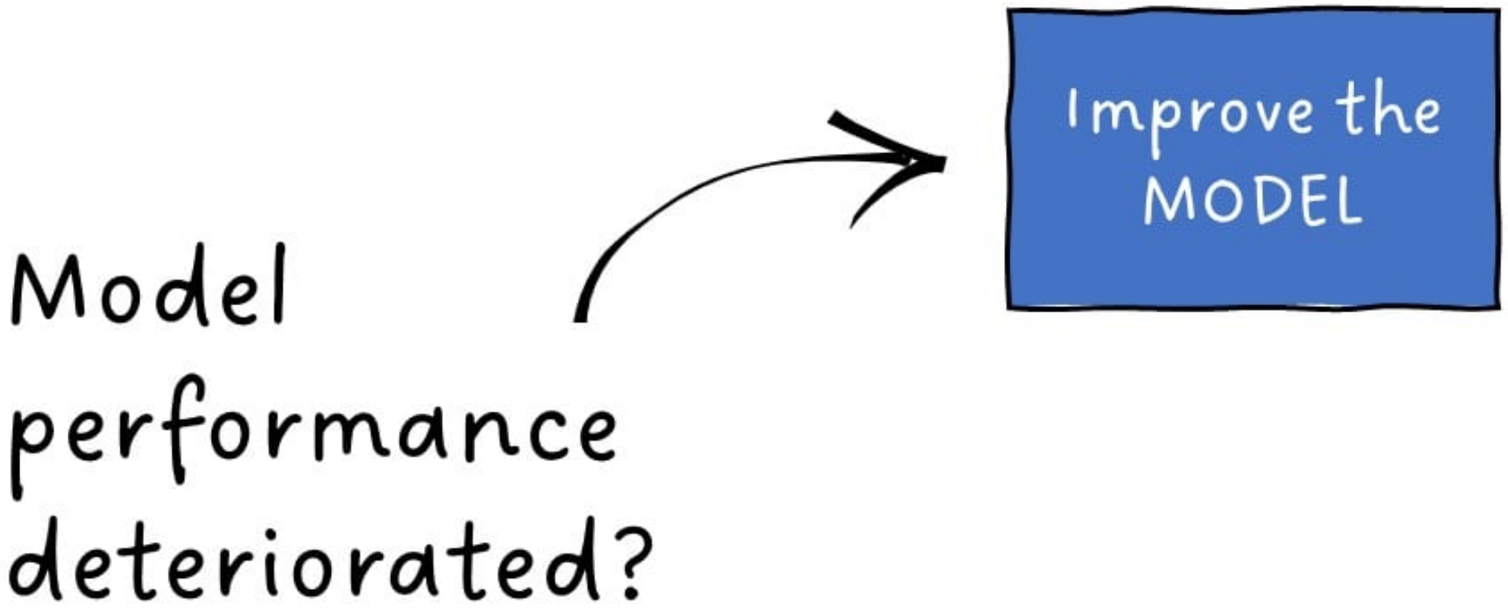
Model performance deteriorated?
Improve the MODEL

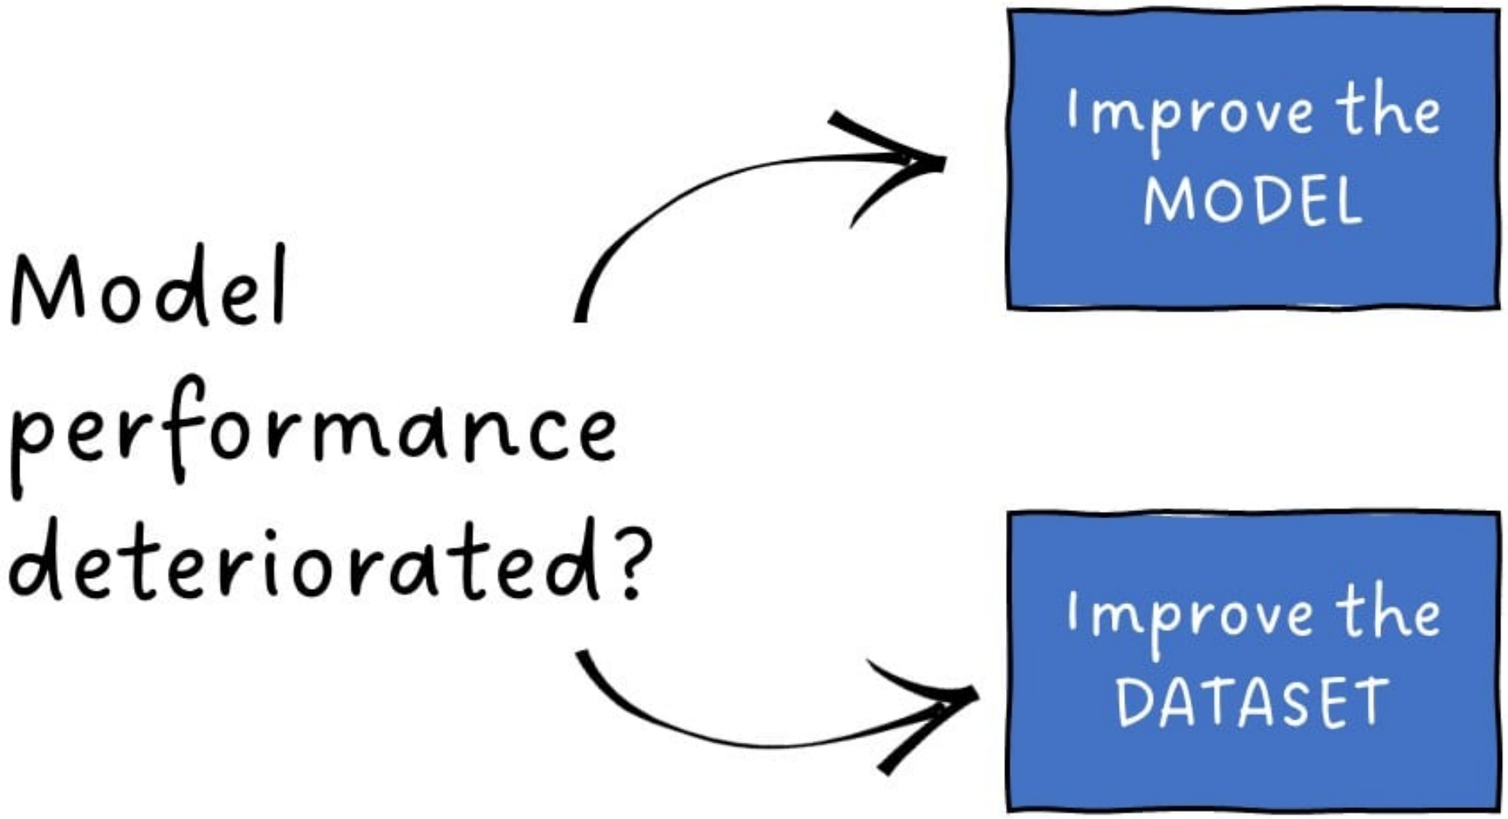
Model performance deteriorated?
Improve the MODEL
Improve the DATASET

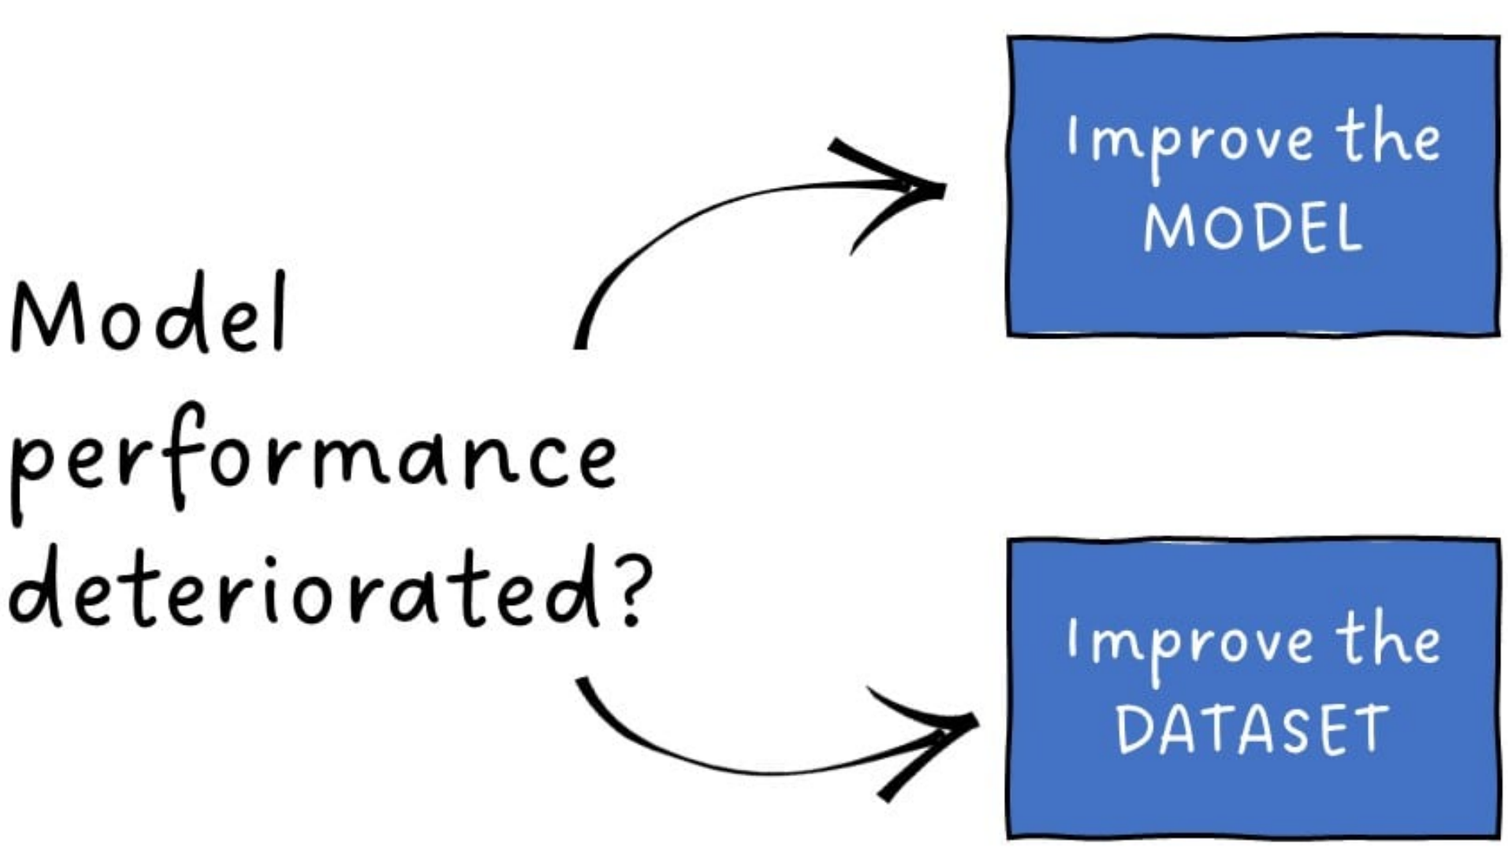

MODEL-centric approach to ML development

- typical for competitions, where dataset is fixed
- try new models
- try new DERIVED features

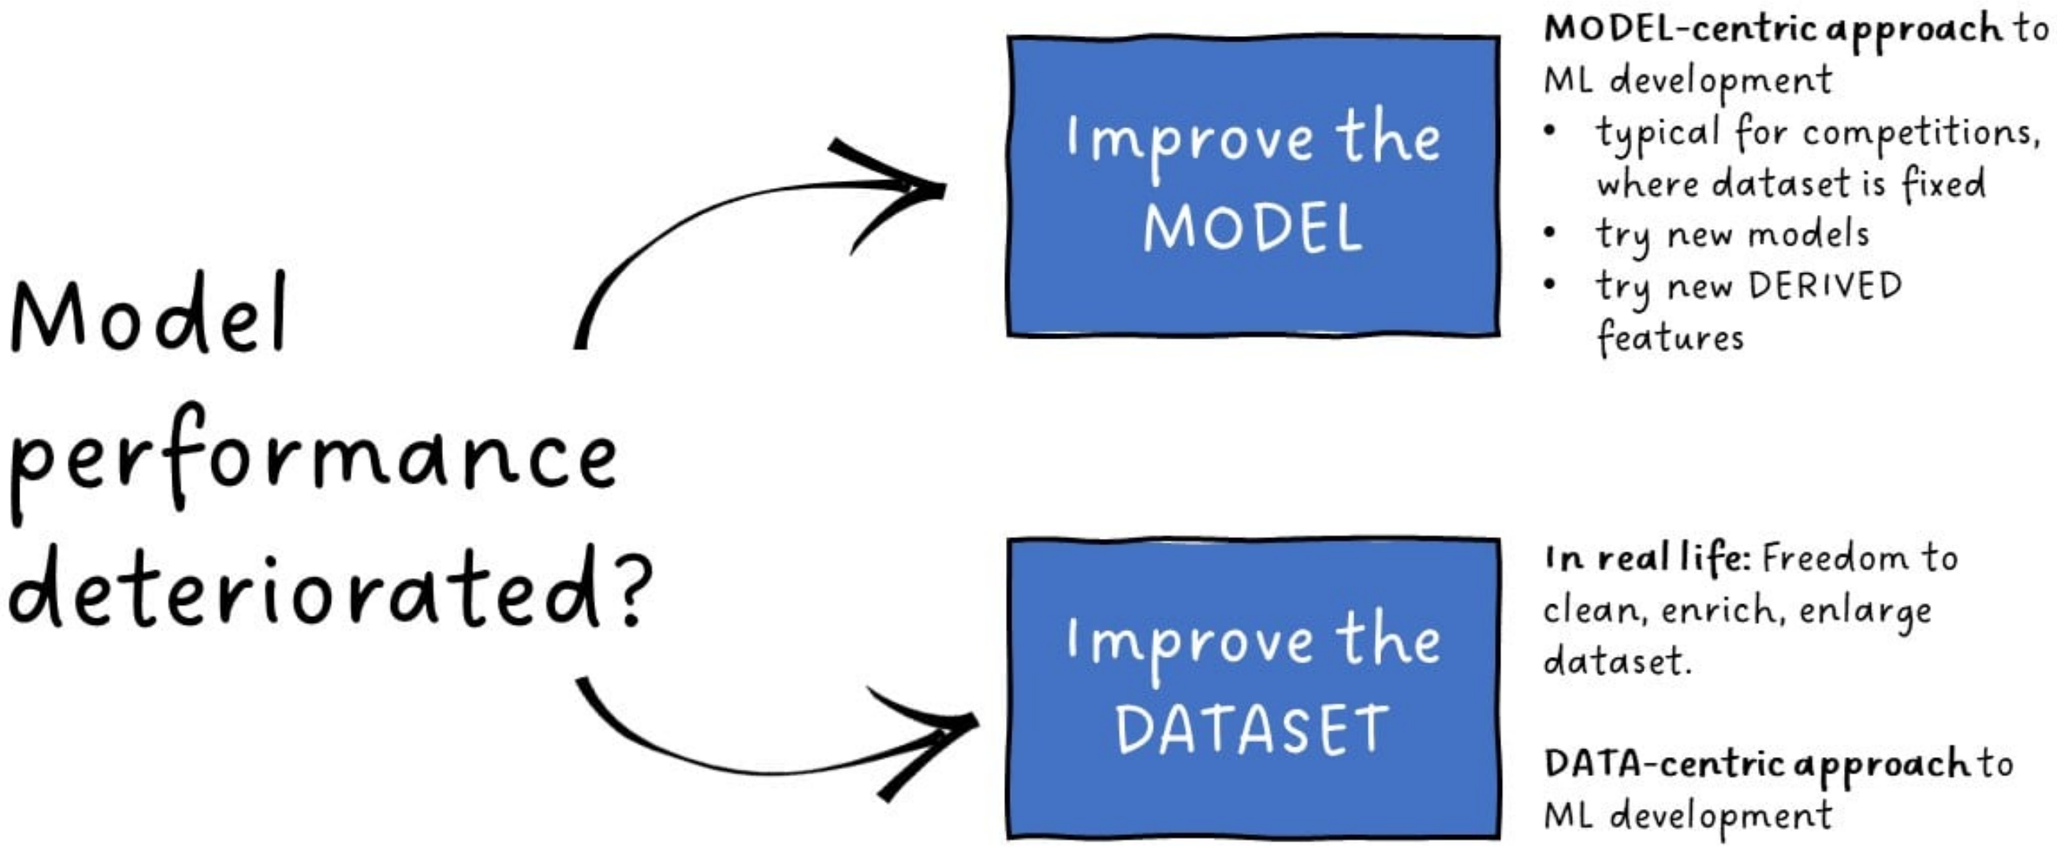

MODEL-centric approach to ML development

- typical for competitions, where dataset is fixed
- try new models
- try new DERIVED features

In real life: Freedom to clean, enrich, enlarge dataset.

DATA-centric approach to ML development

[1] https://datacentricai.org/

Quality above quantity

- Get more features with relevant information
- Get better labels

Label == Target variable value in the training set

Label quality == label closeness to the ground truth

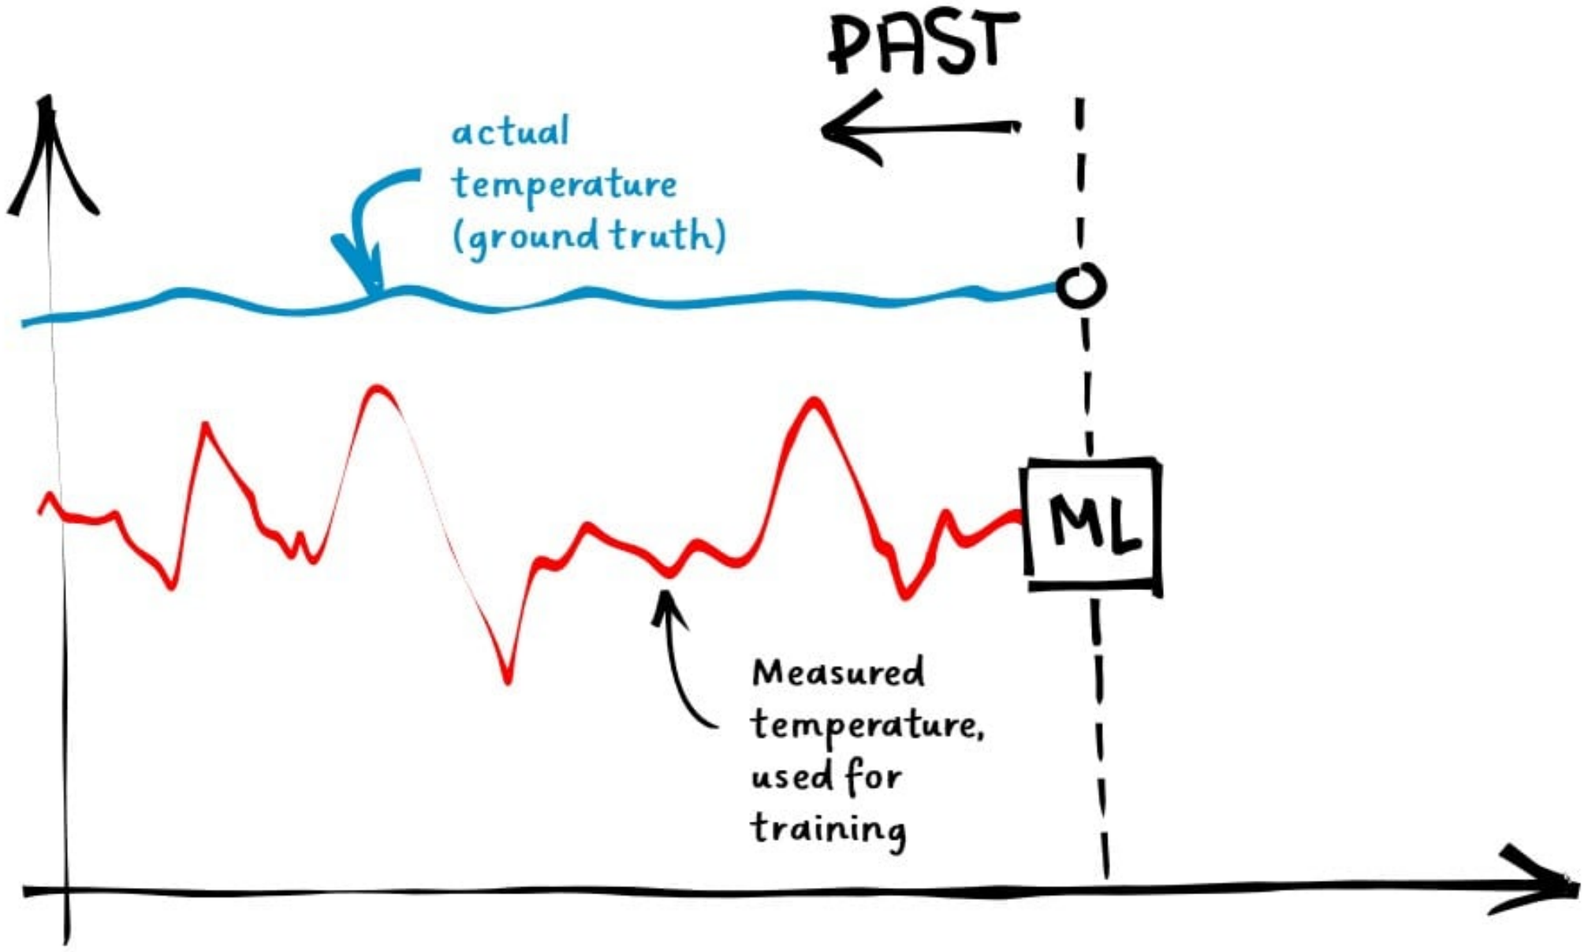
PAST
actual
temperature
(ground truth)
ML
Measured
temperature,
used for
training

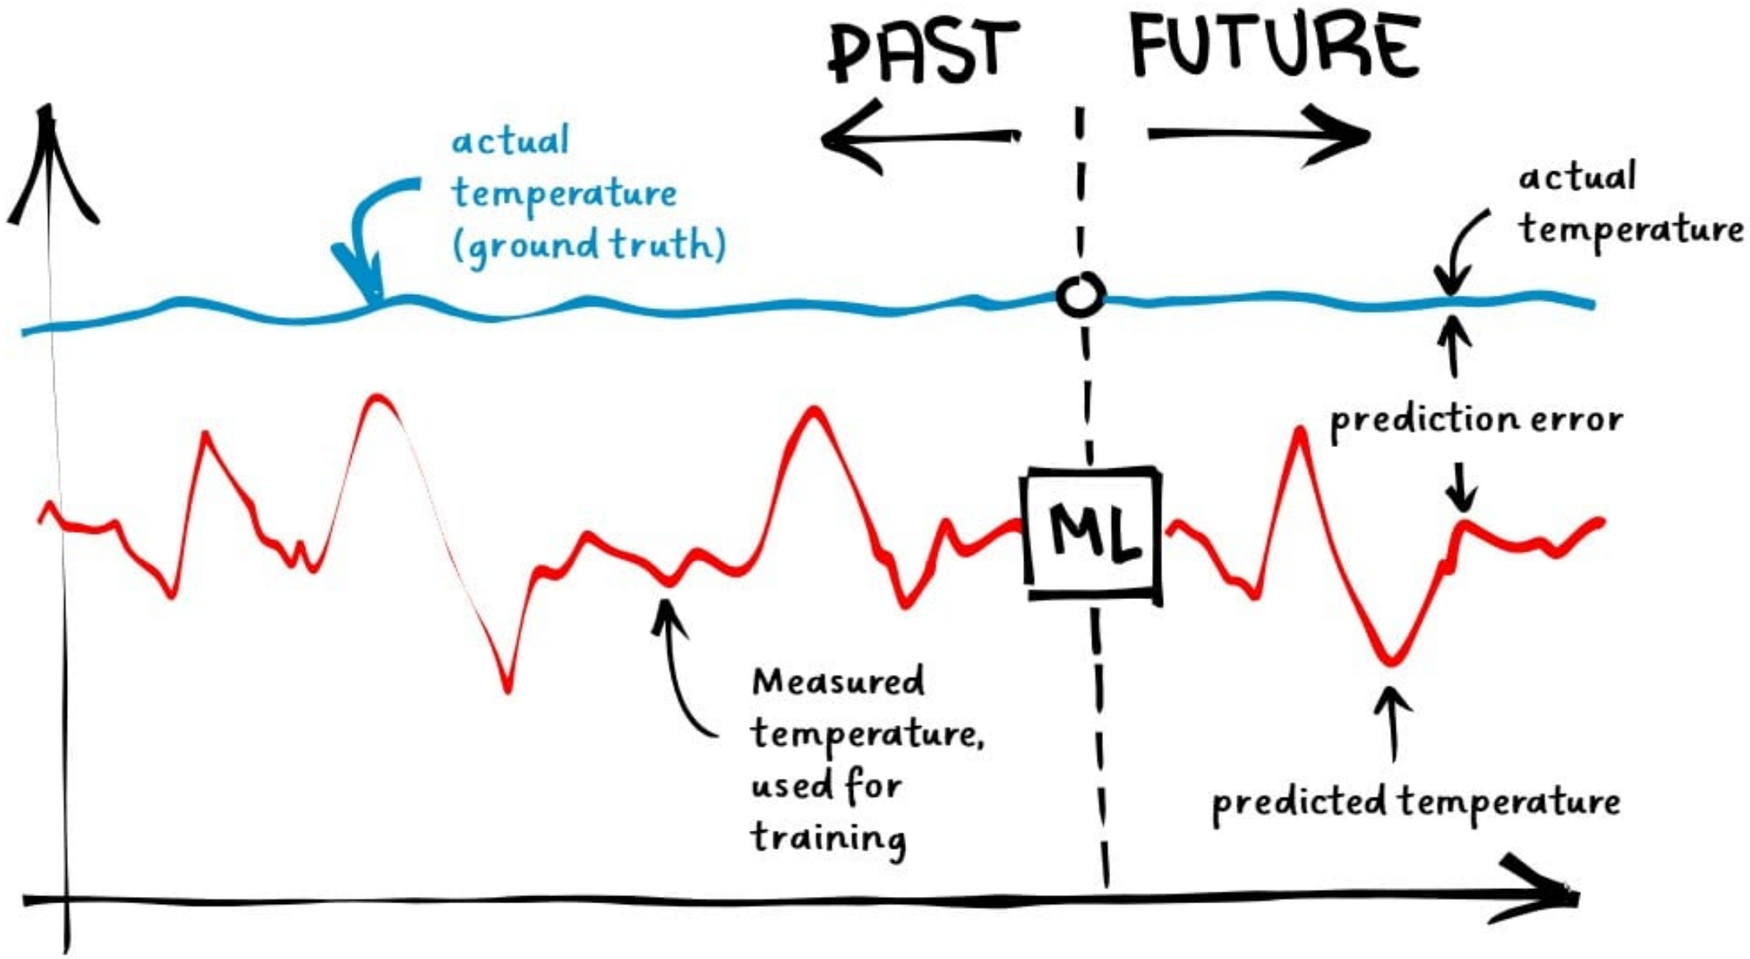
PAST
FUTURE
actual temperature (ground truth)
actual temperature
prediction error
ML
Measured temperature, used for training
predicted temperature

Benefits of labeling tools

Manual labeling is complex, lengthy, error-prone

Use good labeling tools!

- Labeling more efficient
- More accurate
- User interface fit for purpose
- "Label these examples first for maximum impact"
- "It seems you made a mistake here"

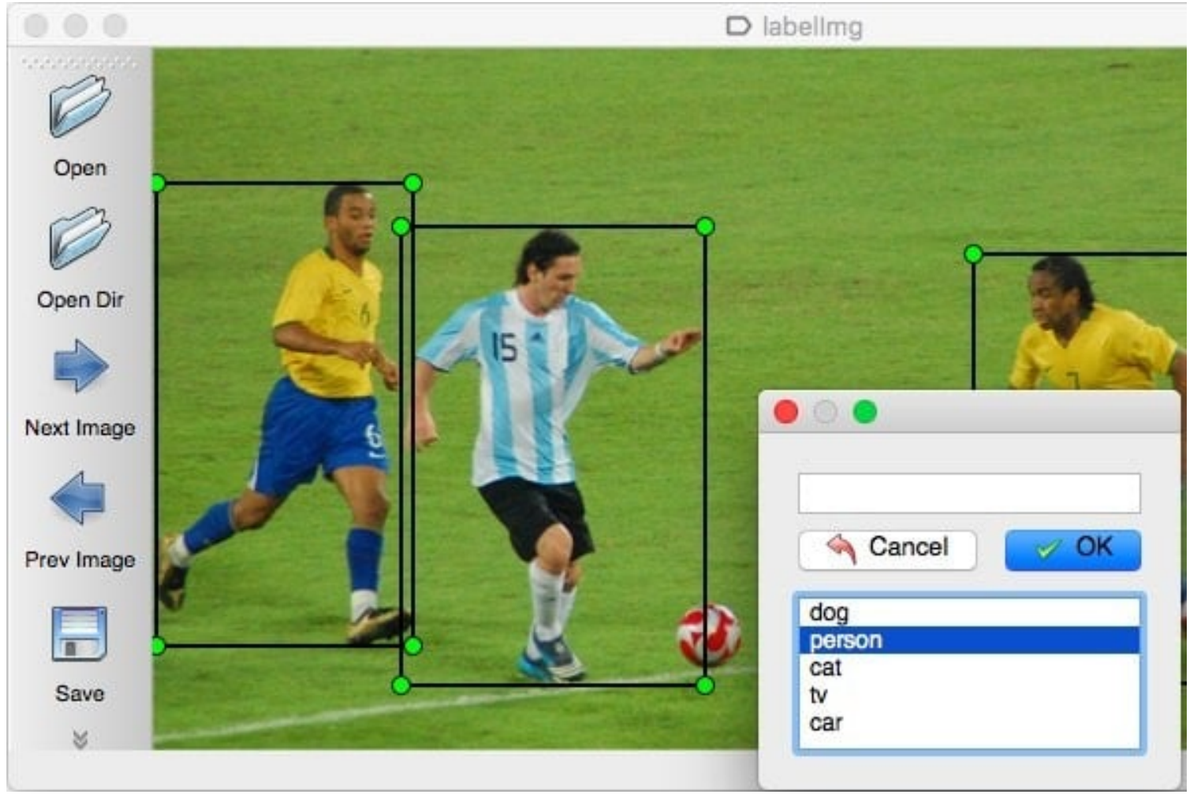

Example: Image labeling tool, for building image classifiers

Human-in-the-loop system

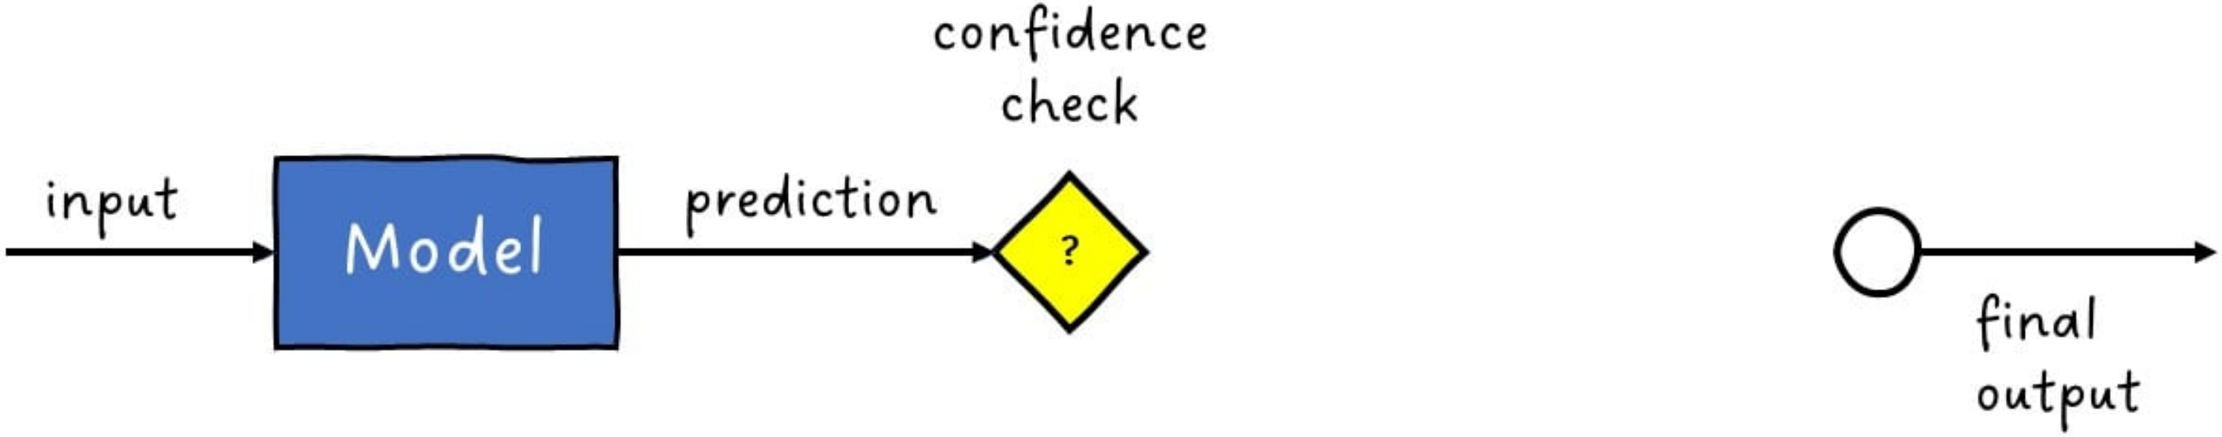

Human-in-the-loop system

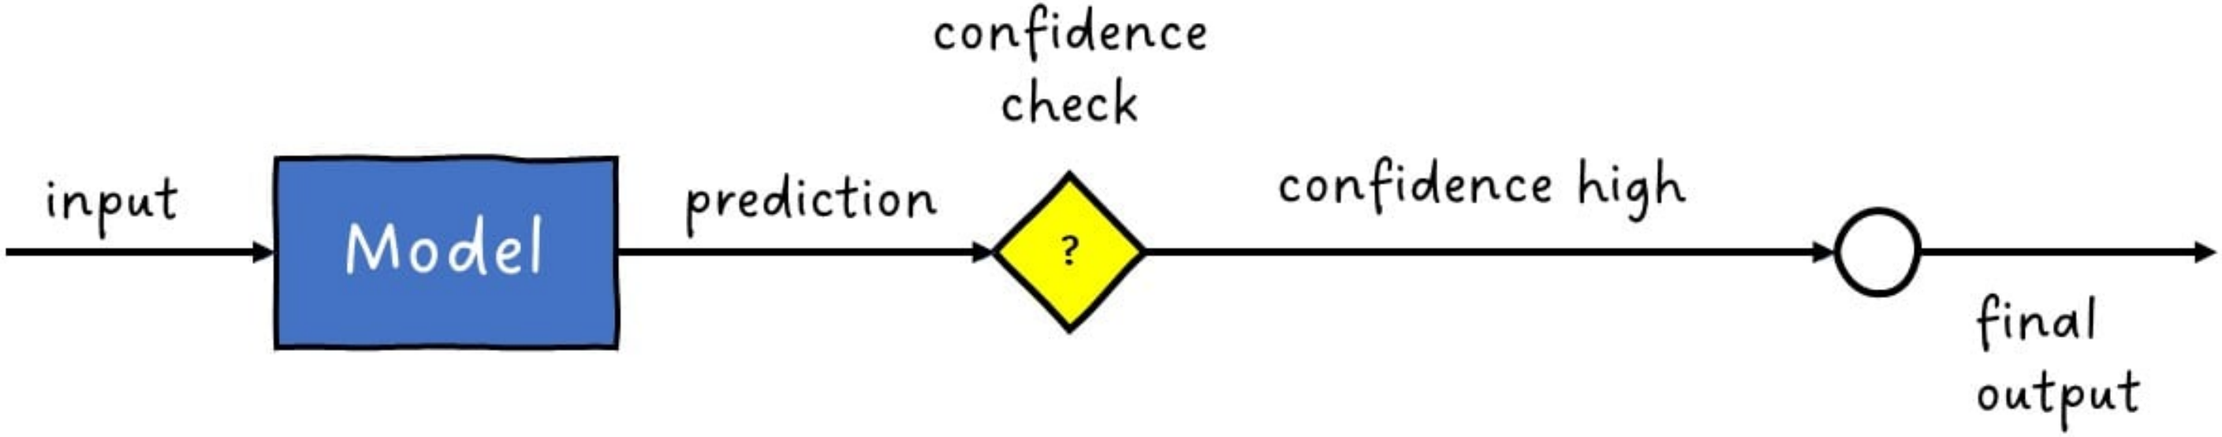

Human-in-the-loop system

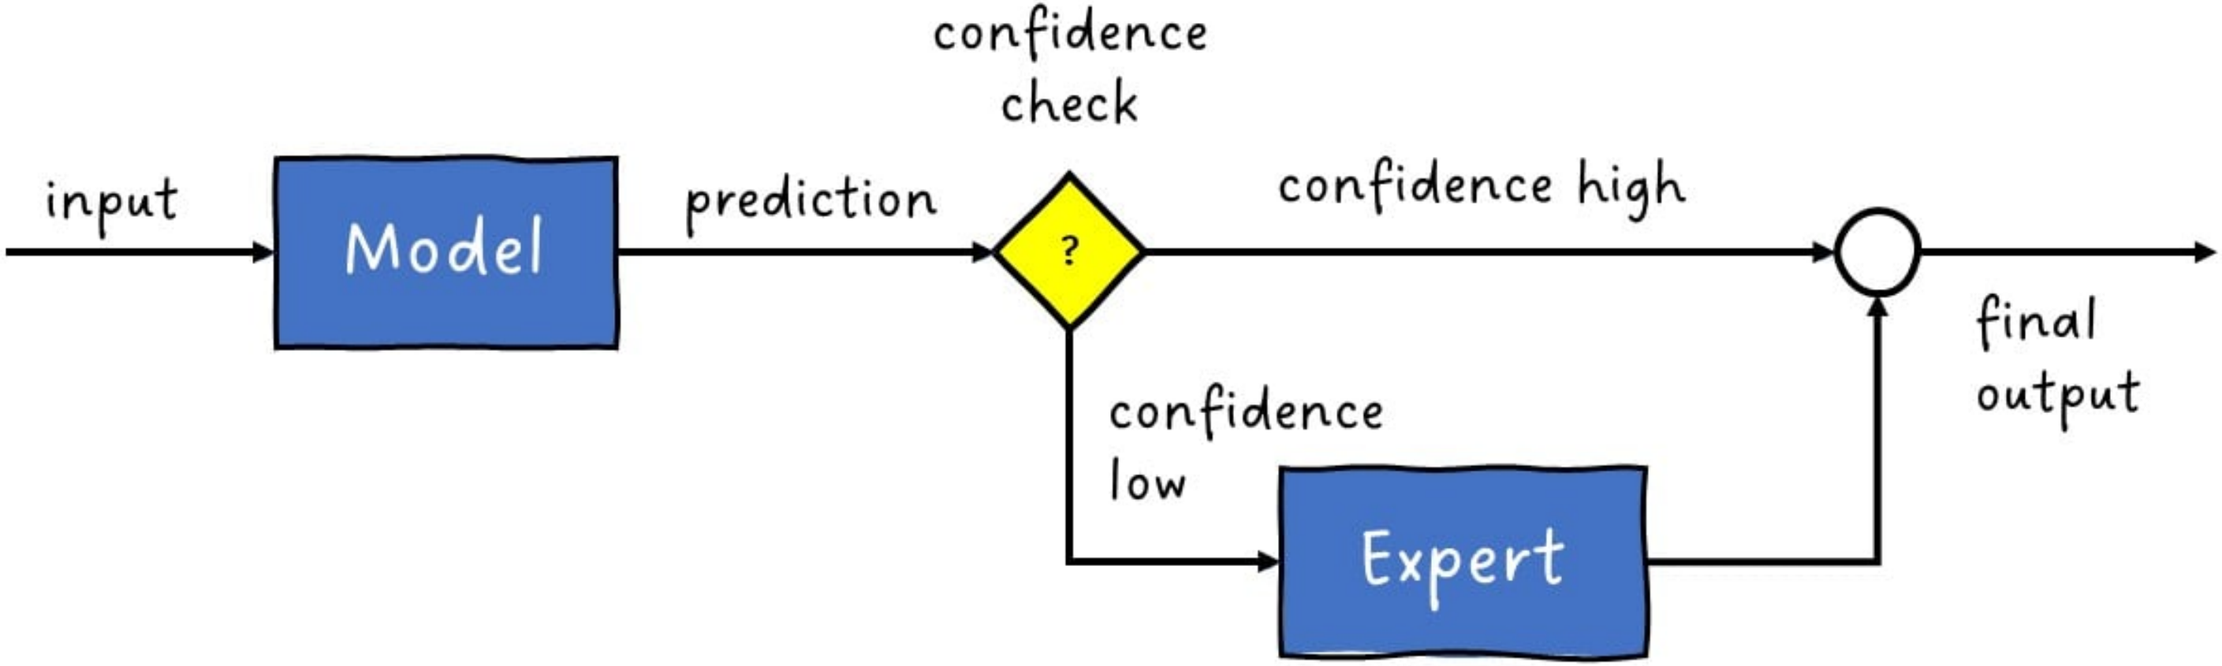

Human-in-the-loop system

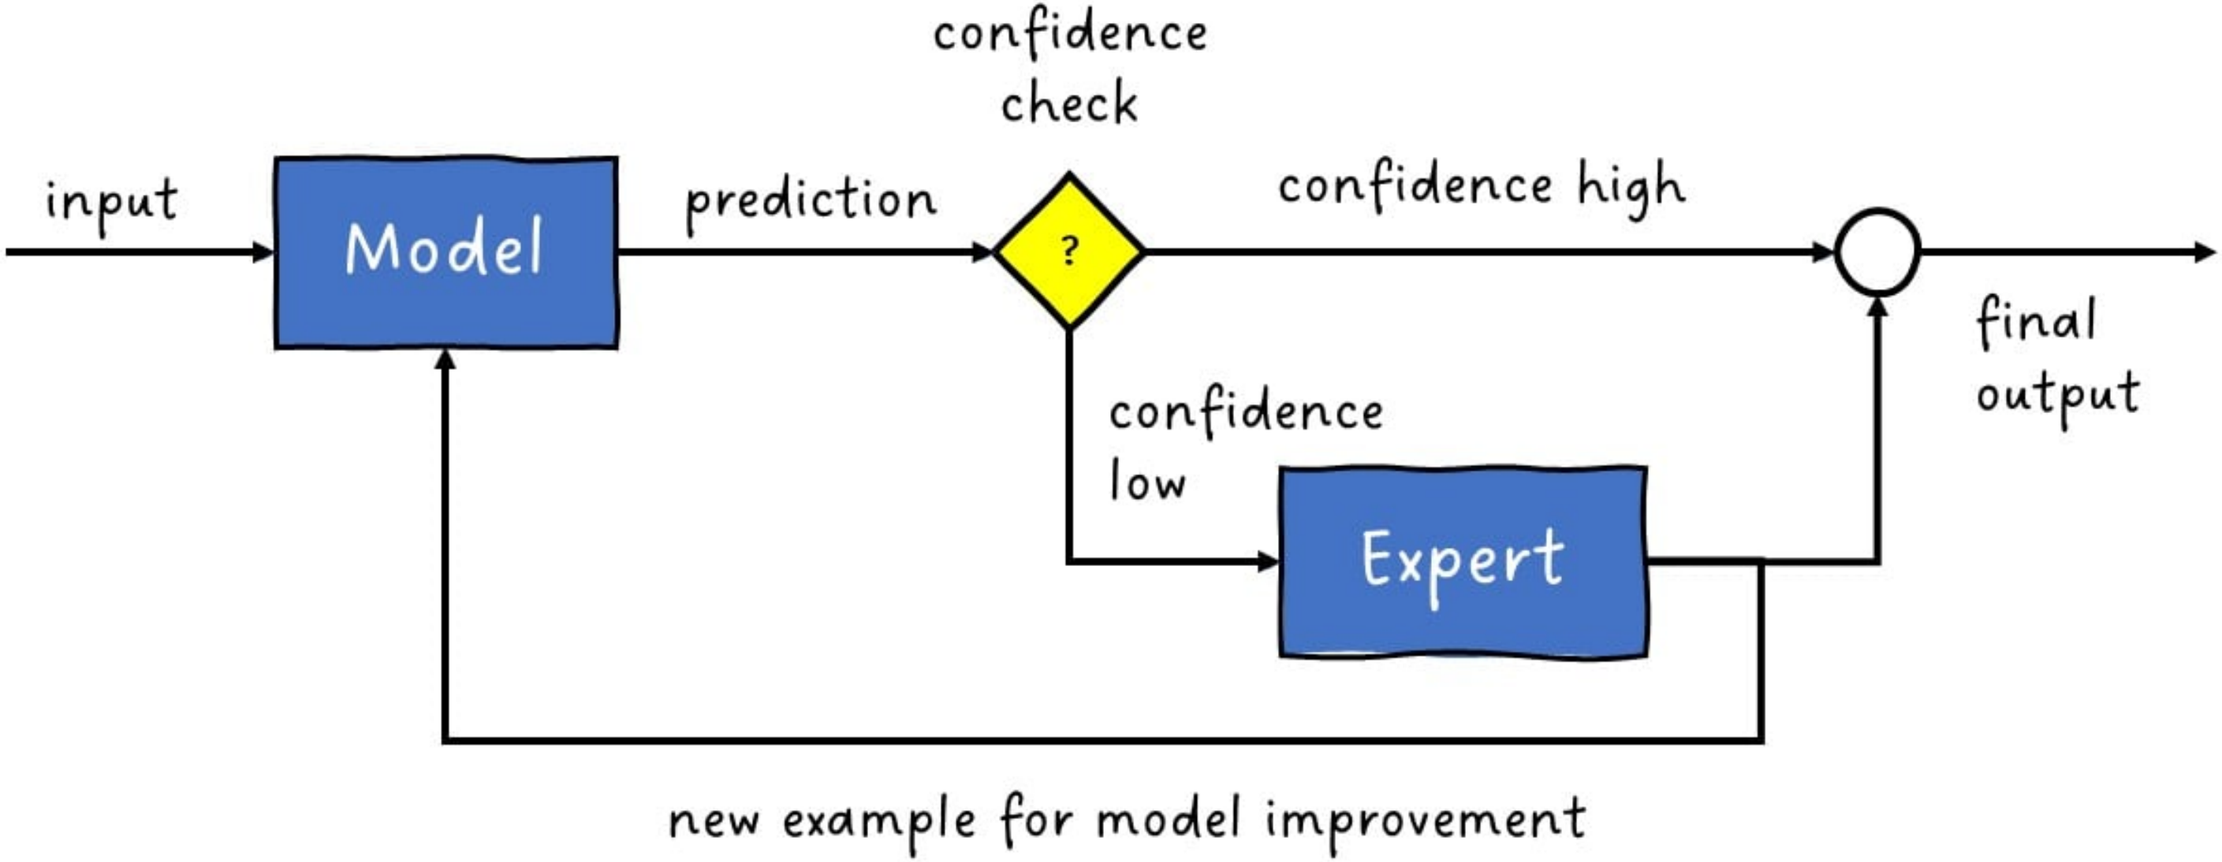

When the labels arrive

- Run ML build pipeline => New model
- Check if new model is better than the old one
  - YES => test and deploy
  - NO => keep searching,

Keep searching

- Try new models
- Try new features
- Try new data sources

Immensely helpful: Metadata store (MLFlow Tracking, etc)

- Document the model selection journey
- Avoid repeating same experiments

In any case: *MLOps helps us maintain our model in the fastest, most efficient way*

Let's practice!

MLOPS DEPLOYMENT AND LIFE CYCLING

Model governance

MLOPS DEPLOYMENT AND LIFE CYCLING

Nemanja Radojkovic
Senior Machine Learning Engineer


NHS
ArtiQ Wins UK Government Funding for its AI Algorithms to Support Timely Diagnosis of Lung Diseases
by ArtiQ Jun 16, 2021


Business Of Law; Internet Law & Cyber-Security
Law Bots: How AI Is Reshaping the Legal Profession
10 Min Read
By: Matthew Stepka | February 21, 2022


Artificial Intelligence and Machine Learning
9 min read
5 Ways AI is Transforming the Finance Industry
Check out the easy ways by which artificial intelligence can transform the finance industry.
Pinakin Ariwala
Updated on Oct 14


2 minute read · November 30, 2022 5:29 PM GMT+1 · Last Updated a day ago
Honda to develop advanced level 3 self-driving technology by 2029
Reuters

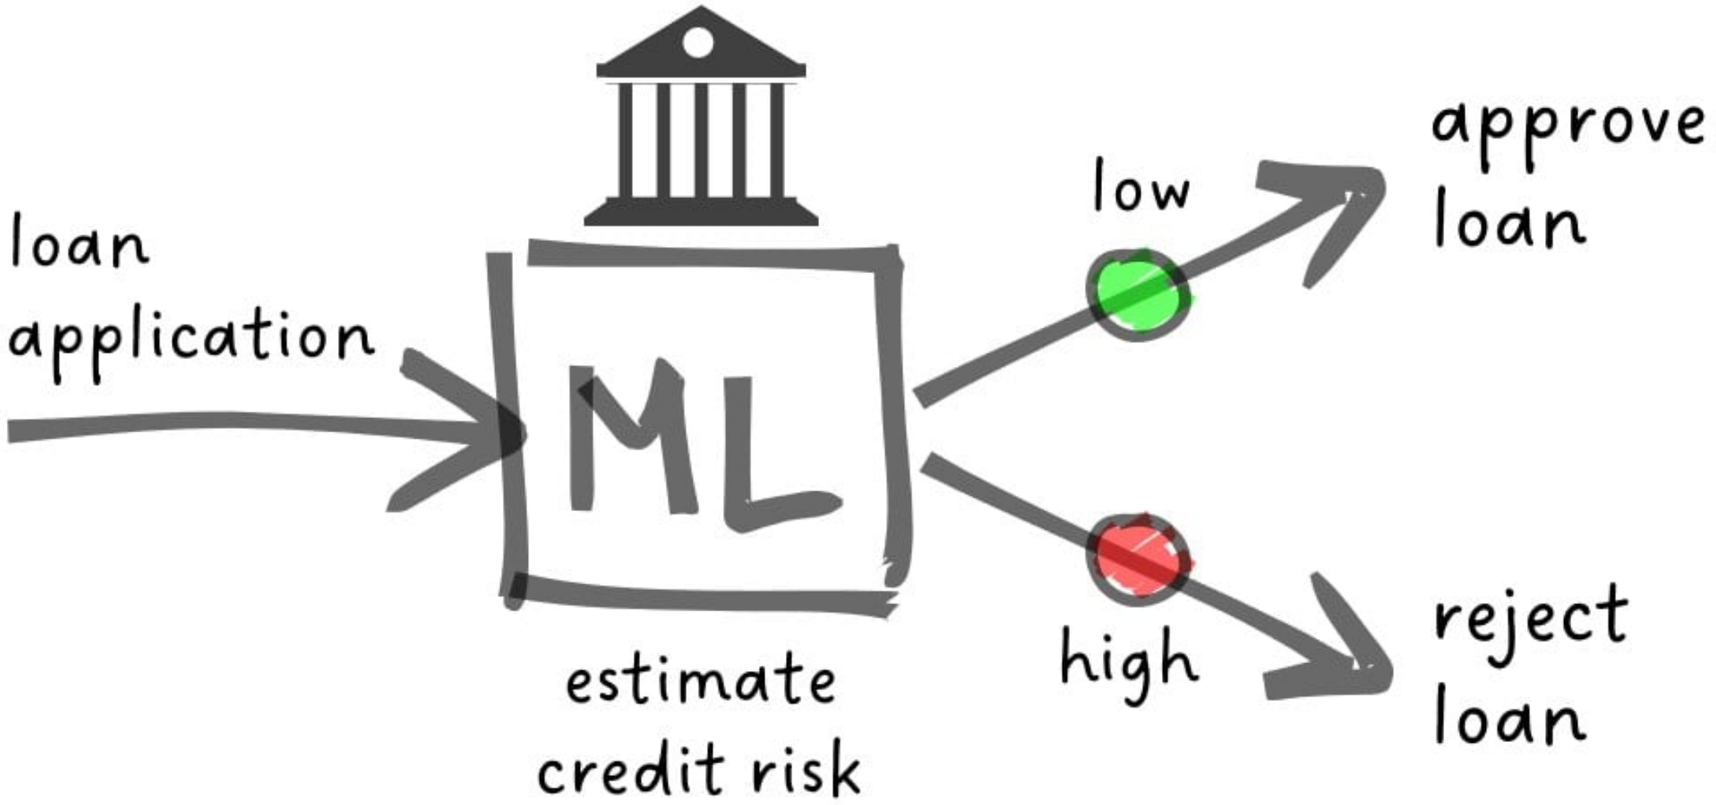
loan
application
ML
estimate
credit risk
low
approve
loan
high
reject
loan

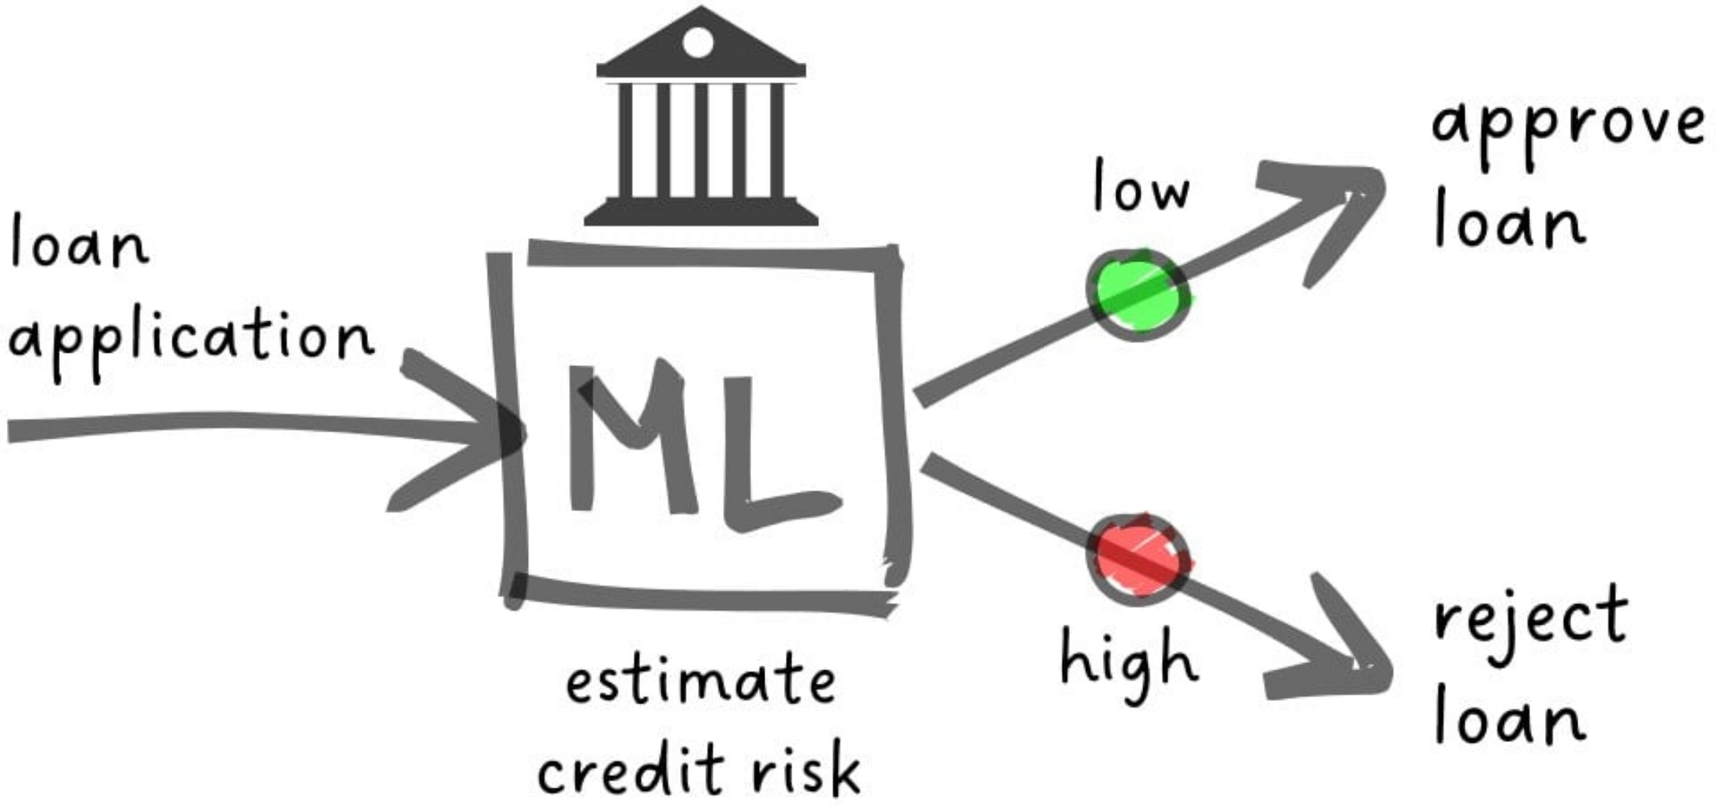

if predicted
but actual
x N

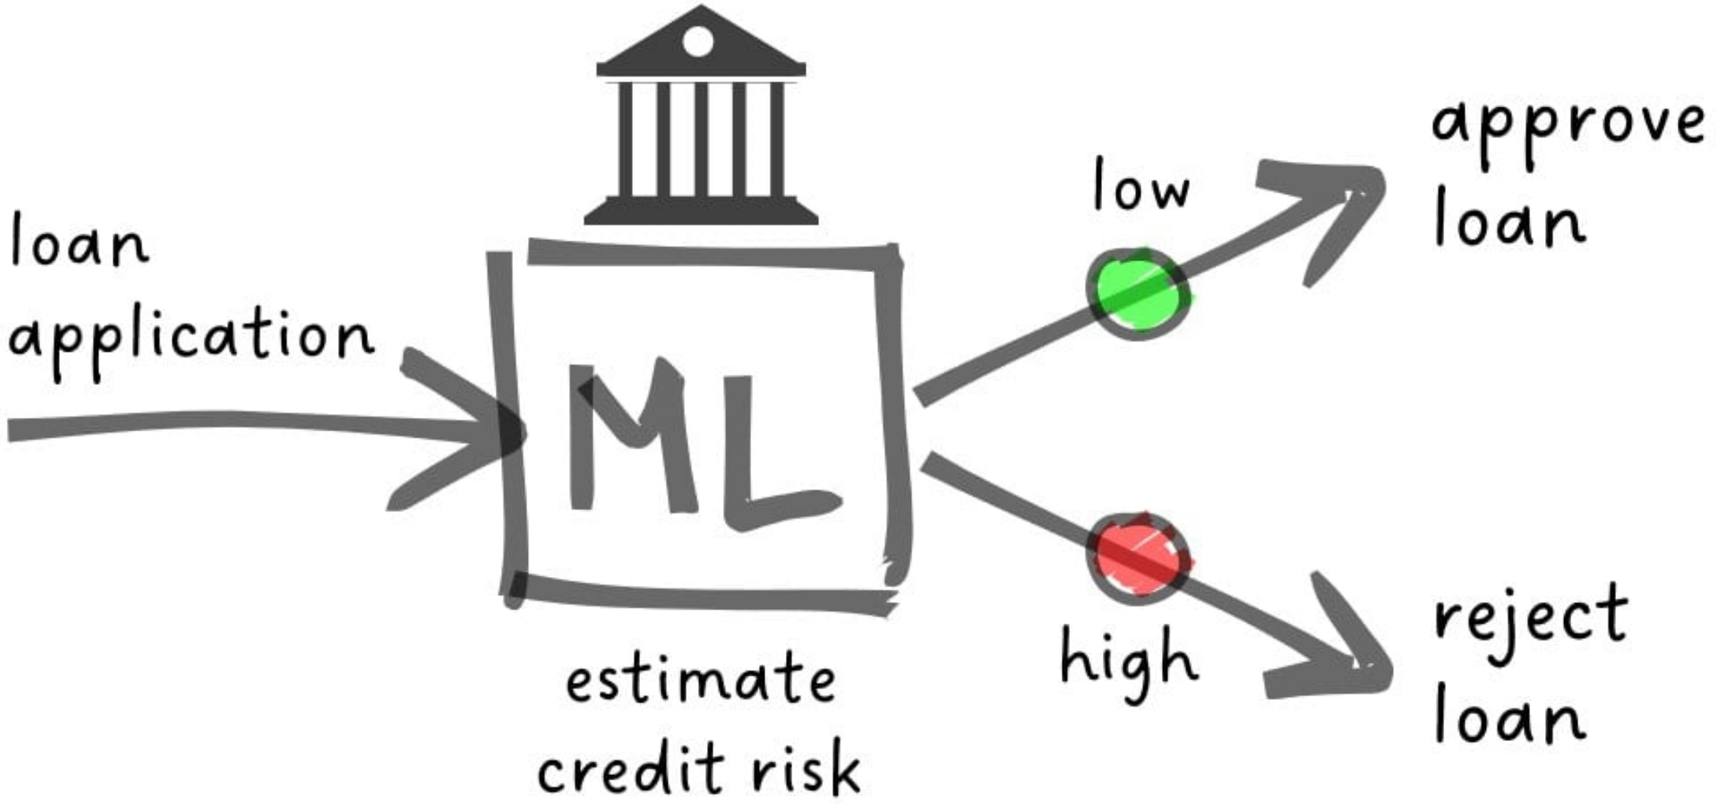
loan
application
ML
estimate
credit risk
low
approve
loan
high
reject
loan


if predicted
but actual
x N

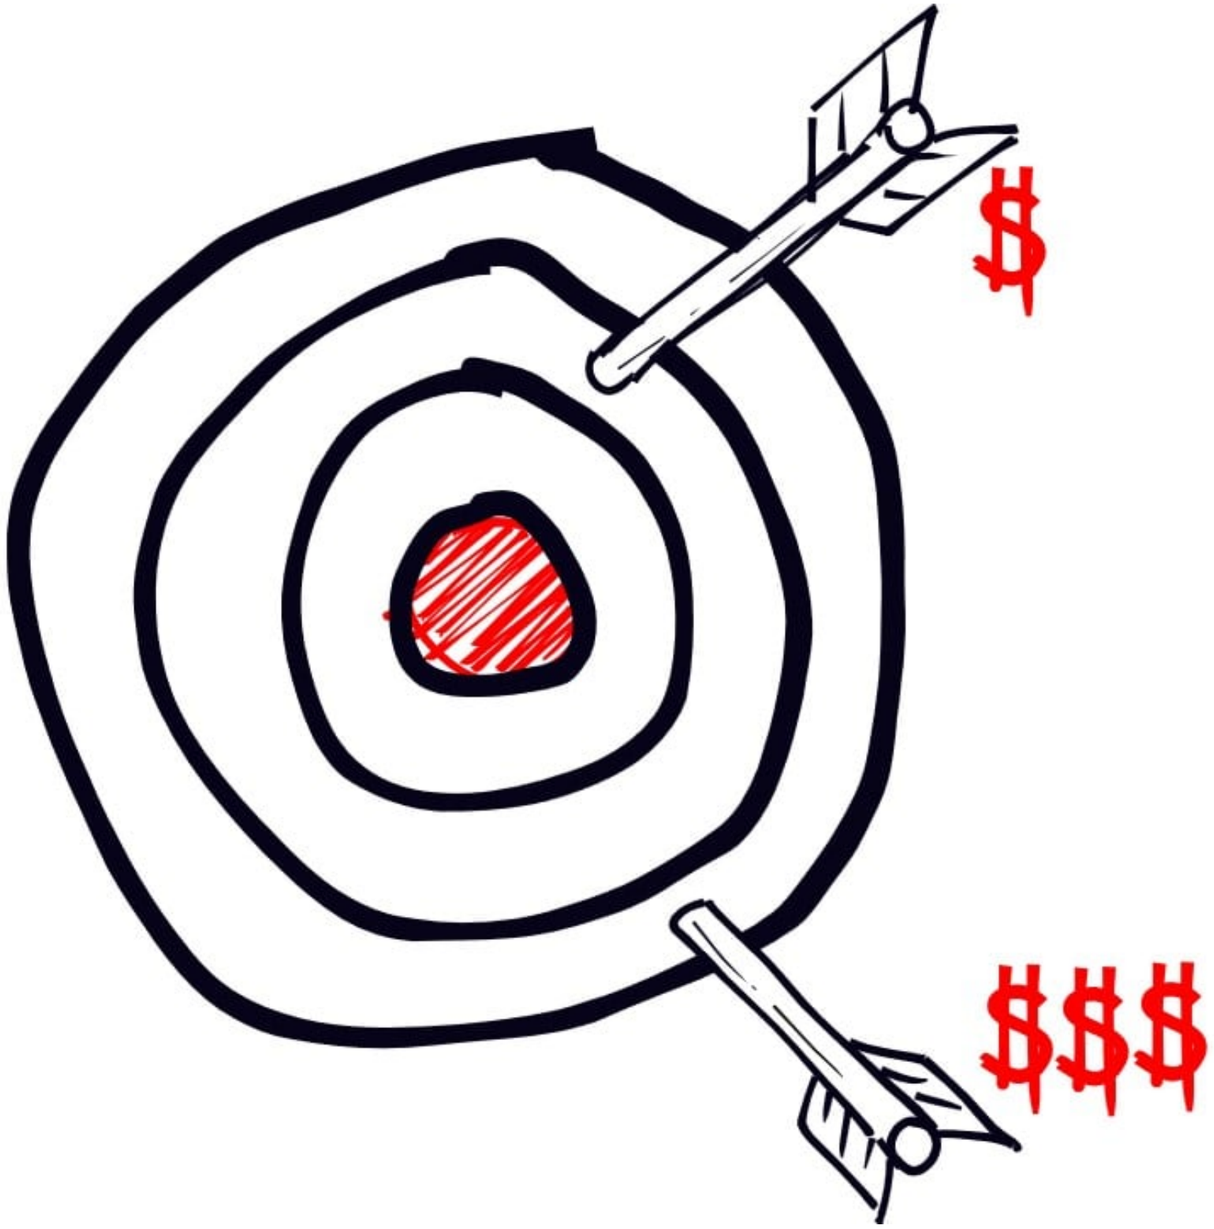

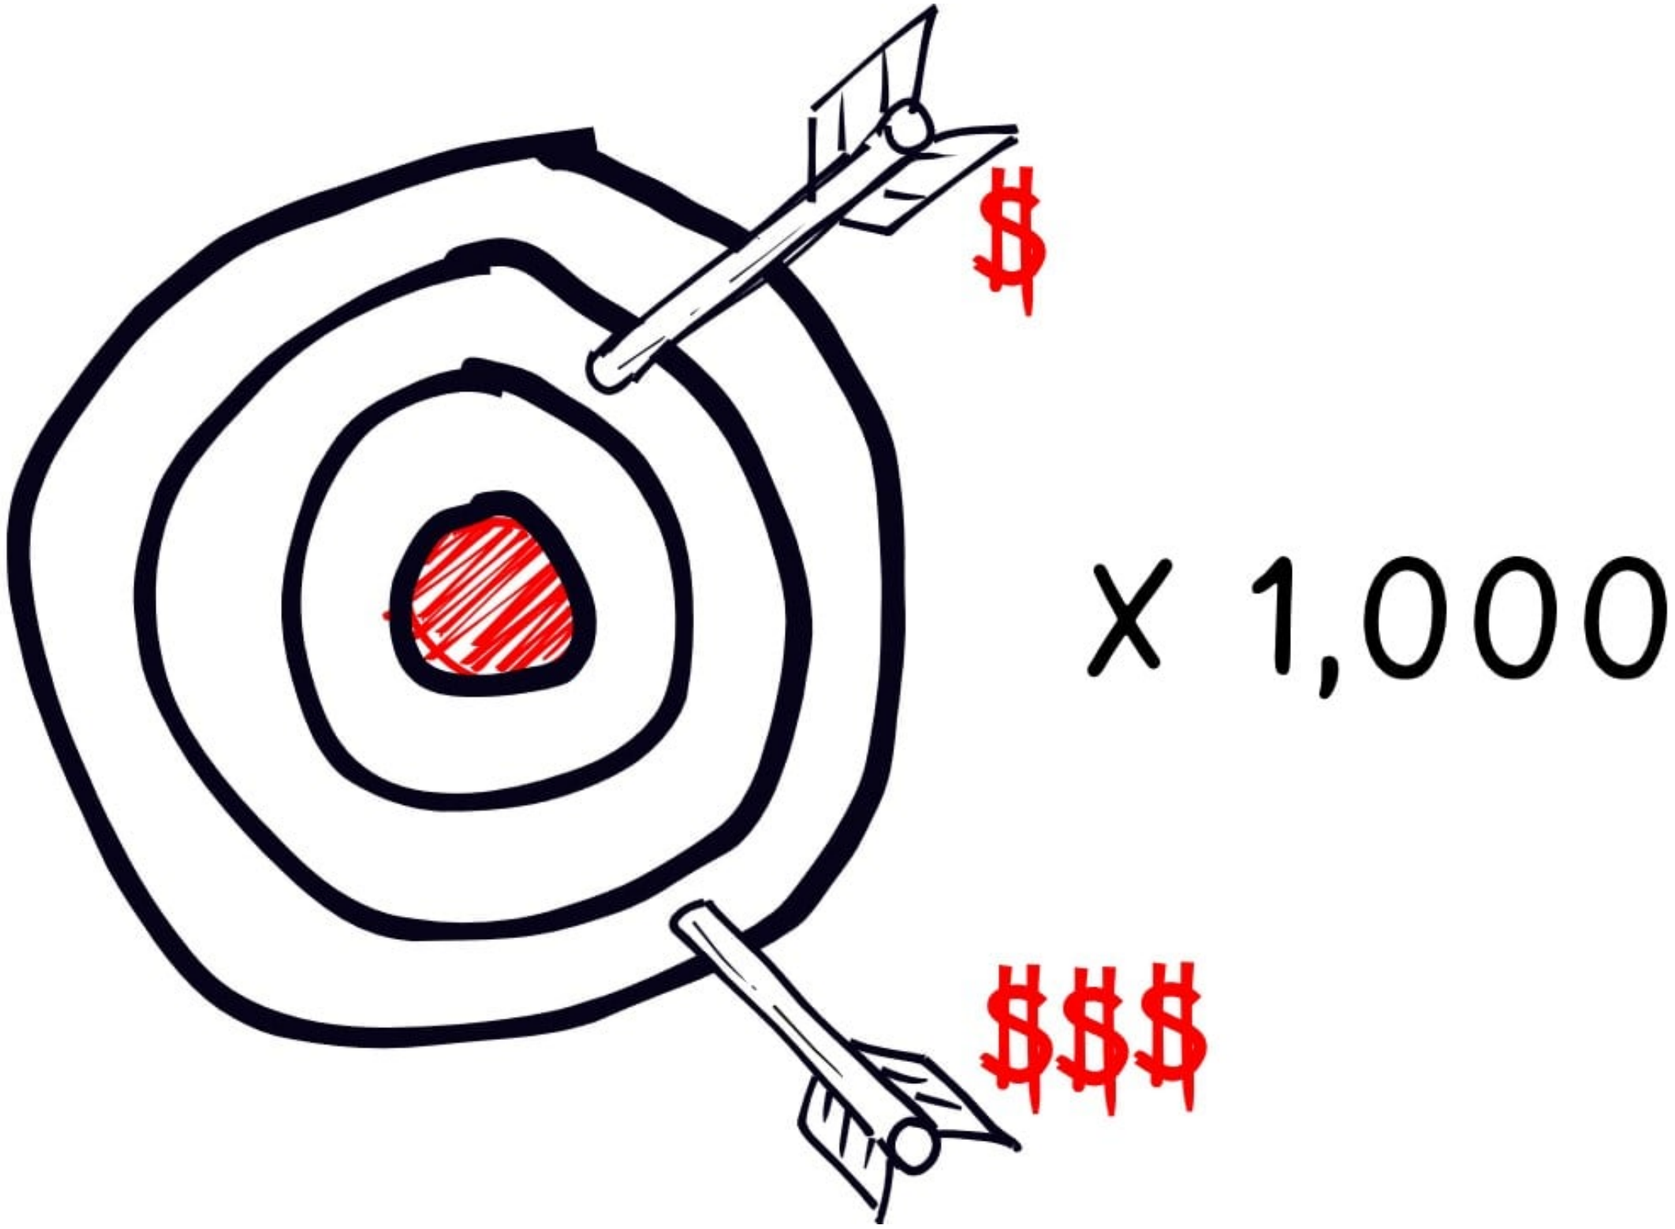
$
$$$
X 1,000

Governance

> AI/ML model governance is the overall process for how an organization controls access, implements policy, and tracks activity for models and their results.

> Effective model governance is the bedrock for minimizing risk to both an organization's bottom line and to its brand.

~ *DataRobot.com*

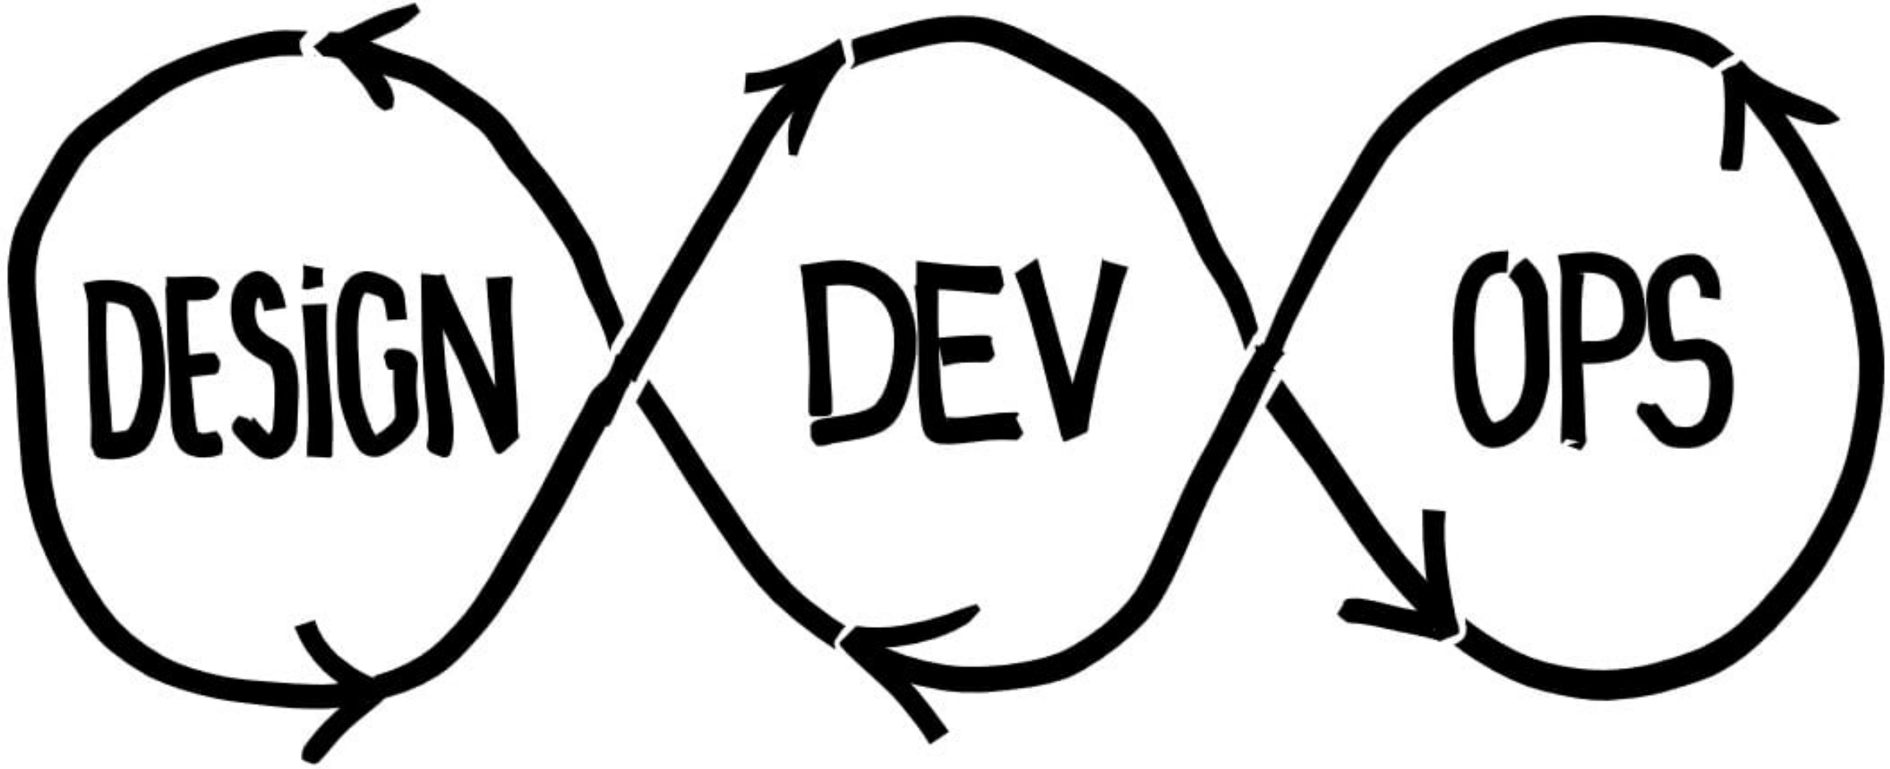
DESIGN
DEV
OPS

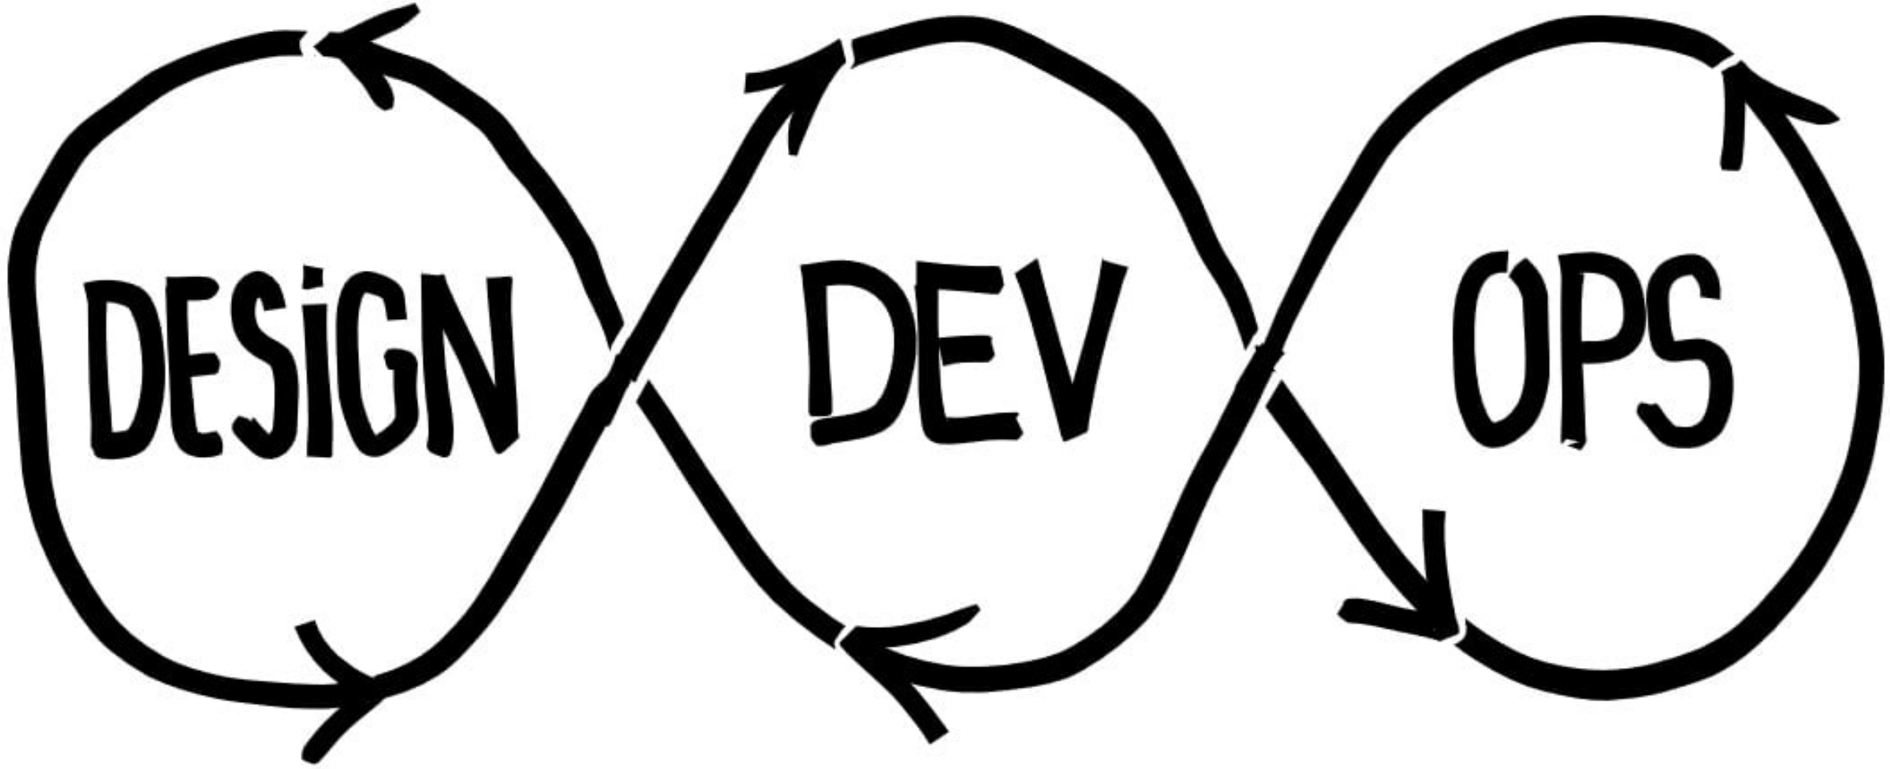

- Is it ethical to use ML?
- Private, sensitive data?
- Unethical model bias?
- ...

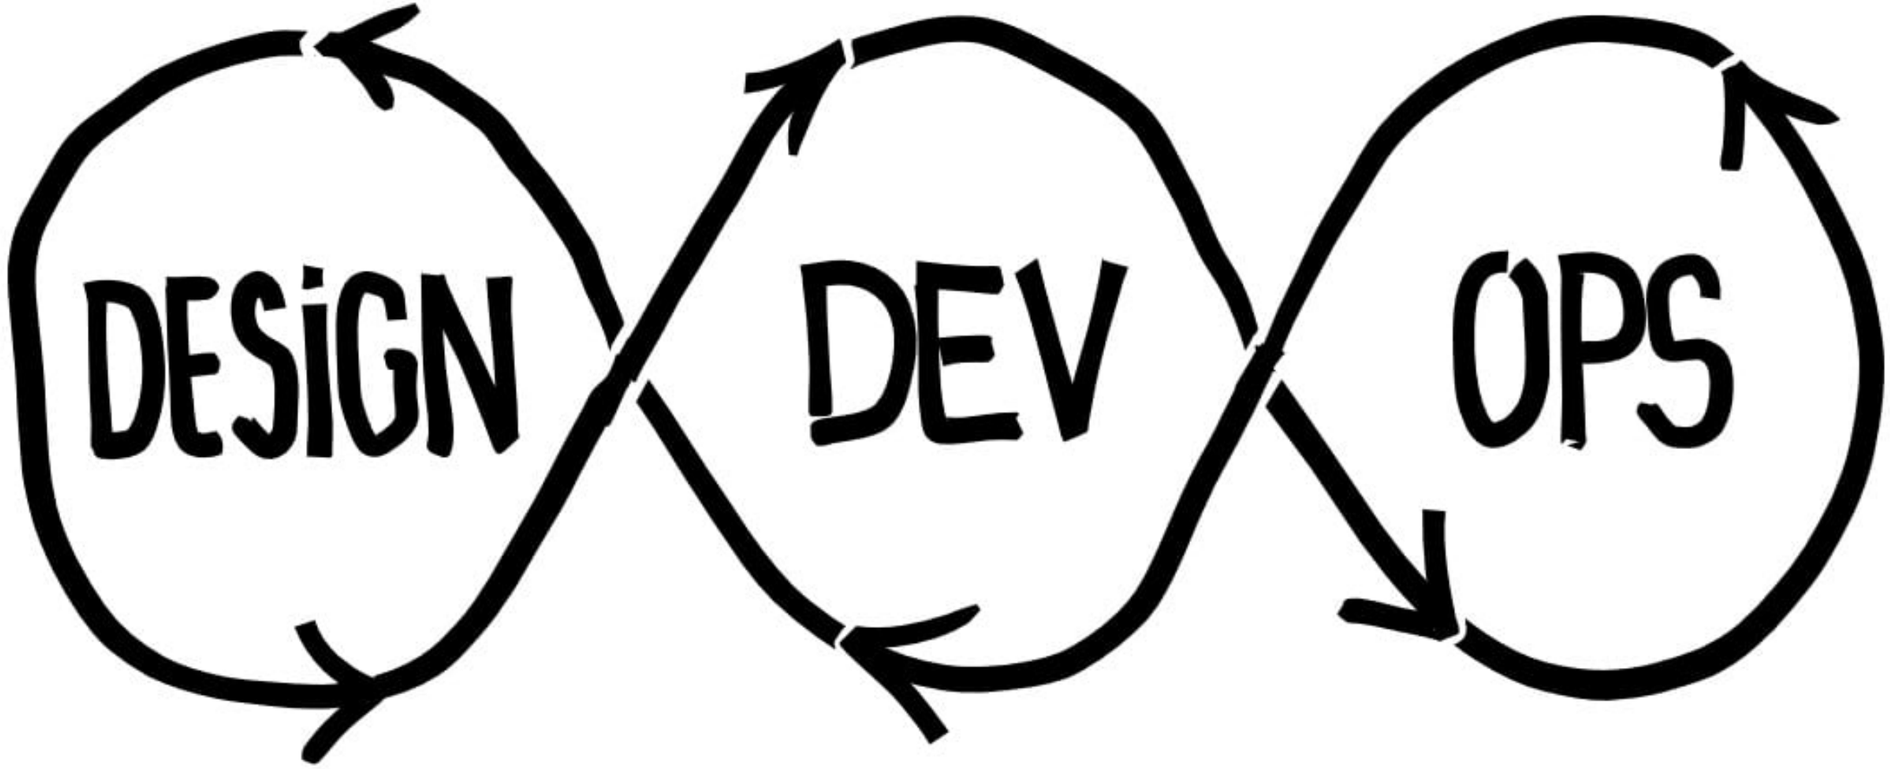

- Is it ethical to use ML?
- Private, sensitive data?
- Unethical model bias?
- ...

- Model selection docs
- Training data prep
- Quality assurance
- Versioning
- Reproducibility
- ...

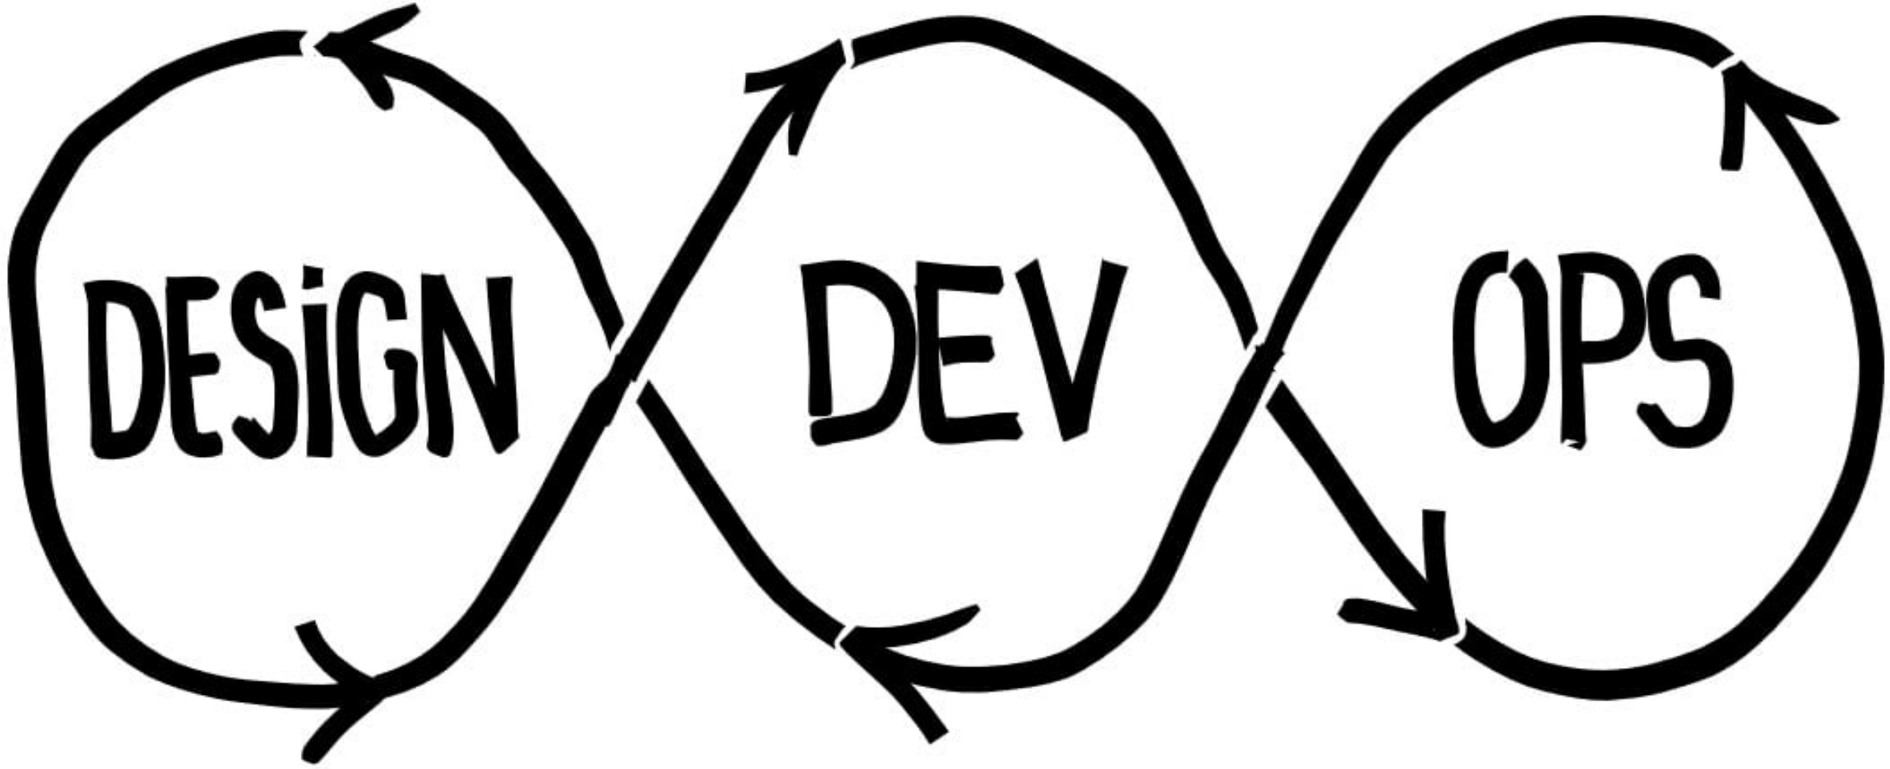

- Is it ethical to use ML?
- Private, sensitive data?
- Unethical model bias?
- ...

- Model selection docs
- Training data prep
- Quality assurance
- Versioning
- Reproducibility
- ...

- Is API secure?
- Monitoring in place?
- Alerting?
- Failure handling?
- ...

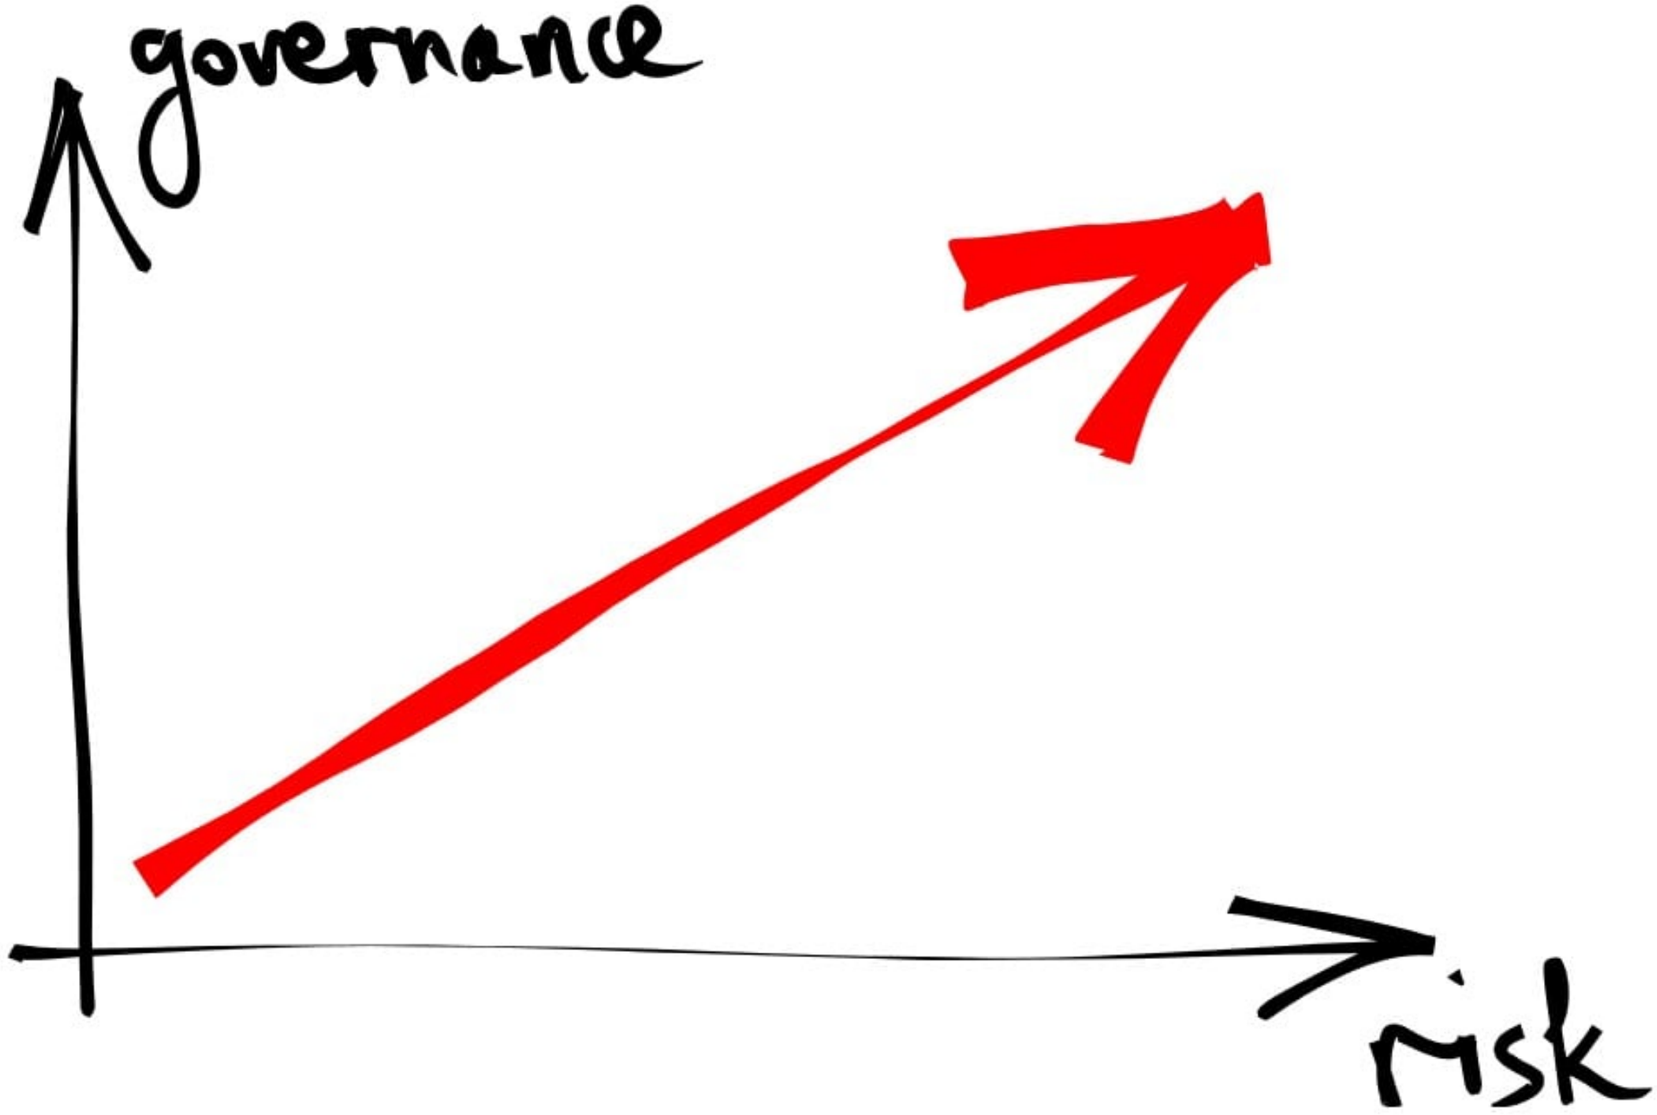
governance
risk

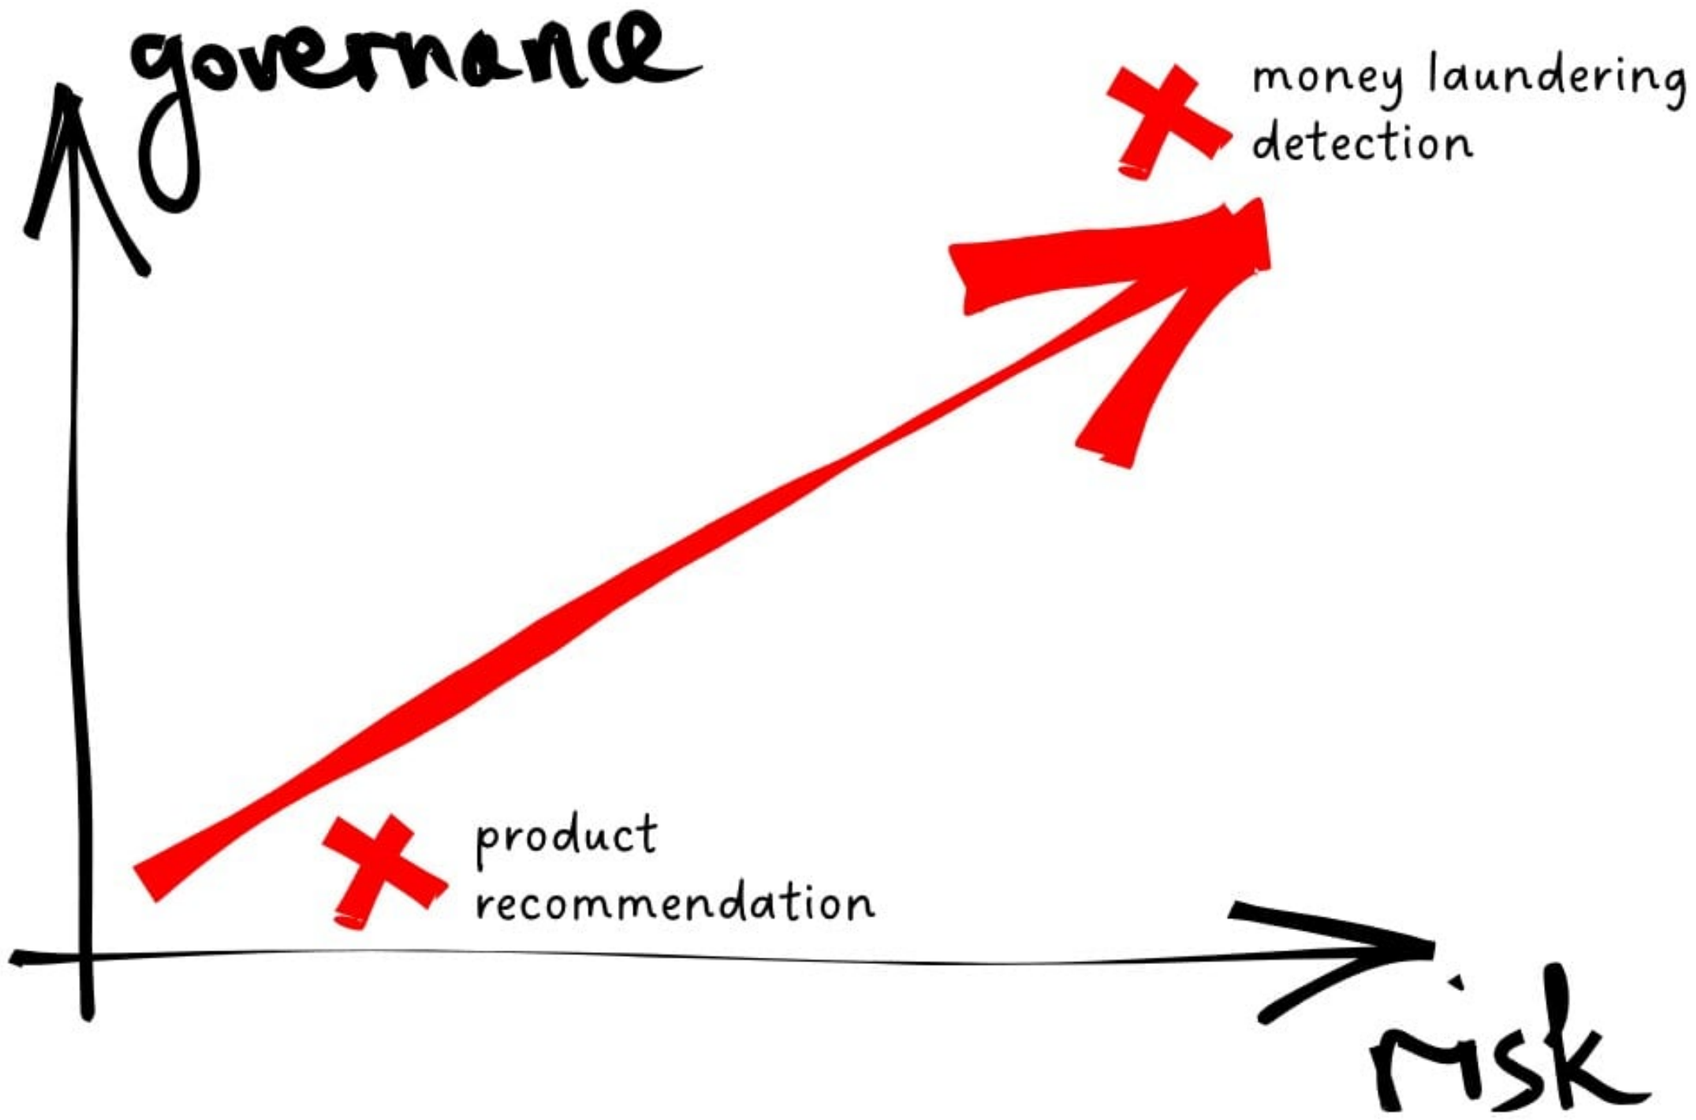
governance
money laundering
detection
product
recommendation
risk

Risk categories

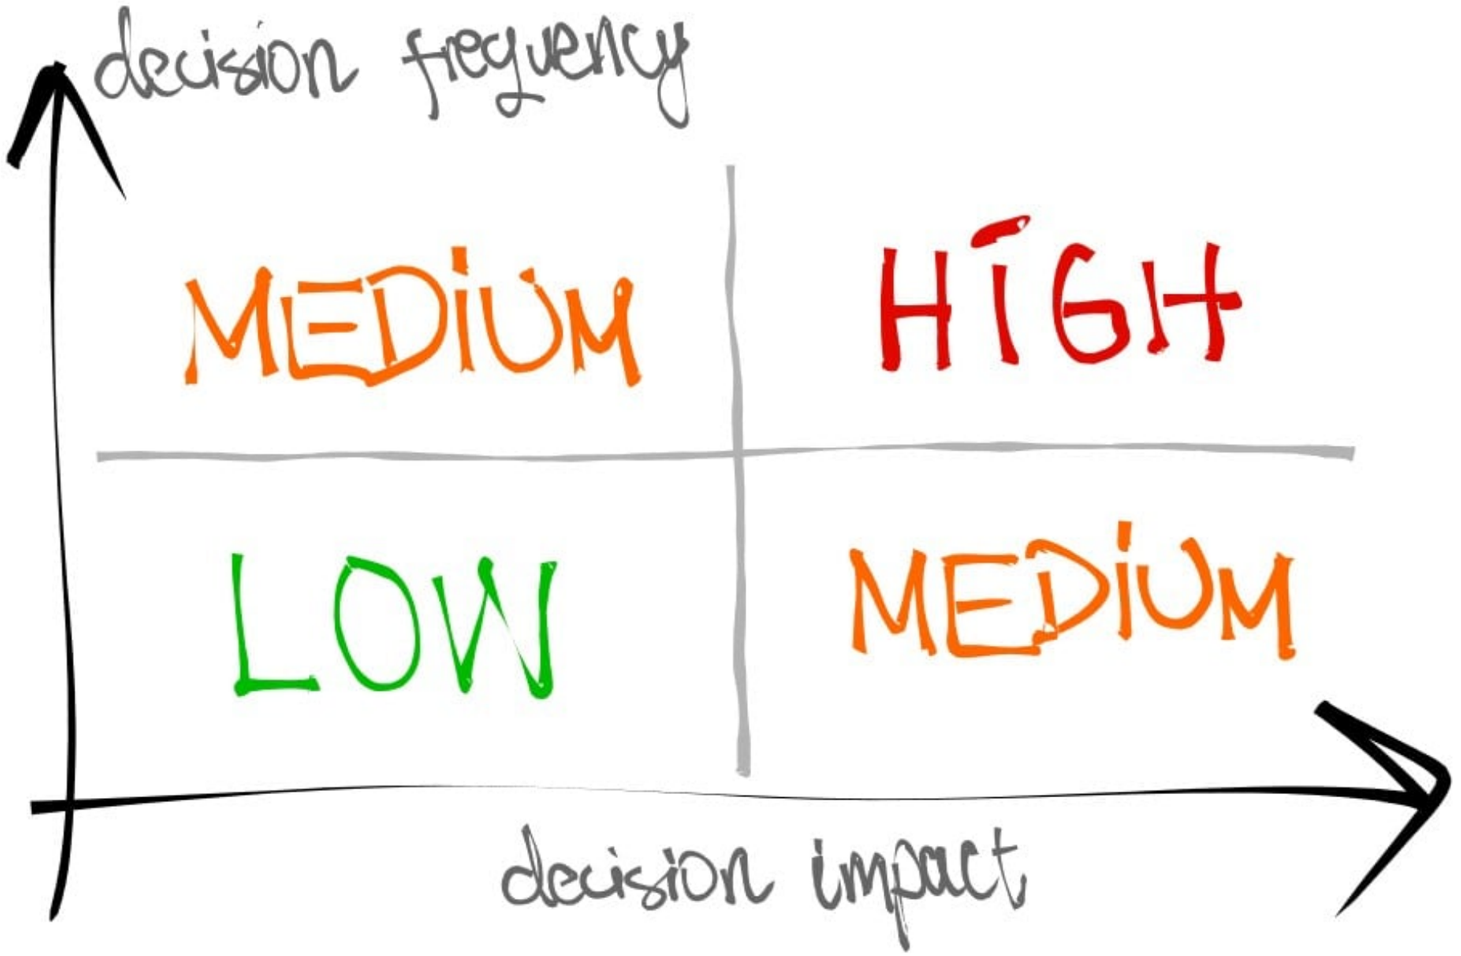

IMPACT

- Financial
- Reputational
- ...

Summary

- Governance means extra steps, but alternative is anarchy.
- "Launching many models fast" is not our goal, but generating business value.
- Reckless ML --> more damage than benefit.
- More models, more obvious the benefit of governance becomes.

Let's practice!

MLOPS DEPLOYMENT AND LIFE CYCLING

Wrap up

MLOPS DEPLOYMENT AND LIFE CYCLING

Nemanja Radojkovic
Senior Machine Learning Engineer

Principles above all

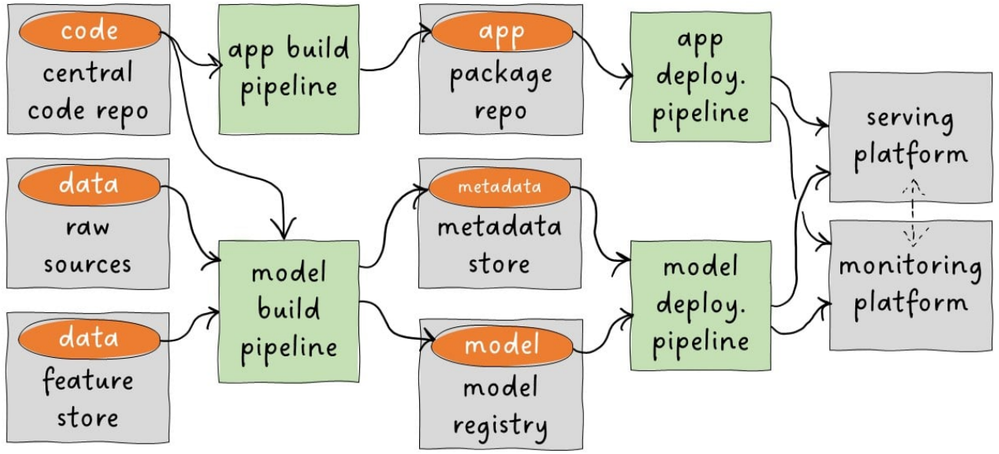

#automation

#collaboration

#efficiency

#transparency

#user-satisfaction

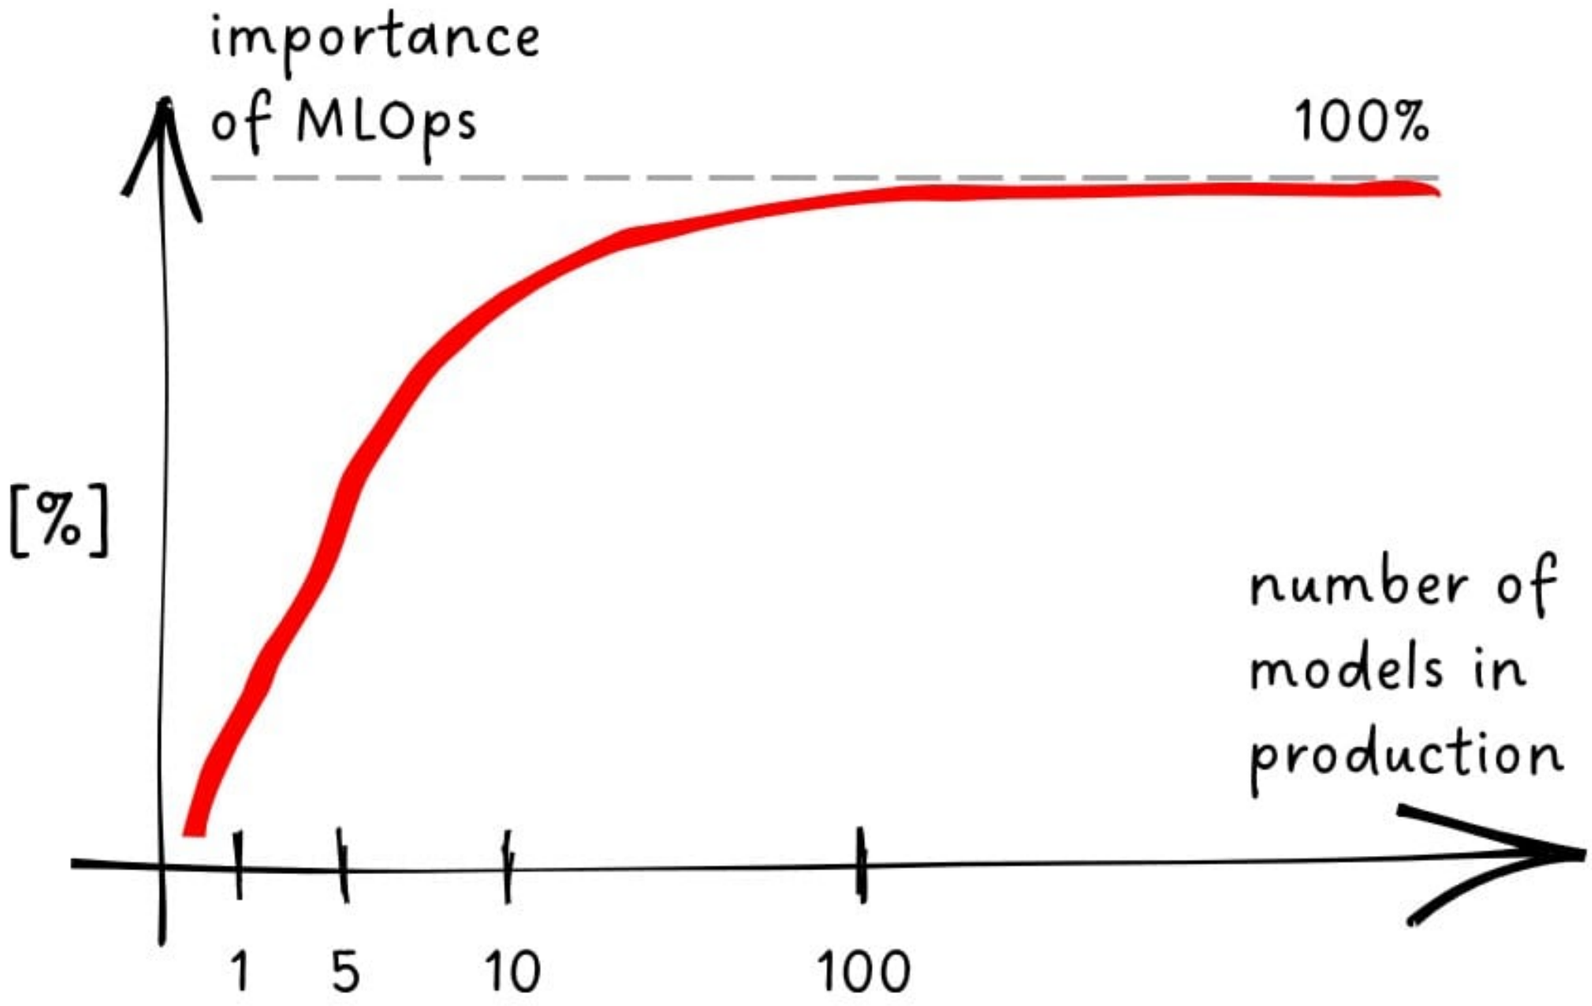
importance
of MLOps
100%
[%]
number of
models in
production
1
5
10
100

Start small and grow

Your MLOps framework today

Your MLOps framework in five years

Thank you!

MLOPS DEPLOYMENT AND LIFE CYCLING