

Image Classification Using **LeNet-5 CNN Architecture** with the CIFAR-10 Dataset

```
# Step 1: Importing the necessary libraries.
import tensorflow as tf
from tensorflow.keras import datasets, layers, models
from tensorflow.keras.utils import to_categorical
import matplotlib.pyplot as plt

# Step 2: Loading and pre-processing of the CIFAR-10 dataset.
(train_images, train_labels), (test_images, test_labels) = datasets.cifar10.load_data()

# Step 3: Normalizing pixel values to be between 0 and 1.
train_images = train_images / 255.0
test_images = test_images / 255.0

# Step 4: One-hot encode the labels.
# This step is required to use the loss function "categorical_crossentropy".
train_labels = to_categorical(train_labels, 10)
test_labels = to_categorical(test_labels, 10)

# Step 5: Defining the class names for CIFAR-10 images.
class_names = ['airplane', 'automobile', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck']

# Step 6: Visualizing a few training images from the CIFAR-10 dataset.
# plt.figure(figsize=(10, 10))
# for i in range(25):
#     plt.subplot(5, 5, i + 1)
#     plt.xticks([])
#     plt.yticks([])
#     plt.grid(False)
#     plt.imshow(train_images[i])
#     plt.xlabel(class_names[train_labels[i].argmax()]) # Using argmax to get the label index
# plt.show()

# Step 7: Building the CNN model (LeNet-5 CNN Architecture).
model = models.Sequential([
    layers.Conv2D(6, (5, 5), activation='tanh', input_shape=(32, 32, 3)),
    layers.AveragePooling2D((2, 2)),
    layers.Conv2D(16, (5, 5), activation='tanh'),
    layers.AveragePooling2D((2, 2)),
    layers.Conv2D(120, (5, 5), activation='tanh'),
    layers.Flatten(),
    layers.Dense(84, activation='tanh'),
    layers.Dense(10, activation='softmax')
])

# Step 8: Printing the model summary.
model.summary()

# Step 9: Compiling the CNN model.
model.compile(optimizer='adam', # Adam uses a default learning rate of 0.001
              loss='categorical_crossentropy',
              metrics=['accuracy'])

# Step 10: Training the CNN model.
history = model.fit(train_images, train_labels, epochs=10, batch_size=32, validation_data=(test_images, test_labels))

# Step 11: Evaluating the performance of the CNN model.
test_loss, test_acc = model.evaluate(test_images, test_labels, verbose=2)
print(f'\nTest accuracy is: {test_acc}')
```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
conv2d (Conv2D)	(None, 28, 28, 6)	456
average_pooling2d (Average Pooling2D)	(None, 14, 14, 6)	0
conv2d_1 (Conv2D)	(None, 10, 10, 16)	2416
average_pooling2d_1 (Average Pooling2D)	(None, 5, 5, 16)	0
conv2d_2 (Conv2D)	(None, 1, 1, 120)	48120
flatten (Flatten)	(None, 120)	0
dense (Dense)	(None, 84)	10164

```

dense_1 (Dense)                (None, 10)                850

=====
Total params: 62006 (242.21 KB)
Trainable params: 62006 (242.21 KB)
Non-trainable params: 0 (0.00 Byte)

```

```

Epoch 1/10
1563/1563 [=====] - 16s 7ms/step - loss: 1.7944 - accuracy: 0.3648 - val_loss: 1.6707 - val_accuracy: 0.4111
Epoch 2/10
1563/1563 [=====] - 7s 5ms/step - loss: 1.5635 - accuracy: 0.4458 - val_loss: 1.4850 - val_accuracy: 0.4704
Epoch 3/10
1563/1563 [=====] - 9s 5ms/step - loss: 1.4360 - accuracy: 0.4908 - val_loss: 1.4489 - val_accuracy: 0.4941
Epoch 4/10
1563/1563 [=====] - 8s 5ms/step - loss: 1.3555 - accuracy: 0.5203 - val_loss: 1.3917 - val_accuracy: 0.5075
Epoch 5/10
1563/1563 [=====] - 8s 5ms/step - loss: 1.2883 - accuracy: 0.5429 - val_loss: 1.3451 - val_accuracy: 0.5264
Epoch 6/10
1563/1563 [=====] - 8s 5ms/step - loss: 1.2369 - accuracy: 0.5613 - val_loss: 1.3323 - val_accuracy: 0.5305
Epoch 7/10
1563/1563 [=====] - 8s 5ms/step - loss: 1.1920 - accuracy: 0.5775 - val_loss: 1.3340 - val_accuracy: 0.5314
Epoch 8/10
1563/1563 [=====] - 8s 5ms/step - loss: 1.1554 - accuracy: 0.5891 - val_loss: 1.3227 - val_accuracy: 0.5343
Epoch 9/10
1563/1563 [=====] - 8s 5ms/step - loss: 1.1226 - accuracy: 0.6000 - val_loss: 1.3260 - val_accuracy: 0.5386
Epoch 10/10
1563/1563 [=====] - 8s 5ms/step - loss: 1.0900 - accuracy: 0.6161 - val_loss: 1.3161 - val_accuracy: 0.5429
313/313 - 1s - loss: 1.3161 - accuracy: 0.5429 - 725ms/epoch - 2ms/step

```

Test accuracy is: 0.542900025844574

```

# Plotting training and validation accuracy
plt.figure(figsize=(12, 4))

plt.subplot(1, 2, 1)
plt.plot(history.history['accuracy'], label='Training Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.ylim([0, 1])
plt.legend(loc='lower right')
plt.title('Training and Validation Accuracy')
plt.grid(True)

# Plotting training and validation loss
plt.subplot(1, 2, 2)
plt.plot(history.history['loss'], label='Training Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.ylim([0, 2]) # Adjusted y-axis limit to better visualize the loss values
plt.legend(loc='upper right')
plt.title('Training and Validation Loss')
plt.grid(True)

plt.show()

```

